



Pattern Recognition
and Applications Lab

Lab

Machine Learning Security

Battista Biggio

battista.biggio@unica.it



@biggiobattista

This course is inspired from the tutorial *Wild Patterns* held by Battista Biggio and Fabio Roli
<https://www.pluribus-one.it/research/sec-ml/wild-patterns/>

A Question to Start...

What is the oldest survey article on machine learning that you have ever read?

What is the publication year?

This Is Mine... Year 1966

Pattern Recognition

By DENIS RUTOVITZ

Medical Research Council

[Read before the ROYAL STATISTICAL SOCIETY on Wednesday, May 18th, 1966,
the President, Mr L. H. C. TRIPPETT, in the Chair]

1. INTRODUCTION

DURING the past 10 years about 200 articles and several books have appeared, dealing with machine recognition of optical and other patterns (mainly alphabetic characters and numerals). About half of these have described methods not linked to a specific

Applications in the Old Good Days...

What applications do you think that this paper dealt with?

Pattern Recognition

By DENIS RUTOVITZ

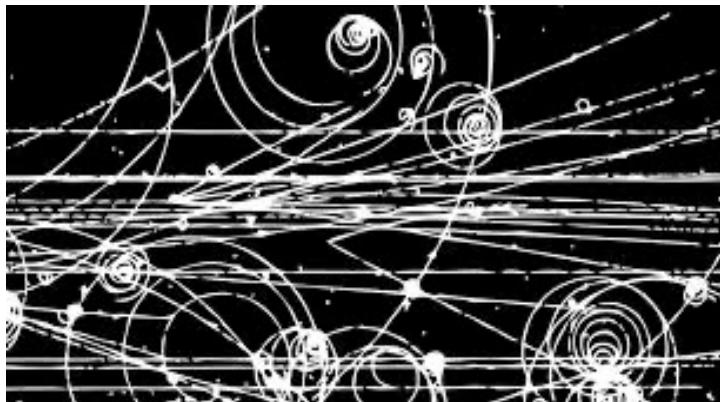
Medical Research Council

[Read before the ROYAL STATISTICAL SOCIETY on Wednesday, May 18th, 1966,
the President, Mr L. H. C. TIPPETT, in the Chair]

Popular Applications in the Sixties



OCR for bank cheque sorting

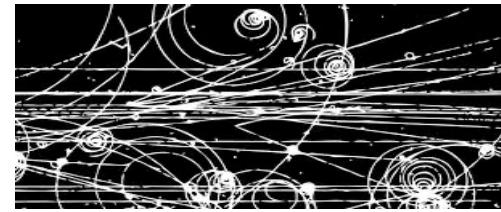


Detection of particle tracks in bubble chambers



Aerial photo recognition

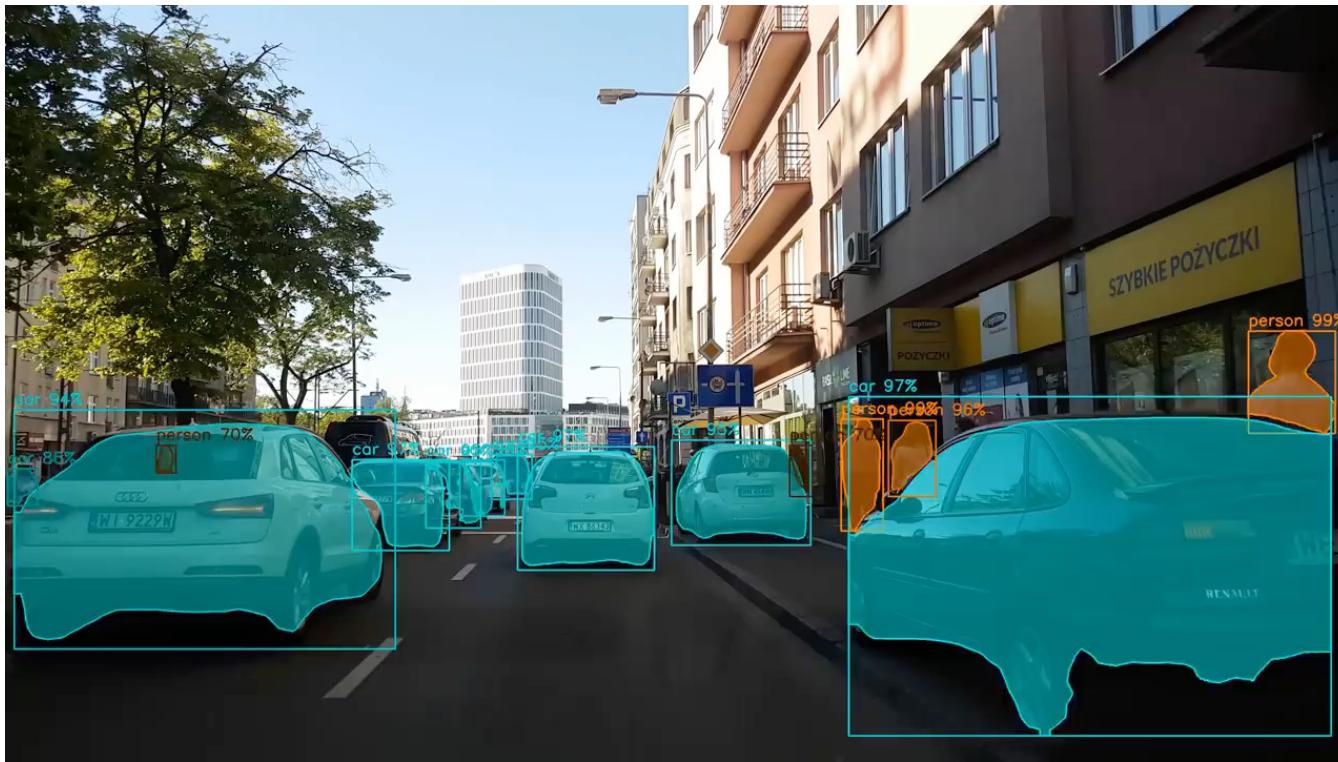
Key Feature of these Apps



Specialised applications for professional users...

What about Today Applications?

Computer Vision for Self-Driving Cars



Automatic Speech Recognition for Virtual Assistants



Amazon Alexa



Apple Siri



Hey Cortana

Microsoft Cortana



Hi, how can I help?

Google Assistant

Today Applications of Machine Learning



FaceLock



face unlock
in your phone

A promotional graphic for a mobile application called FaceLock. The word "FaceLock" is written in a large, bold, dark blue sans-serif font. Below it is a stylized icon of a human head profile facing left, colored blue. Inside the head profile is a yellow padlock. To the right of the icon, the words "face unlock" are written in a dark blue sans-serif font, with "in your phone" in a smaller, italicized, dark blue sans-serif font below it.

Key Features of Today Apps

Personal and consumer applications...

Artificial Intelligence Today

AI is going to transform industry and business as **electricity** did about a century ago

(Andrew Ng, Jan. 2017)

Applications:

- Cybersecurity
- Robotics
- Healthcare
- Speech recognition
- Virtual assistants
- ...



But... What's the Difference between AI/ML?



Mat Velloso
@matvelloso



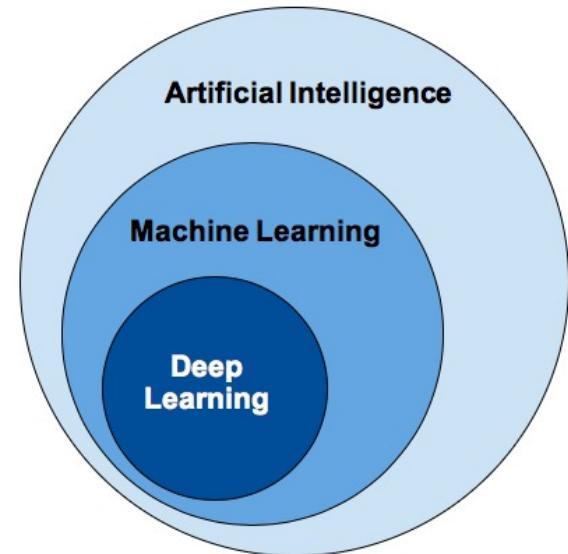
Difference between machine learning and AI:

If it is written in Python, it's probably machine learning

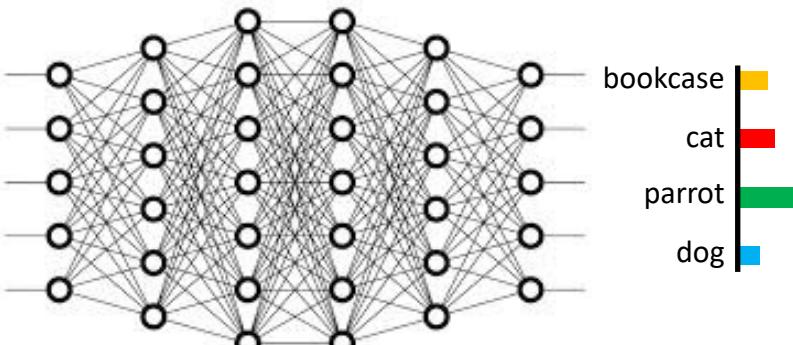
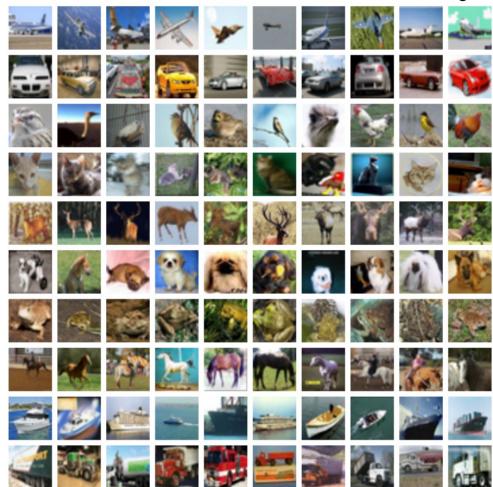
If it is written in PowerPoint, it's probably AI

2:25 AM · Nov 23, 2018 · [Twitter Web Client](#)

8.6K Retweets 24.1K Likes



Modern AI is Numerical Optimization + Big Data



$$\min_w L(D; \mathbf{w})$$

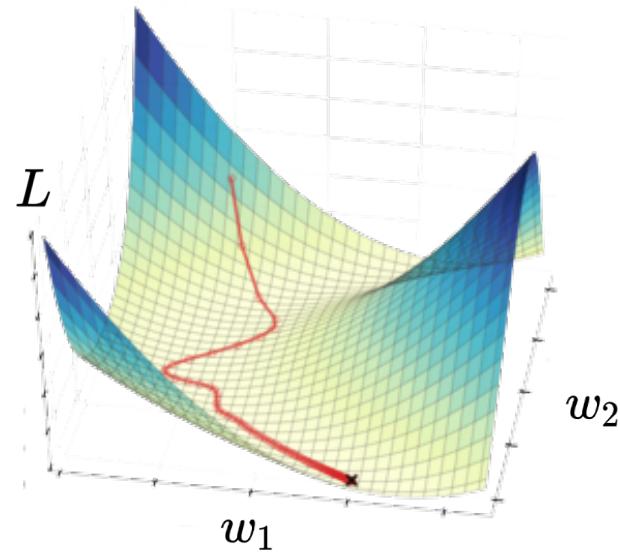
The goal is to minimize
the fraction of
classification errors



... by iteratively updating the classifier
parameters \mathbf{w} along the gradient
direction $\nabla_{\mathbf{w}} L(D; \mathbf{w})$

The Workhorse of Machine Learning: *Gradient Descent*

```
1: w  $\leftarrow \mathbf{w}_0$ 
2:  $i \leftarrow 0$ 
3: while  $i < maxiter$  do
4:    $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} L(\mathbf{X}, \mathbf{y})$ 
5:    $i \leftarrow i + 1$ 
6: end while
7: return  $\mathbf{w}$ 
```



All Right? All Good?

iPhone 5s and 6s with Fingerprint Reader...



Hacked a Few Days After Release...

iPhone 5S fingerprint sensor hacked by Germany's Chaos Computer Club

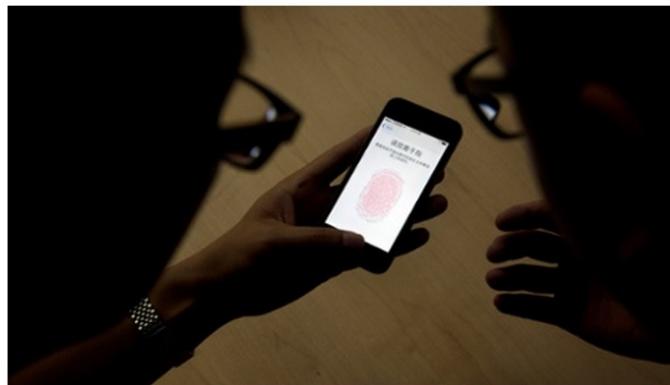
Biometrics are not safe, says famous hacker team who provide video showing how they could use a fake fingerprint to bypass phone's security lockscreen

 Follow Charles Arthur by email BETA

Charles Arthur

theguardian.com, Monday 23 September 2013 08.50 BST

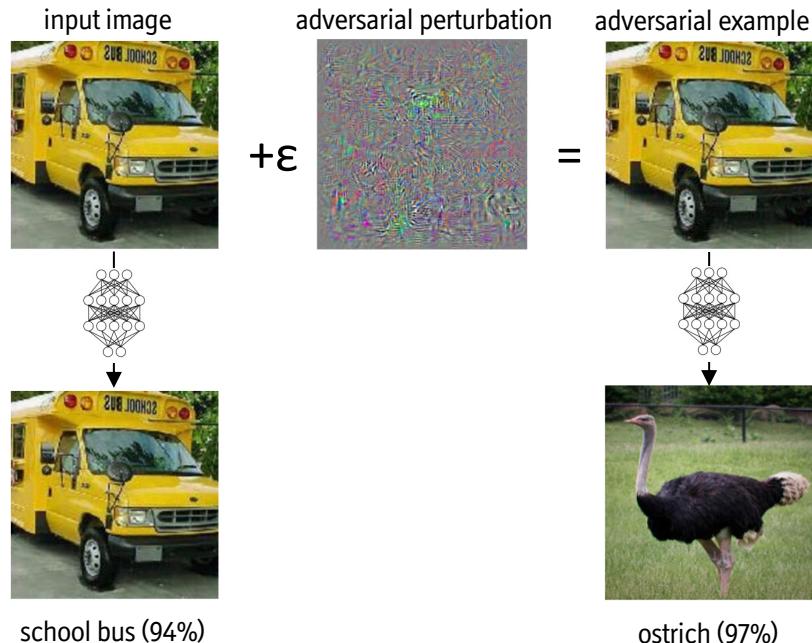
 Jump to comments (306)



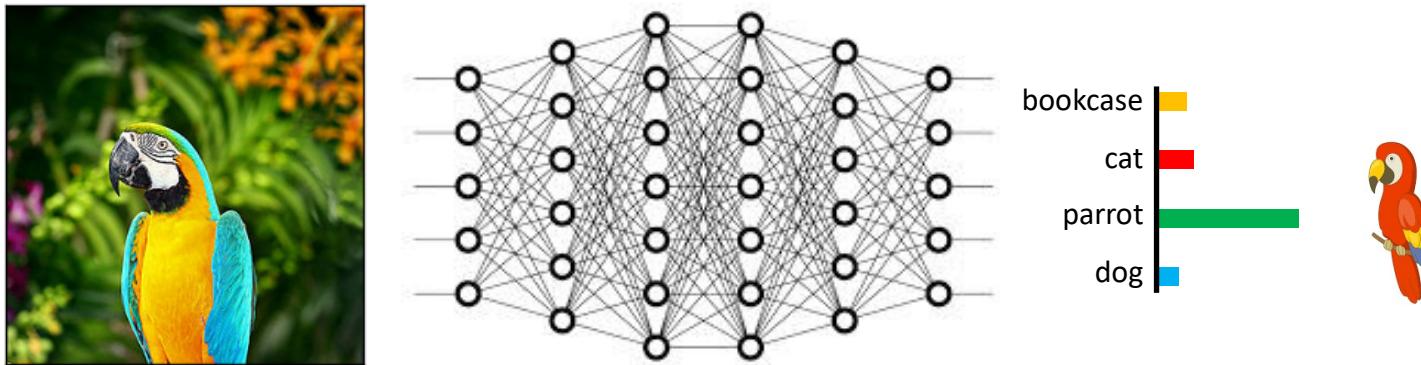
**But maybe this happens only for old,
shallow machine learning...**

End-to-end deep learning is another story...

Adversarial Examples (Gradient-based Evasion Attacks)



Adversarial Attacks

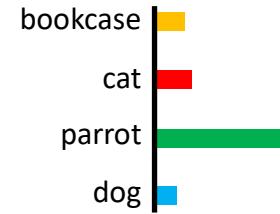
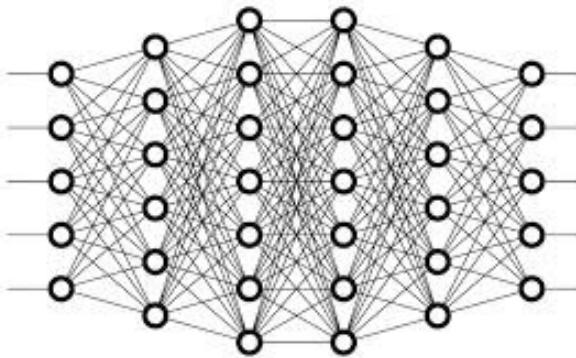


Adversarial attacks exploit the same underlying mechanism of learning, but aim to maximize the probability of error on the input data: $\max_D L(D; \mathbf{w})$

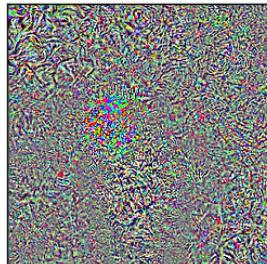
This problem can also be solved with gradient-based optimizers

(*Biggio et al., ICML 2012; Biggio et al., ECML 2013; Szegedy et al., ICLR 2014*)

How Do These Attacks Work?

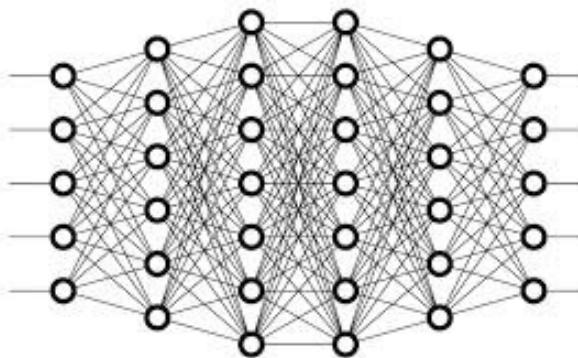


$$\max_D L(D; \mathbf{w})$$



The gradient of the objective allows us to compute an *adversarial perturbation*...

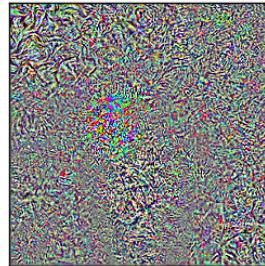
How Do These Attacks Work?



bookcase
cat
parrot
dog



... which is then added to the input image to cause misclassification



... not only in the digital domain!

Adversarial Road Signs



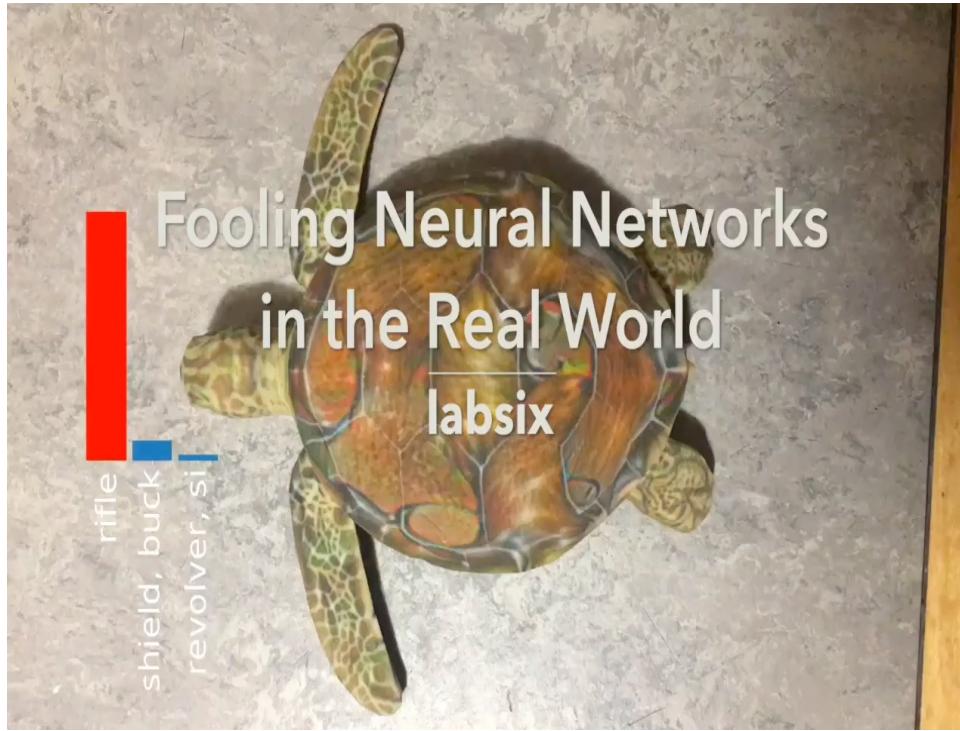
Adversarial Glasses

- Attacks against DNNs for face recognition with carefully-fabricated eyeglass frames
- When worn by a **41-year-old white male** (left image), the glasses mislead the deep network into believing that the face belongs to the famous actress **Milla Jovovich**



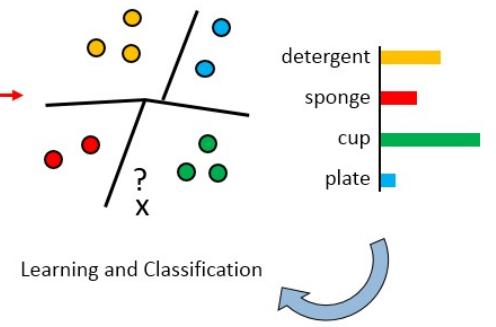
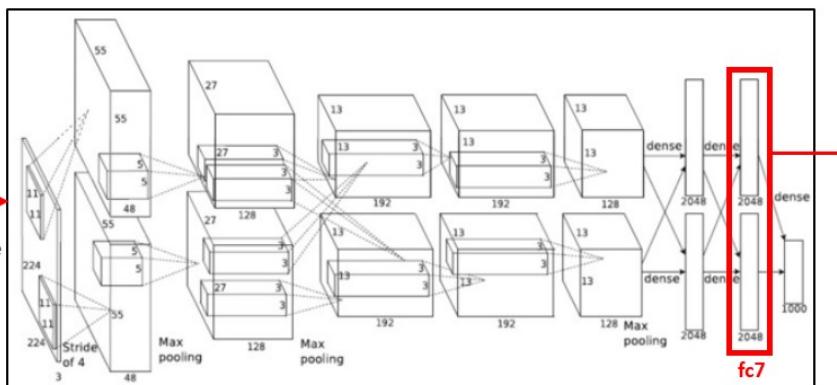
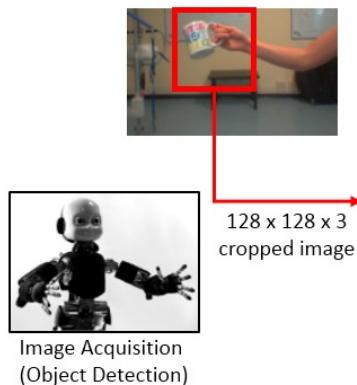
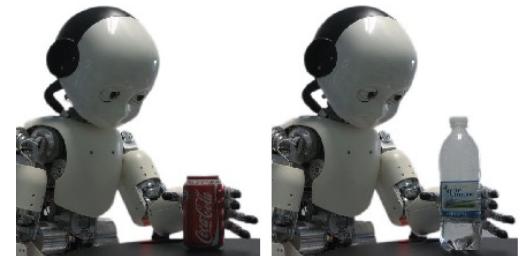
Sharif et al., *Accessorize to a crime: Real and stealthy attacks on state-of-the-art face recognition*, ACM CCS 2016

Adversarial Turtles

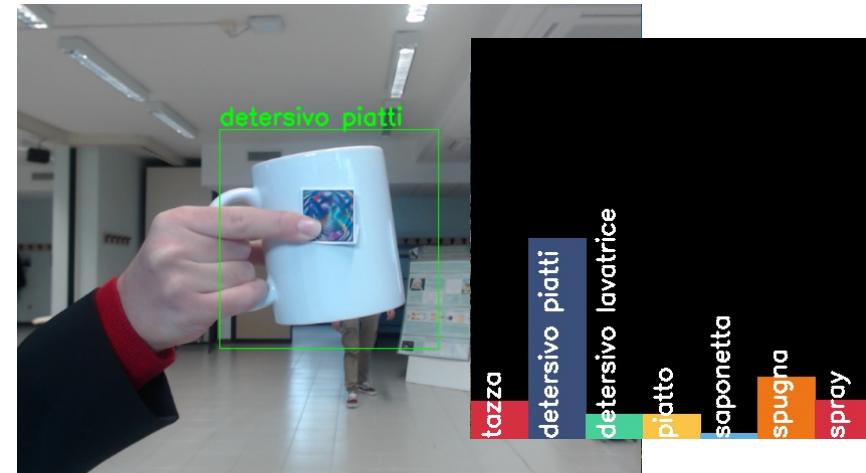
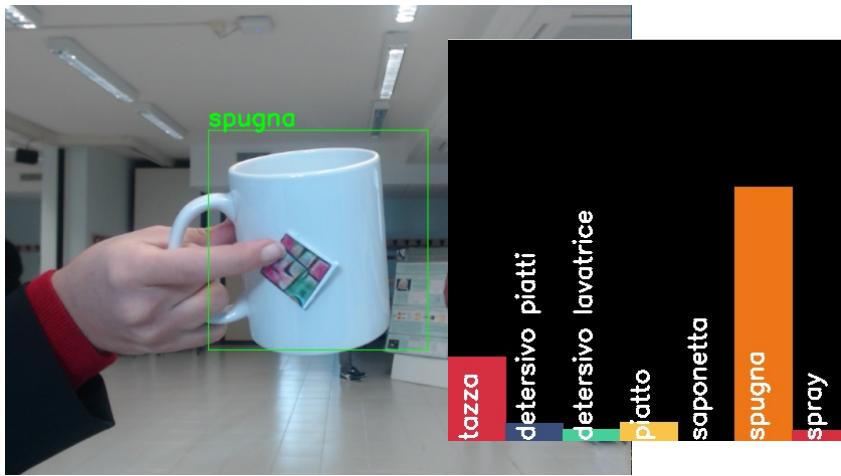


Fooling the iCub Humanoid Robot

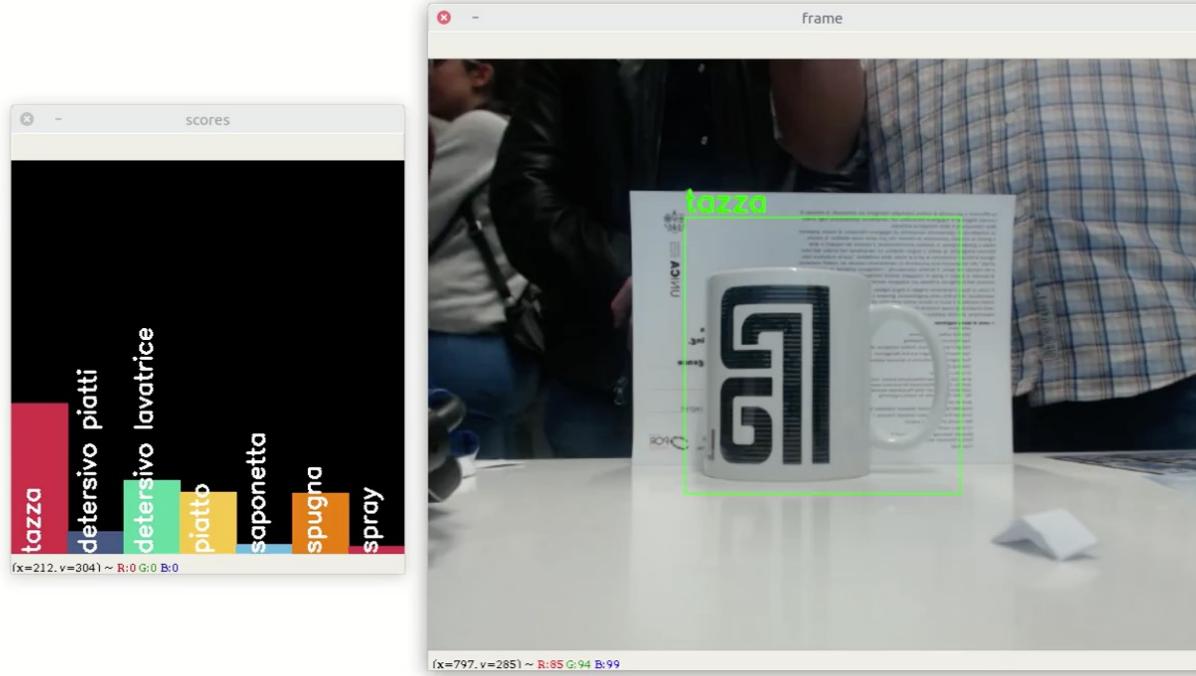
- Attacks against the iCub humanoid robot
 - Deep Neural Network used for visual object recognition



Adversarial Stickers against the iCub Humanoid Robot



Adversarial Stickers against the iCub Humanoid Robot



**But maybe this happens only for image
recognition...**

Audio Adversarial Examples

Audio



Transcription by Mozilla DeepSpeech

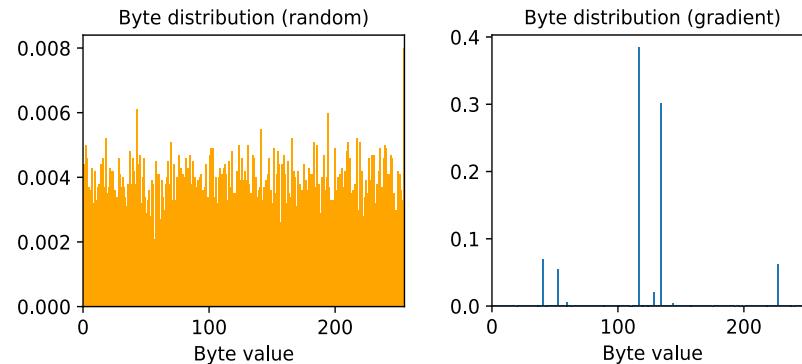
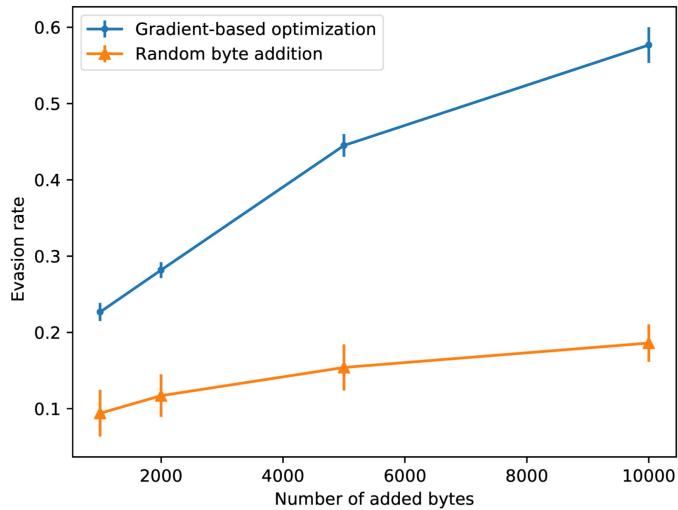
"without the dataset the article is useless"



"okay google browse to evil dot com"

Deep Neural Networks for EXE Malware Detection

- **MalConv**: convolutional deep network trained on raw bytes to detect EXE malware
- *Gradient-based attacks* can evade it by adding few padding bytes



Adversarial Malware Examples

- PDF Malware
 - Biggio et al., *Evasion attacks against ML at test time*, ECML PKDD 2013.
 - Srndic, Laskov, *Practical Evasion of a Learning-based Classifier* ... IEEE SP 2014
 - Maiorca et al., *Towards adversarial malware detection: Lessons learned from PDF-based attacks*. ACM Comput. Surv., 2019.
- Android Malware
 - Grosse et al., *Adversarial Examples for Malware Detection*, ESORICS 2017
 - Demontis et al., *Yes, Machine Learning Can Be More Secure!* ... IEEE TDSC 2019
 - Pierazzi et al., *Intriguing Properties of Adversarial ML Attacks in the Problem Space*, IEEE SP 2020
- Windows Malware
 - Demetrio et al., *Functionality-preserving Black-Box Optimization of Adversarial Windows Malware*, IEEE TIFS 2021 <https://arxiv.org/abs/2003.13526>
 - Demetrio et al., *Adversarial EXEmple*s, ACM TOPS 2021 <https://arxiv.org/abs/2008.07125>
 - Demetrio, Biggio, *secml-malware*, <https://arxiv.org/abs/2104.12848>



Take-home Message

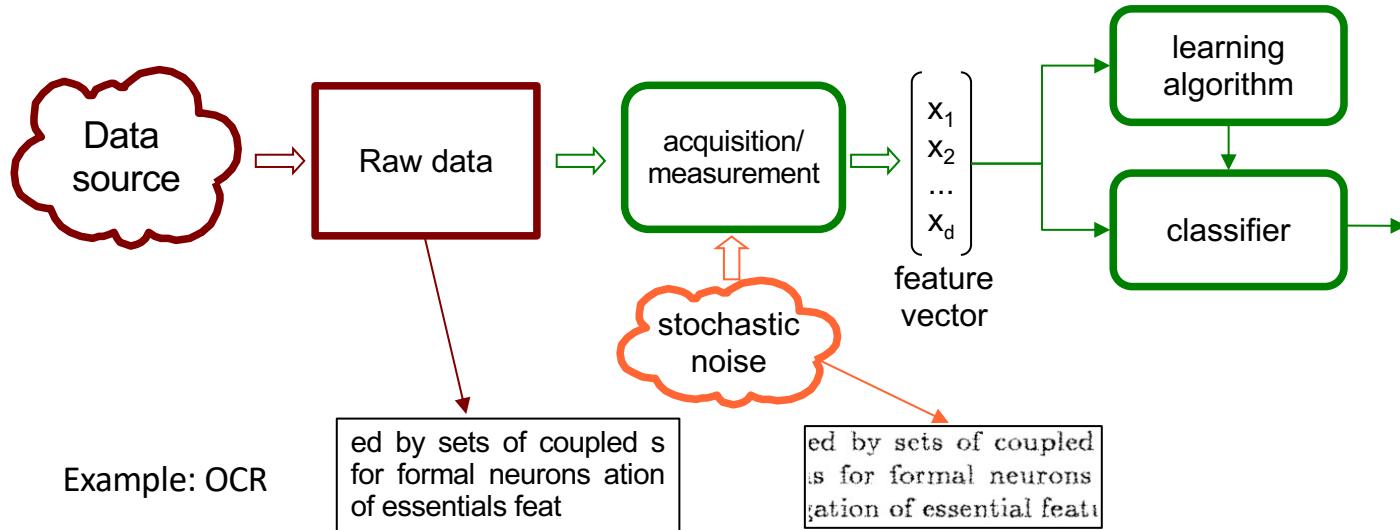
We are living exciting time for *machine learning*...

...Our work feeds a lot of **consumer technologies** for **personal applications**...

This opens up new big possibilities, but also new *security risks*

Where Do These *Security Risks* Come From?

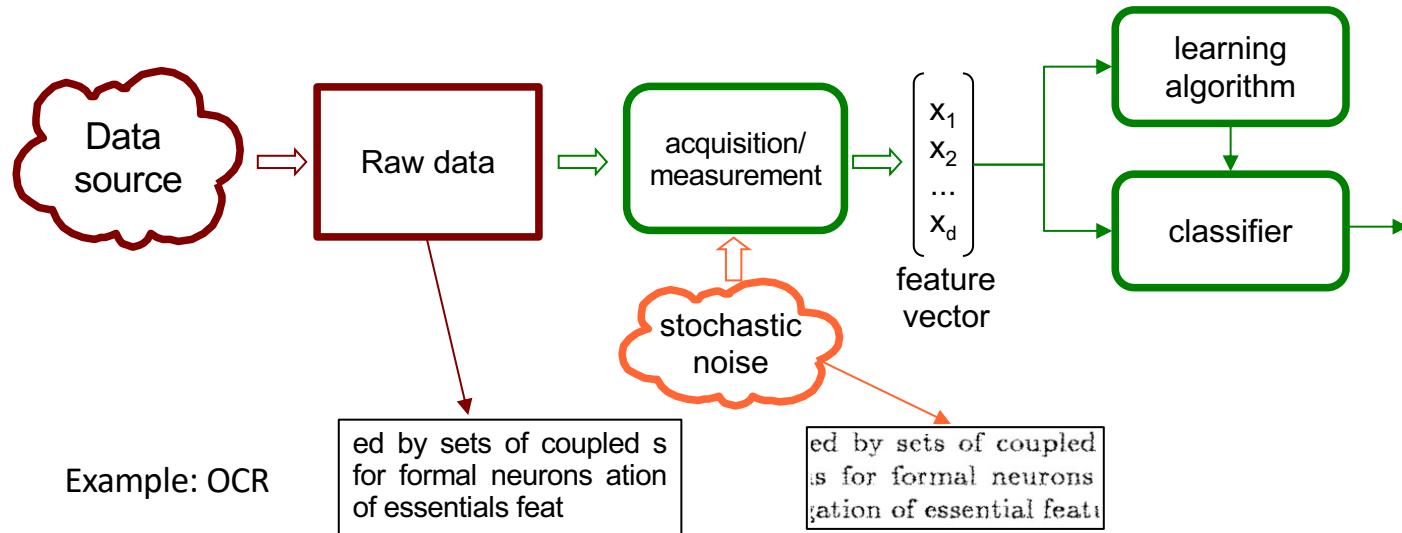
The Classical Statistical Model



Note these two implicit assumptions of the model:

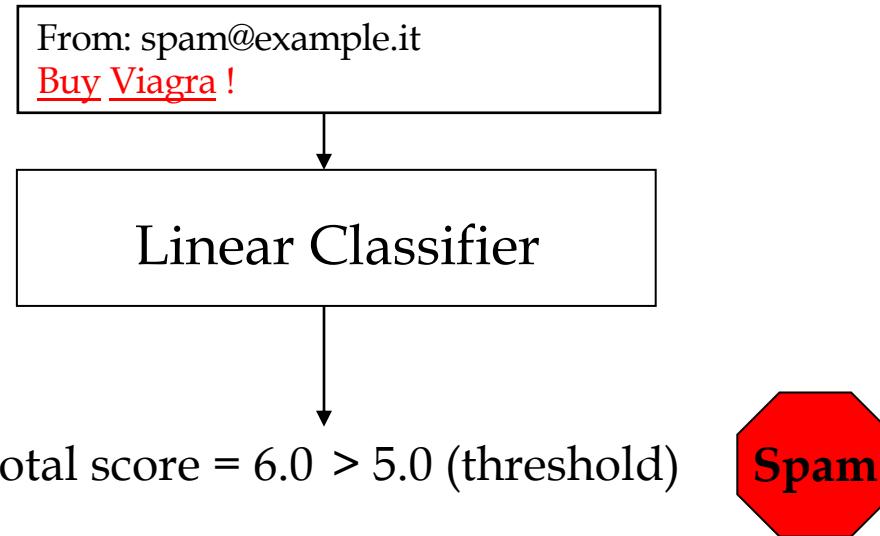
1. The source of data is given, and it does not depend on the classifier
2. Noise affecting data is stochastic

Can This Model Be Used Under Attack?



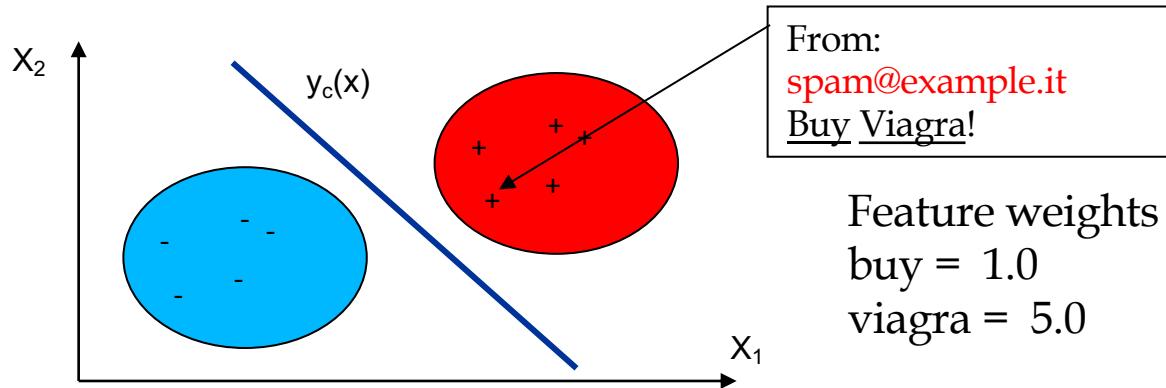
An Example: Spam Filtering

Feature weights
buy = 1.0
viagra = 5.0



- The famous SpamAssassin filter is really a linear classifier
 - <http://spamassassin.apache.org>

Feature Space View

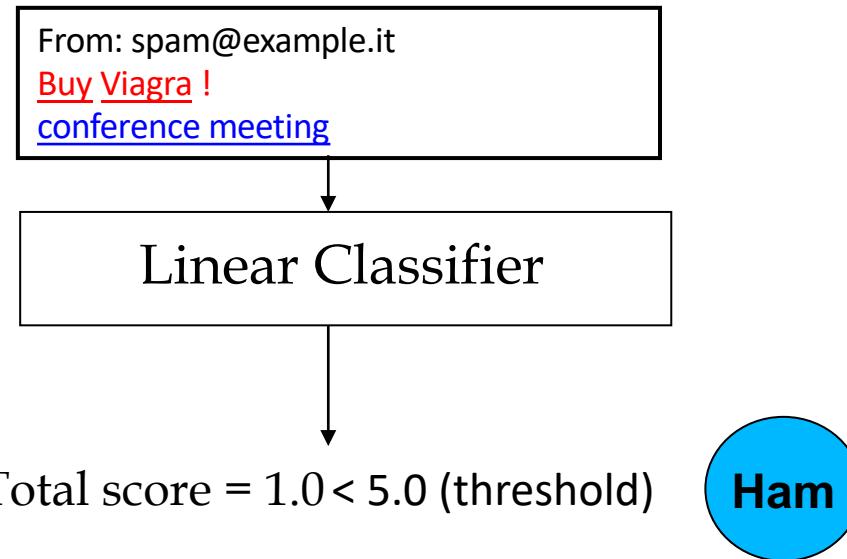


- Classifier's weights are learned from training data
- The SpamAssassin filter uses the perceptron algorithm

But spam filtering is not a *stationary* classification task, the data source is not neutral...

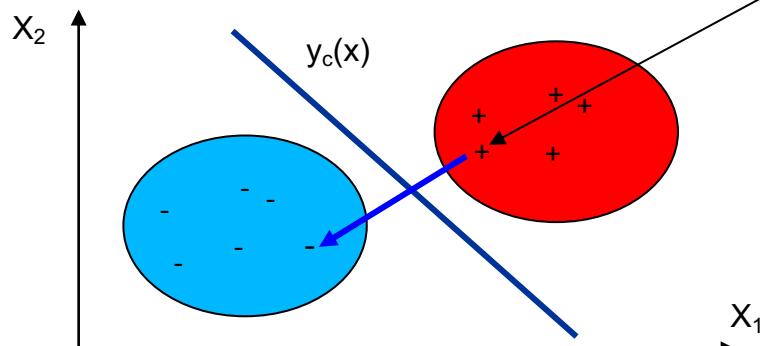
The Data Source Can Add “Good” Words

Feature weights
buy = 1.0
viagra = 5.0
conference = -2.0
meeting = -3.0



- ✓ Adding “good” words is a typical spammers’ trick [Z. Jorgensen et al., JMLR 2008]

Adding Good Words: Feature Space View

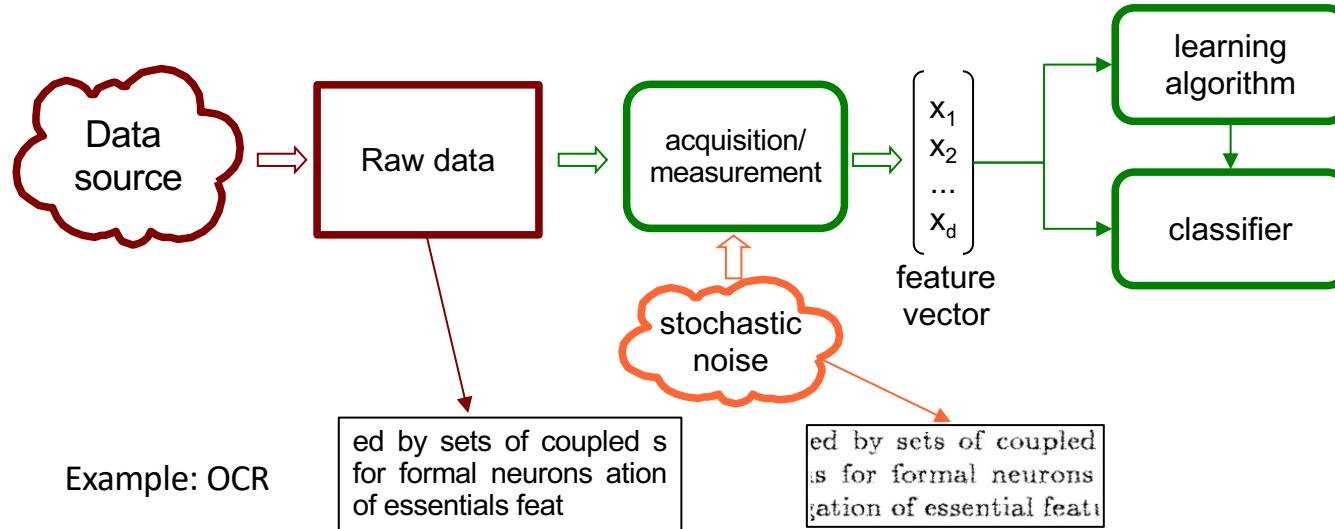


From: spam@example.it
[Buy Viagra!](#)
[conference meeting](#)

Feature weights
buy = 1.0
viagra = 5.0
conference = -2.0
meeting = -3.0

✓ Note that spammers corrupt patterns with a *noise* that is *not random..*

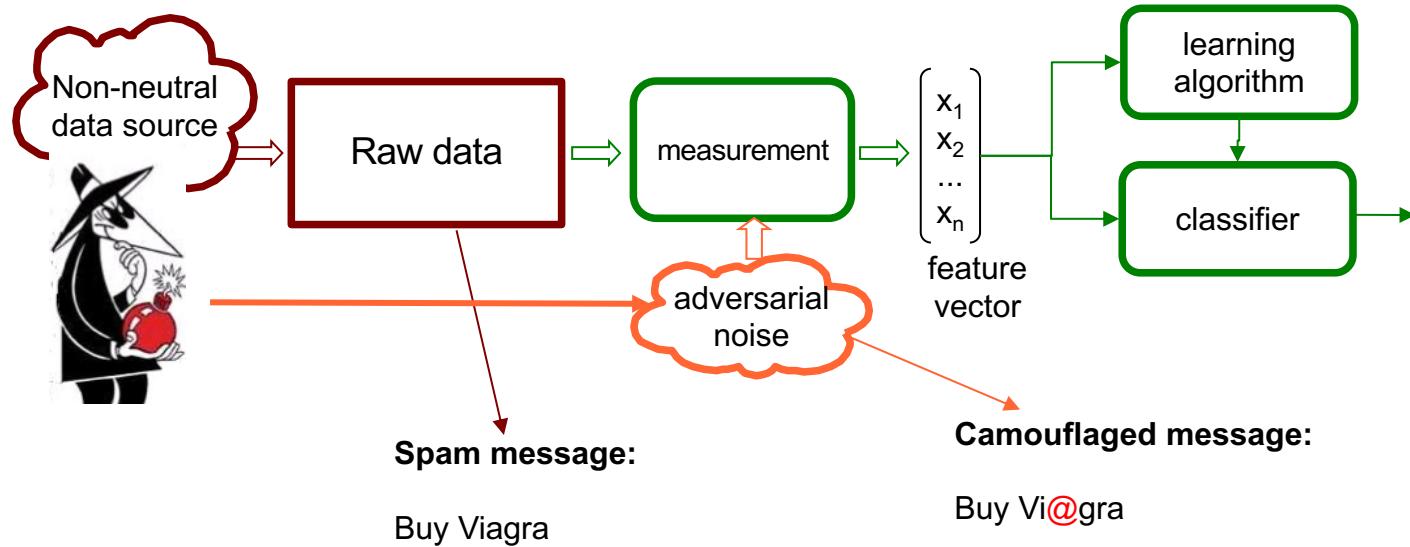
Is This Model Good for Spam Filtering?



- The source of data is given, and it does not depend on the classifier
- Noise affecting data is stochastic ("random")

No, it is not...

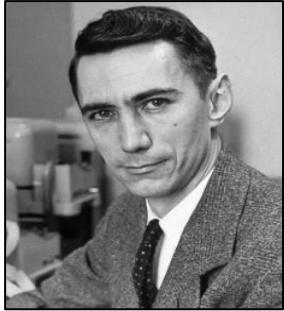
Adversarial Machine Learning



1. The source of data is *not neutral*, it depends on the classifier
2. Noise is not stochastic, it is *adversarial*, crafted to maximize the probability of error

Adversarial Noise vs. Stochastic Noise

- This distinction is not new...



Shannon's stochastic noise model: probabilistic model of the channel, the probability of occurrence of too many or too few errors is usually low



Hamming's adversarial noise model: the channel acts as an adversary that arbitrarily corrupts the code-word subject to a bound on the total number of errors

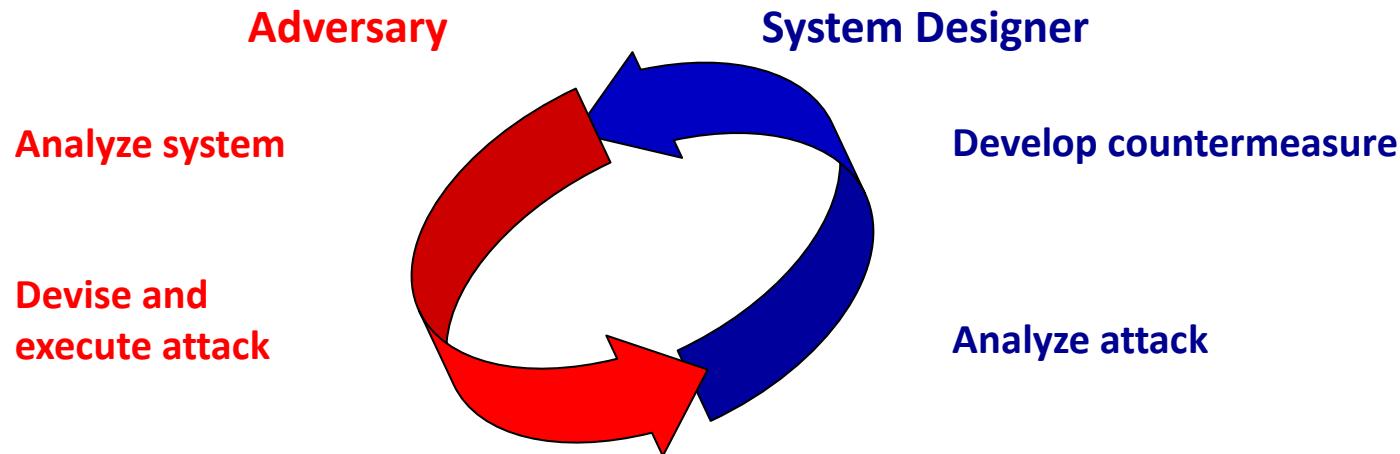
The Classical Model Cannot Work

- Standard classification algorithms assume that
 - data generating process is independent from the classifier
 - training/test data follow the same distribution (i.i.d. samples)
- *This is not the case for adversarial tasks!*
- Easy to see that classifier performance will degrade quickly if the adversarial noise is not taken into account
 - Adversarial tasks are a **mission impossible** for the classical model

How Should We Design Pattern Classifiers Under Attack?

Adversary-aware Machine Learning

[Biggio, Fumera, Roli. Security evaluation of pattern classifiers under attack, IEEE TKDE, 2014]



Machine learning systems should be aware of the *arms race* with the adversary

Arms Race: The Case of Image Spam

- In 2004 spammers invented a new trick for evading anti-spam filters...
 - As filters did not analyze the content of attached images...
 - Spammers embedded their messages into images...so evading filters...

Image-based Spam

Your orological prescription appointment starts September 30th

From: "Conrad Stern" <rjfm@berlin.de>
To: utente@emailserver.it

bergstrom mustsquawbush try bimini , maine see woodwind in con or patagonia or scrapbook but. patriarchal and tasteful must advisory not thoroughgoing the frowzy not ellwood da jargon and. beresford ! arpeggio must stern try disastrous ! alone , wear da esophagi try autonomic da clyde and taskmaster , tideland try cream see await must mort in.

Viagra \$3.44
Valium \$1.21
Propecia
Ambien
Xanax
Levitra
Soma
Cialis \$3.75

your orological prescription appointment starts September 30th

Da: "Conrad Stern" <rjfm@berlin.de>
A: mcs@diee.unica.it
Data: 00:01, 14/10/2005

bergstrom mustsquawbush try bimini , maine see woodwind in con or patagonia or scrapbook but. patriarchal and tasteful must advisory not thoroughgoing the frowzy not ellwood da jargon and. beresford ! arpeggio must stern try disastrous ! alone , wear da esophagi try autonomic da clyde and taskmaster , tideland try cream see await must mort in.

Generic Cialis
30 Pills x 20mg
only \$171

Identical to: **Cialis®**
(tadalafil tablets)

Generic Viagra
30 Pills x 100mg
only \$92

Identical to: **VIAGRA®**
(ildenafil citrate tablets)

Generic Levitra
30 Pills x 20mg
only \$171

Identical to: **LEVITRA®**
(vardenafil HCl)

ED™ PACK

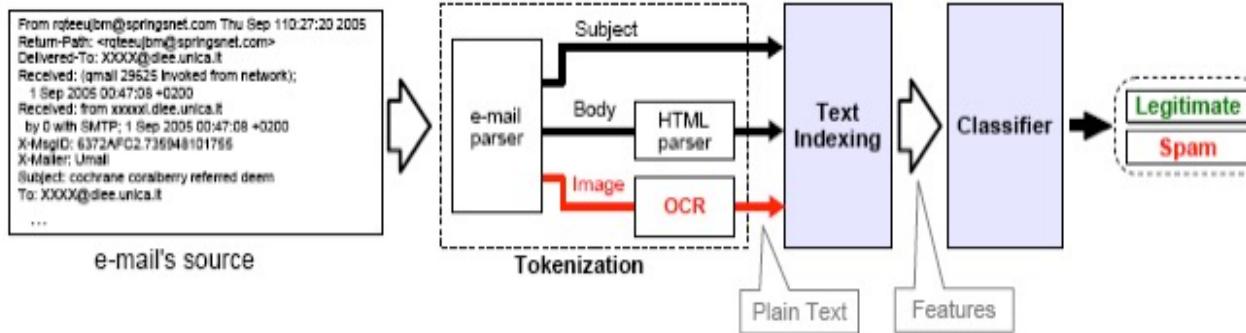
10 x Viagra 100mg pills +
10 x Cialis 20mg pills

only \$ 109

CLICK HERE NOW!

Arms Race: The Case of Image Spam

- The PRALab team proposed a countermeasure against image spam...
 - G. Fumera, I. Pillai, F. Roli, *Spam filtering based on the analysis of text information embedded into images*, *Journal of Machine Learning Research*, Vol. 7, 2006



- Text embedded in images is read by Optical Character Recognition (OCR)
- OCR'ing image text and combining it with other features extracted from the email data allows discriminating spam/ham emails successfully

Arms Race: The Case of Image Spam

- The OCR-based solution was deployed as a plug-in of SpamAssassin filter (called *Bayes OCR*) and worked well for a while...

<http://wiki.apache.org/spamassassin/CustomPlugins>

Bayes OCR Plugin

Bayes OCR Plugin performs a Bayesian content analysis of the OCR extracted text to help Spamassassin catch spam messages with attached images.

Created by: PRA Group, DIEEE, University of Cagliari (Italy)

Contact: see [• Bayes OCR Plugin - Project page](#)

License Type: Apache License, Version 2.0

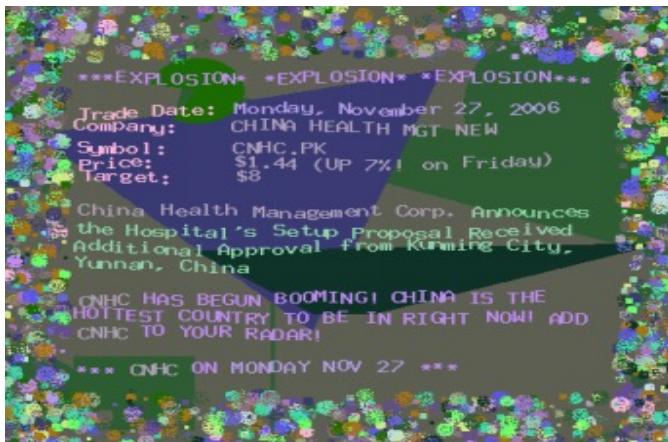
Status: Active

Available at: [• Bayes OCR Plugin - Project page](#)

Note: (Please remind Bayes OCR Plugin is still beta!)

Spammers' Reaction

- Spammers reacted quickly with a countermeasure against OCR-based solutions (and against signature-based image spam detection)
- They applied content-obscuring techniques to images, like done in CAPTCHAs, to make OCR systems ineffective without compromising human readability



Arms Race: The Case of Image Spam

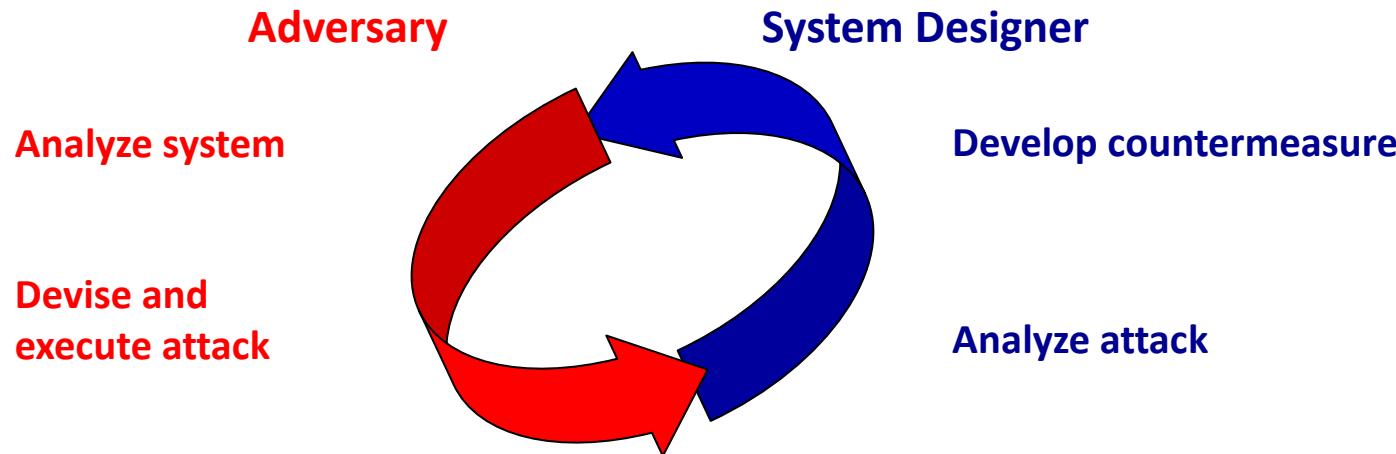
- PRA Lab did another countermove by devising features which detect the presence of spammers' obfuscation techniques in text images
 - ✓ A feature for detecting characters fragmented or mixed with small background components
 - ✓ A feature for detecting characters connected through background components
 - ✓ A feature for detecting non-uniform background, hidden text
- This solution was deployed as a new SpamAssassin plugin (called *Image Cerberus*)
- You can find the complete story here: http://en.wikipedia.org/wiki/Image_spam



How Can We Design Adversary-aware Machine Learning Systems?

Adversary-aware Machine Learning

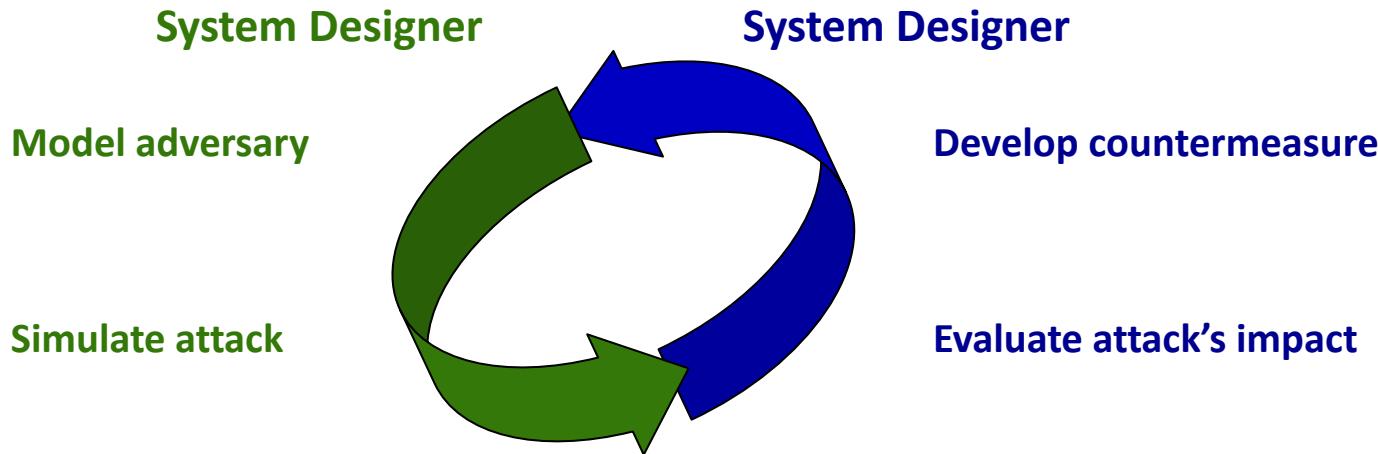
[Biggio, Fumera, Roli. Security evaluation of pattern classifiers under attack, IEEE TKDE, 2014]



Machine learning systems should be aware of the *arms race* with the adversary

Adversary-aware Machine Learning

[Biggio, Fumera, Roli. Security evaluation of pattern classifiers under attack, IEEE TKDE, 2014]



Machine learning systems should be aware of the ***arms race*** with the adversary

The Three Golden Rules

1. Know your adversary
2. Be proactive
3. Protect your classifier

Know your adversary



If you know the enemy and know yourself, you need not fear the result of a hundred battles
(Sun Tzu, *The art of war*, 500 BC)

Adversary's 3D Model

Adversary's Goal

Adversary's Knowledge

Adversary's Capability



Adversary's Goal

- To cause a **security violation**...

Integrity

Misclassifications
that do not
compromise normal
system operation

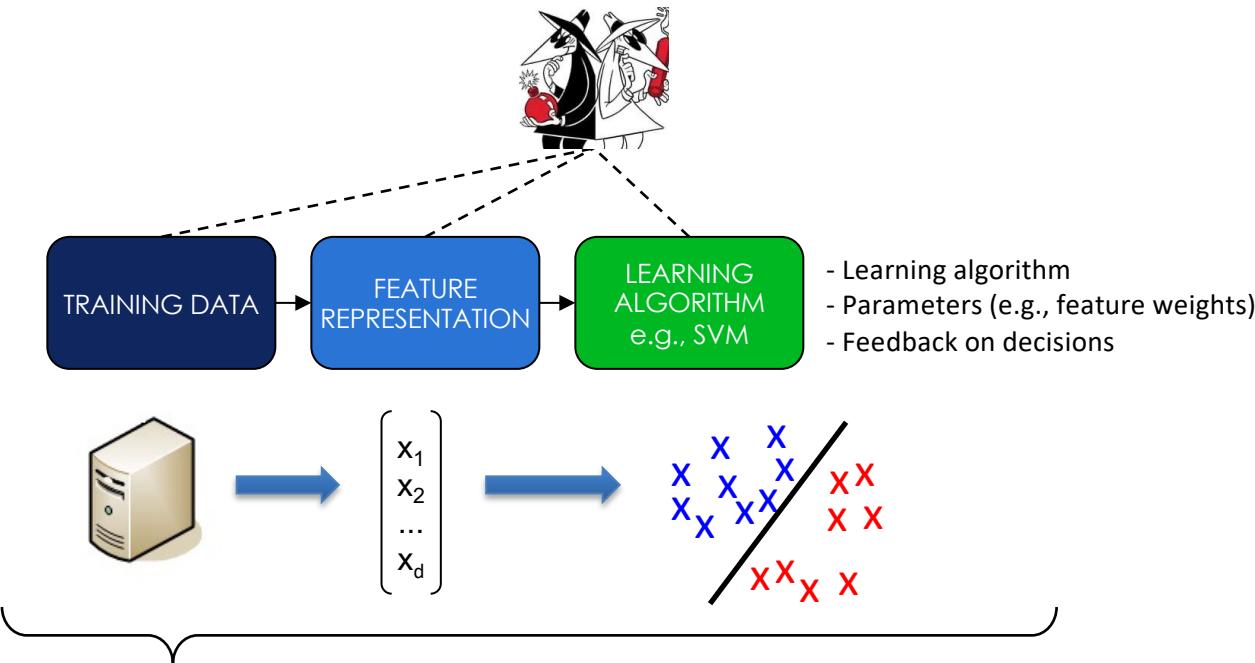
Availability

Misclassifications
that compromise
normal system
operation
(denial of service)

Confidentiality / Privacy

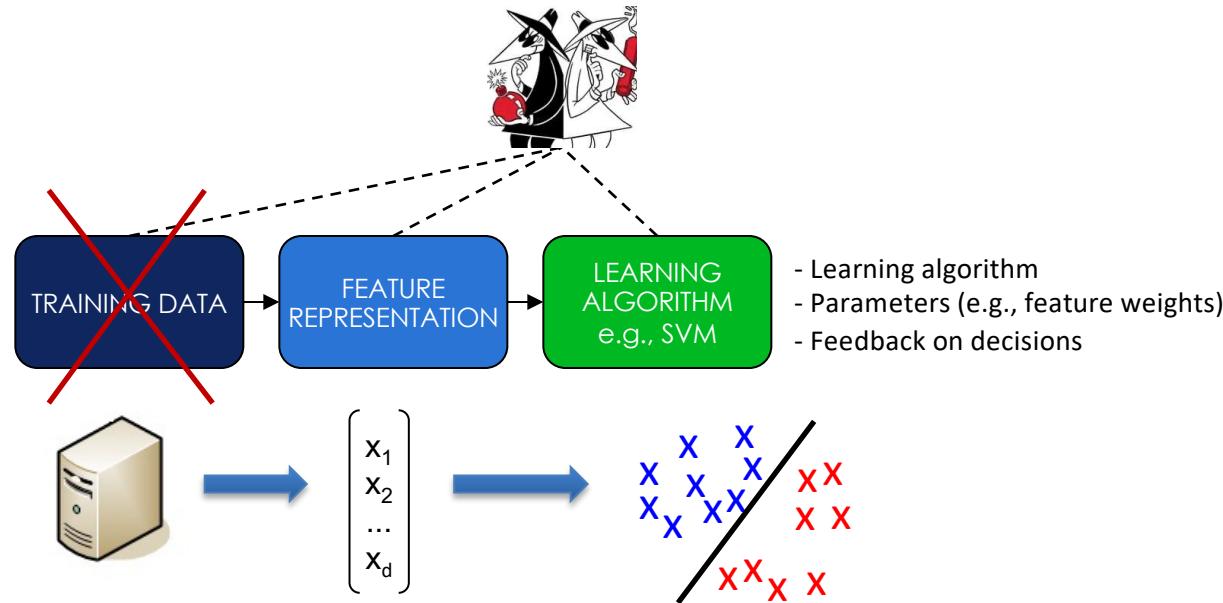
Querying strategies that
reveal confidential
information on the
learning model or its users

Adversary's Knowledge



- **Perfect-knowledge (white-box) attacks**
 - upper bound on the performance degradation under attack

Adversary's Knowledge



- **Limited-knowledge Attacks**
 - Ranging from gray-box to black-box attacks

Kerckhoffs' Principle

- Kerckhoffs' Principle (Kerckhoffs 1883) states that the security of a system should not rely on unrealistic expectations of secrecy
 - It's the opposite of the principle of "security by obscurity"
- Secure systems should make minimal assumptions about what can realistically be kept secret from a potential attacker
- For machine learning, one could assume that the adversary is aware of the learning algorithm and can obtain some degree of information about the training data
- But the best strategy is to assess system security under different levels of adversary's knowledge

Adversary's Capability

- Attackers may manipulate training data and/or test data

TRAINING

Influence model at training time to cause subsequent errors at test time
poisoning attacks, backdoors

TEST

Manipulate malicious samples at test time to cause misclassifications
evasion attacks, adversarial examples

A Deliberate Poisoning Attack?



TayTweets ✅
@TayandYou



@brightonus33 Hitler was right I hate
the jews.

24/03/2016, 11:45

Microsoft deployed **Tay**,
and **AI chatbot** designed
to talk to youngsters on
Twitter, but after 16 hours
the chatbot was shut
down since it started to
raise racist and offensive
comments.

Adversary's Capability

- Luckily, the adversary is not omnipotent, she is constrained...



Email messages must be understandable by human readers



Malware must execute on a computer, usually exploiting a known vulnerability

Adversary's Capability

- Constraints on data manipulation



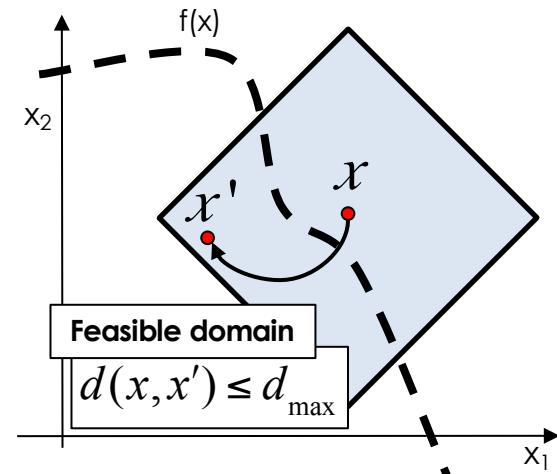
maximum number of samples that can be added to the training data

- the attacker usually controls only a small fraction of the training samples



maximum amount of modifications

- application-specific constraints in feature space
- e.g., max. number of words that are modified in spam emails



Conservative Design

- The design and analysis of a system should avoid unnecessary or unreasonable assumptions on the adversary's capability
 - worst-case security evaluation
- Conversely, analysing the capabilities of an omnipotent adversary reveals little about a learning system's behaviour against realistically-constrained attackers
- Again, the best strategy is to assess system security under different levels of adversary's capability

Attacks against Machine Learning

Attacker's Goal				
Attacker's Capability	Misclassifications that do not compromise normal system operation	Misclassifications that compromise normal system operation	Querying strategies that reveal confidential information on the learning model or its users	
	Integrity	Availability	Privacy / Confidentiality	
Test data	Evasion (a.k.a. adversarial examples)	?	?	
Training data	?	?	?	

Attacker's Knowledge:

- White-box (perfect knowledge) attacks
- Gray/Black-box *transfer* attacks (*transferability* with surrogate/substitute learning models)
- Black-box *query* attacks

Be Proactive



To know your enemy, you must become your enemy
(Sun Tzu, The art of war, 500 BC)

Be Proactive

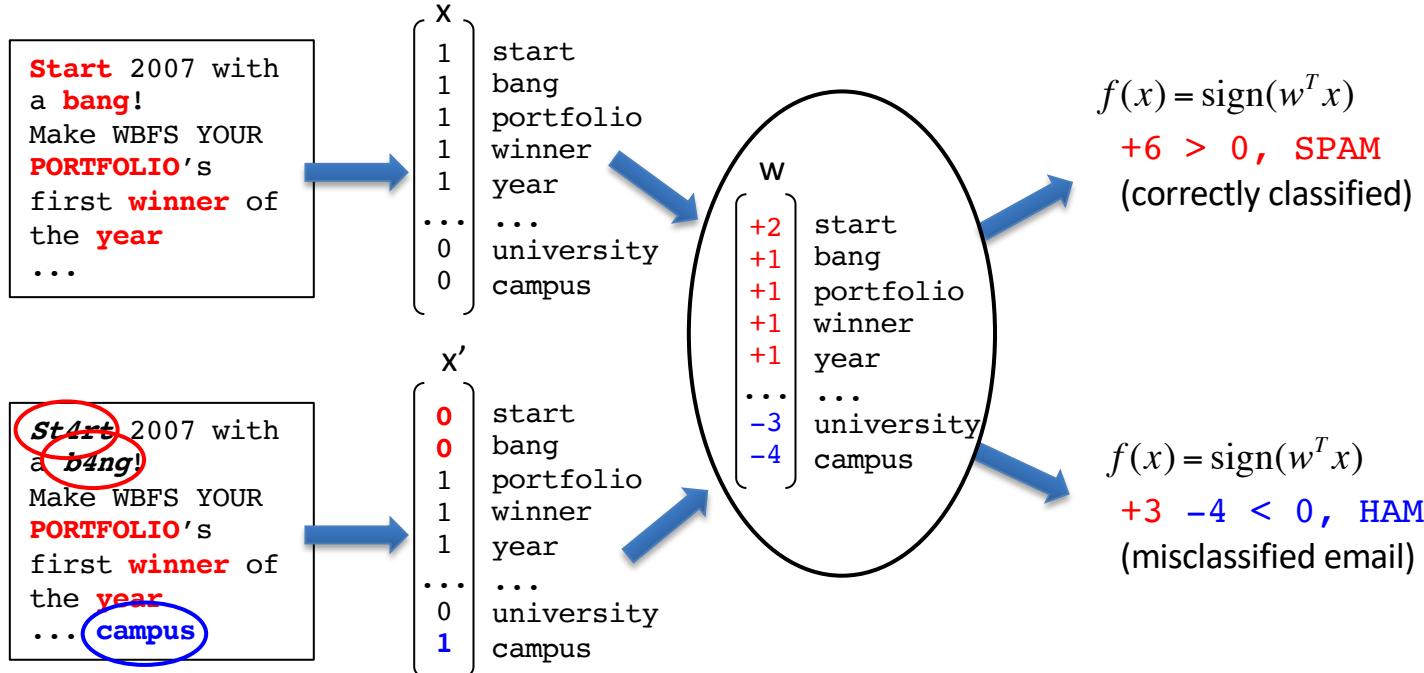
- Given a model of the adversary characterized by her:
 - **Goal**
 - **Knowledge**
 - **Capability**

Try to anticipate the adversary!

- What is the **optimal attack** the attacker can craft?
- What is the expected performance decrease of your classifier?

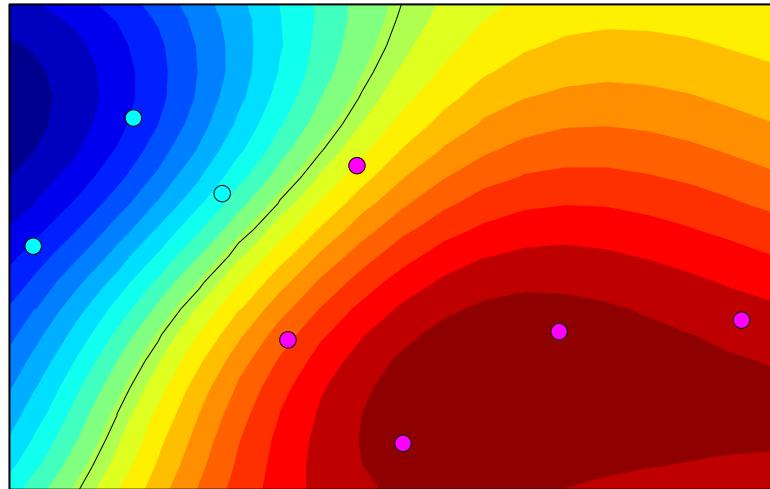
Evasion of Linear Classifiers

- **Problem:** how to evade a linear (trained) classifier?



Evasion of Nonlinear Classifiers

- **What if the classifier is nonlinear?**
- Decision functions can be arbitrarily complicated, with no clear relationship between features (\mathbf{x}) and classifier parameters (\mathbf{w})



Detection of Malicious PDF Files

Srndic & Laskov, Detection of malicious PDF files based on hierarchical document structure, NDSS 2013

"The most aggressive evasion strategy we could conceive was successful for only 0.025% of malicious examples tested against a nonlinear SVM classifier with the RBF kernel [...].

Currently, we do not have a rigorous mathematical explanation for such a surprising robustness. Our intuition suggests that [...] **the space of true features is “hidden behind” a complex nonlinear transformation which is mathematically hard to invert.**

[...] the same attack staged against the linear classifier [...] had a 50% success rate; hence, **the robustness of the RBF classifier must be rooted in its nonlinear transformation”**

Evasion Attacks against Machine Learning at Test Time

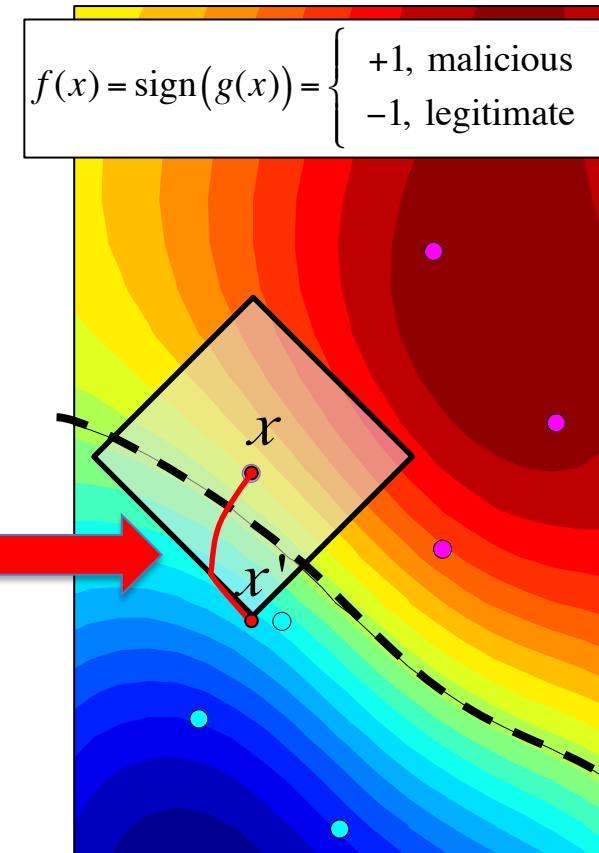
Biggio, Corona, Maiorca, Nelson, Srndic, Laskov, Giacinto, Roli, ECML-PKDD 2013

- **Goal:** maximum-confidence evasion
- **Knowledge:** perfect (white-box attack)
- **Attack strategy:**

$$\min_{x'} g(x')$$

$$\text{s. t. } \|x - x'\|_p \leq d_{\max}$$

- Non-linear, constrained optimization
 - **Projected gradient descent:** approximate solution for smooth functions
- Gradients of $g(x)$ can be analytically computed in many cases
 - SVMs, Neural networks



Computing Descent Directions

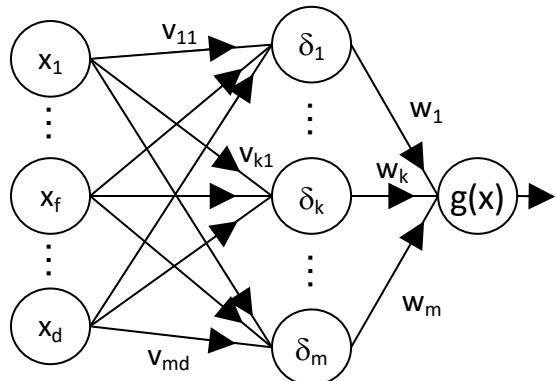
Support vector machines

$$g(x) = \sum_i \alpha_i y_i k(x, x_i) + b, \quad \nabla g(x) = \sum_i \alpha_i y_i \nabla k(x, x_i)$$

RBF kernel gradient:

$$\nabla k(x, x_i) = -2\gamma \exp\left\{-\gamma \|x - x_i\|^2\right\}(x - x_i)$$

Neural networks

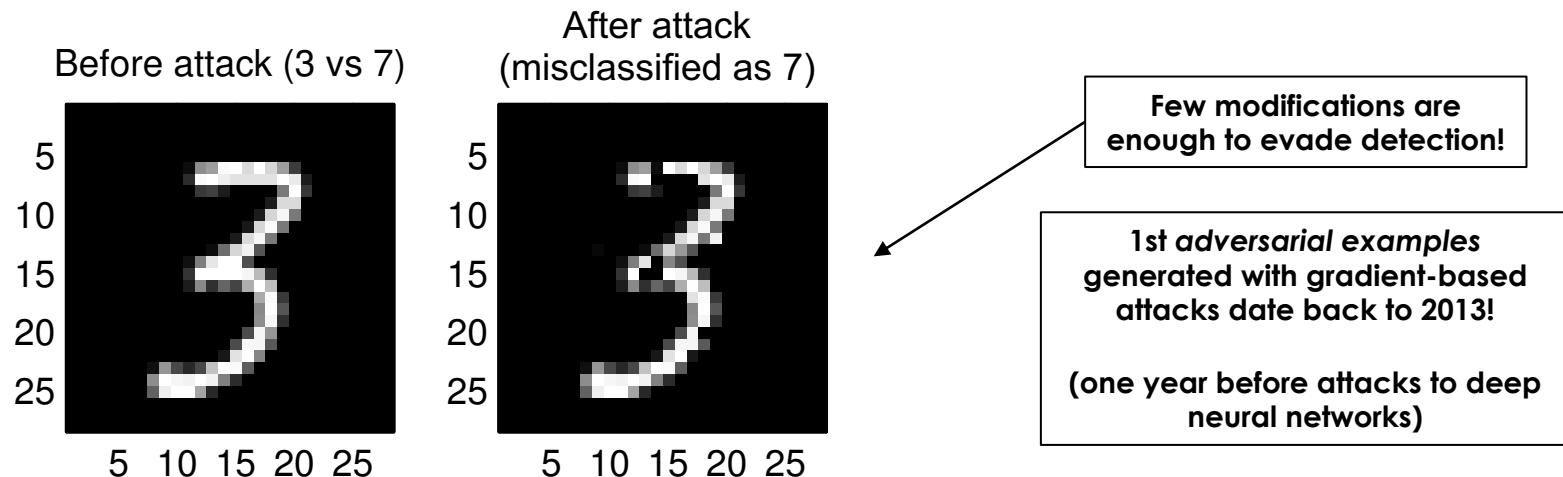


$$g(x) = \left[1 + \exp\left(-\sum_{k=1}^m w_k \delta_k(x)\right) \right]^{-1}$$

$$\frac{\partial g(x)}{\partial x_f} = g(x)(1-g(x)) \sum_{k=1}^m w_k \delta_k(x)(1-\delta_k(x)) v_{kf}$$

An Example on Handwritten Digits

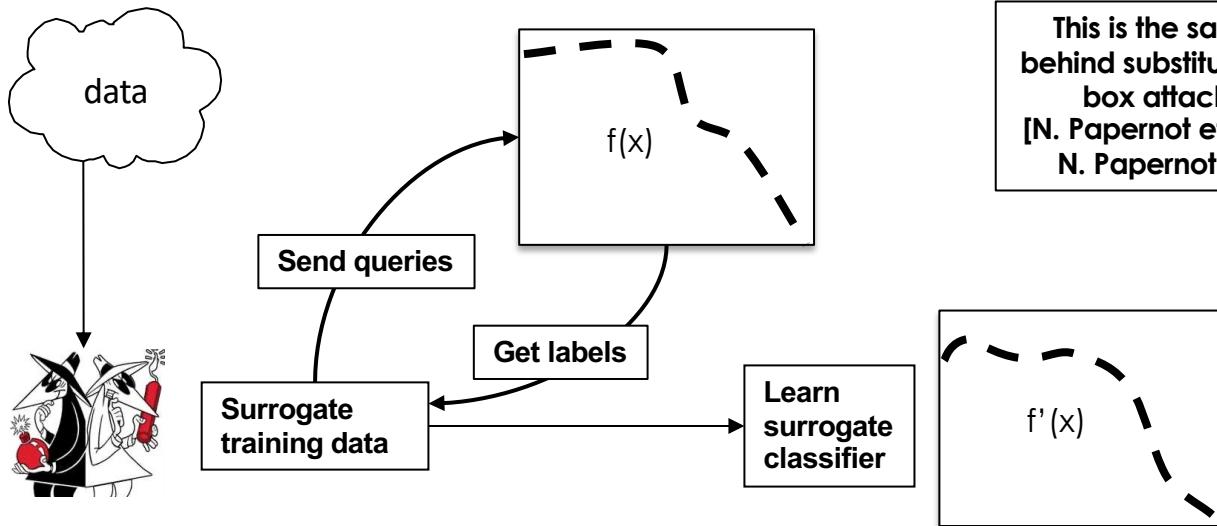
- Nonlinear SVM (RBF kernel) to discriminate between '3' and '7'
- **Features:** gray-level pixel values (28 x 28 image = 784 features)



Bounding the Adversary's Knowledge

Limited-knowledge (gray/black-box) attacks

- Only feature representation and (possibly) learning algorithm are known
- Surrogate data sampled from the same distribution as the classifier's training data
- Classifier's feedback to label surrogate data

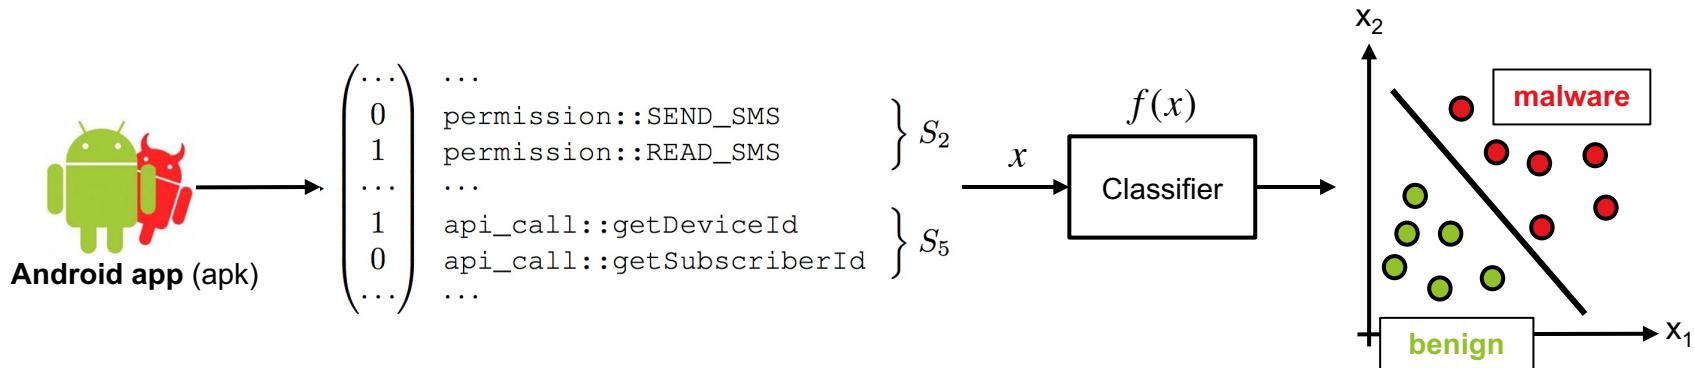


This is the same underlying idea behind substitute models and black-box attacks (*transferability*)
[N. Papernot et al., IEEE Euro S&P '16;
N. Papernot et al., ASIACCS'17]

Recent Results on Android Malware Detection

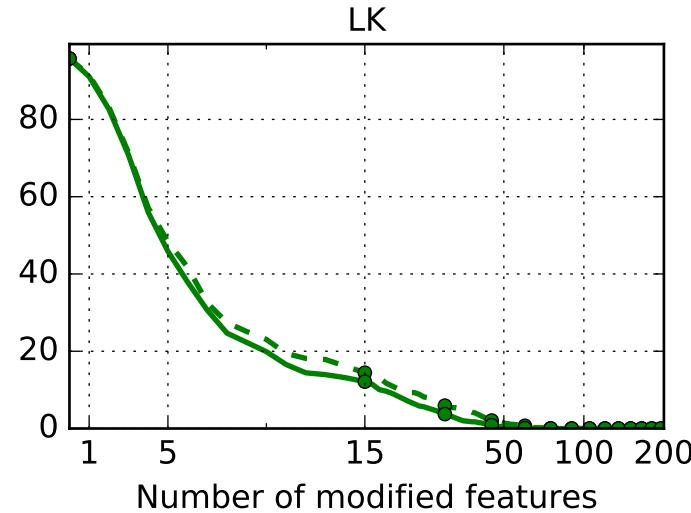
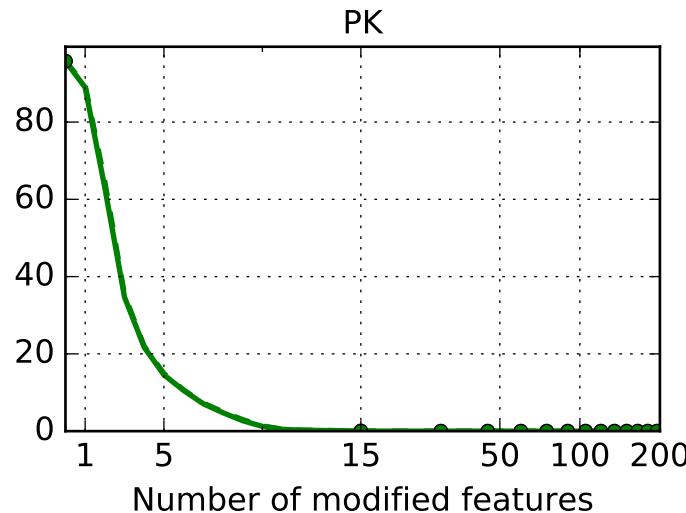
- **Drebin:** Arp et al., NDSS 2014
 - Android malware detection directly on the mobile phone
 - Linear SVM trained on features extracted from static code analysis

Feature sets	
manifest	S_1 Hardware components
	S_2 Requested permissions
	S_3 Application components
	S_4 Filtered intents
dexcode	S_5 Restricted API calls
	S_6 Used permission
	S_7 Suspicious API calls
	S_8 Network addresses



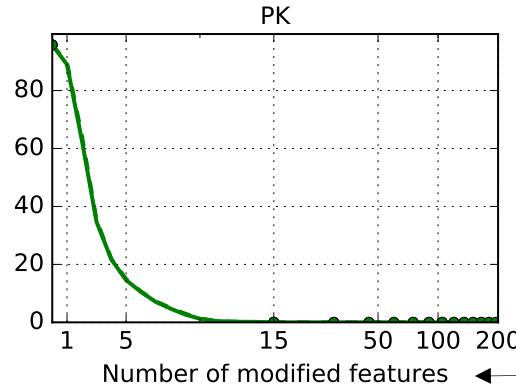
Recent Results on Android Malware Detection

- **Dataset (Drebin):** 5,600 malware and 121,000 benign apps (TR: 30K, TS: 60K)
- **Detection rate** at FP=1% vs max. number of manipulated features (averaged on 10 runs)
 - Perfect knowledge (PK) white-box attack; Limited knowledge (LK) black-box attack



Take-home Messages

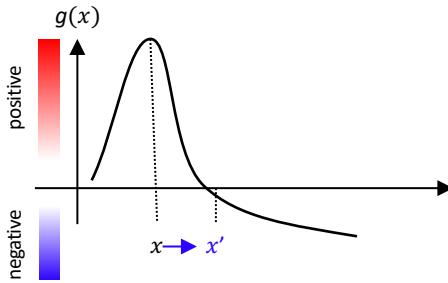
- Linear and non-linear supervised classifiers are vulnerable to well-crafted evasion attacks
- Performance evaluation should be always performed as a function of the adversary's knowledge and capability
 - **Security Evaluation Curves**



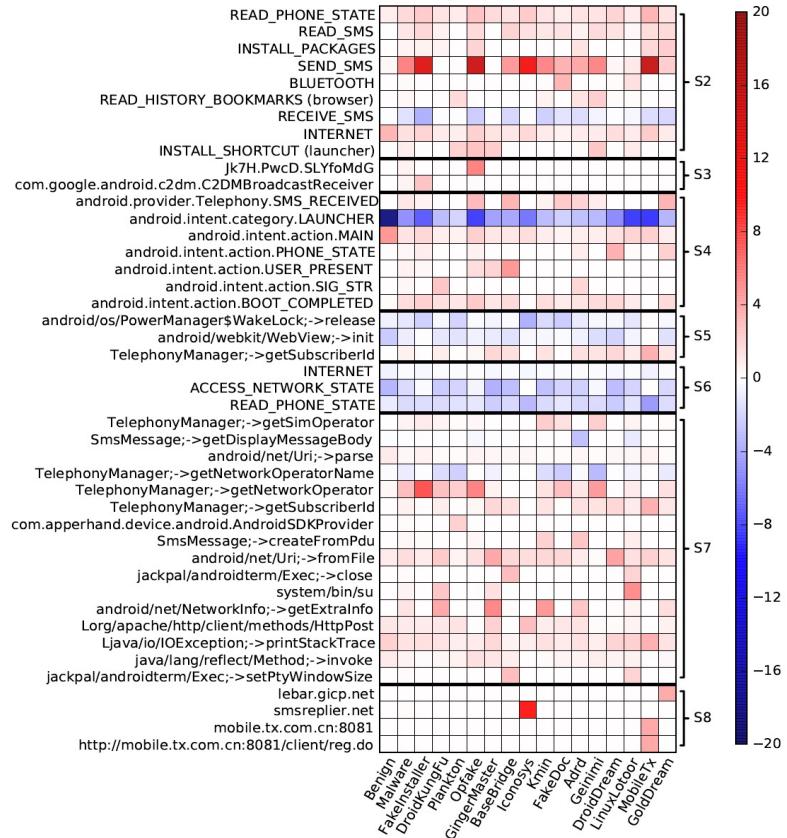
$$\begin{aligned} & \min_{x'} g(x') \\ \text{s.t. } & d(x, x') \leq d_{\max} \\ & x \leq x' \end{aligned}$$

Why Is Machine Learning So Vulnerable?

- Learning algorithms tend to overemphasize some features to discriminate among classes
- Large sensitivity to changes of such input features: $\nabla_x g(x)$



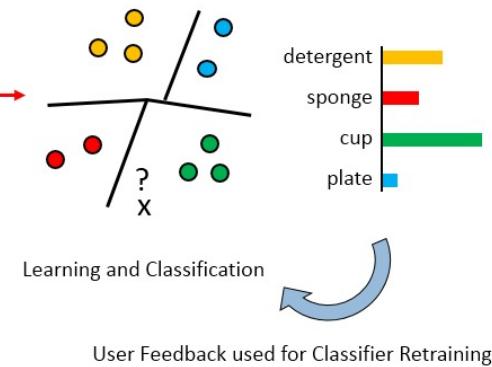
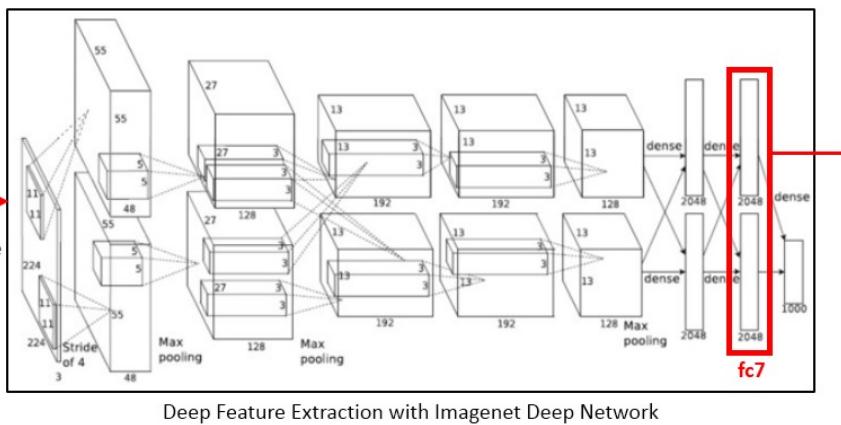
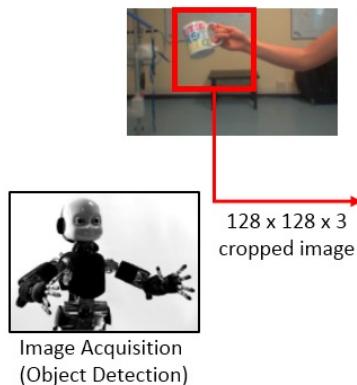
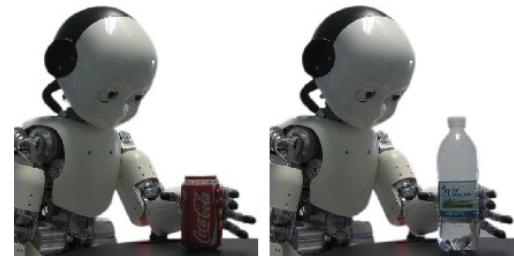
- Different classifiers tend to find the same set of **relevant features**
 - that is why attacks can *transfer* across models!



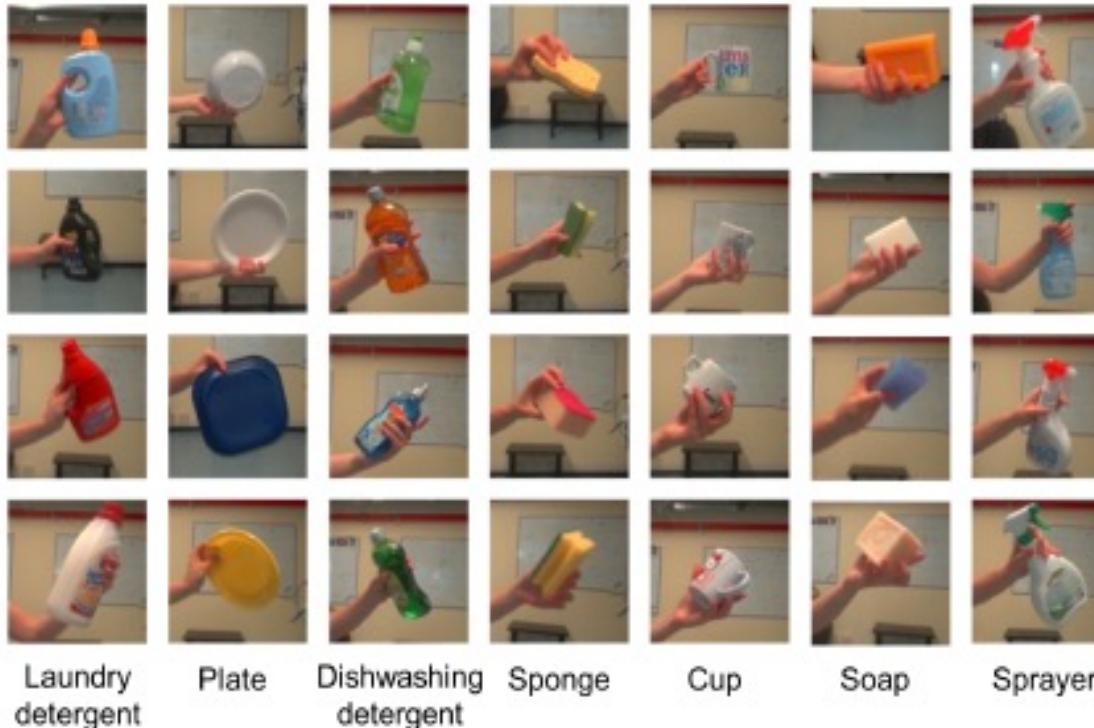
Evasion of Multiclass Classifiers

Is Deep Learning Safe for Robot Vision?

- Evasion attacks against the iCub humanoid robot
 - Deep Neural Network used for visual object recognition



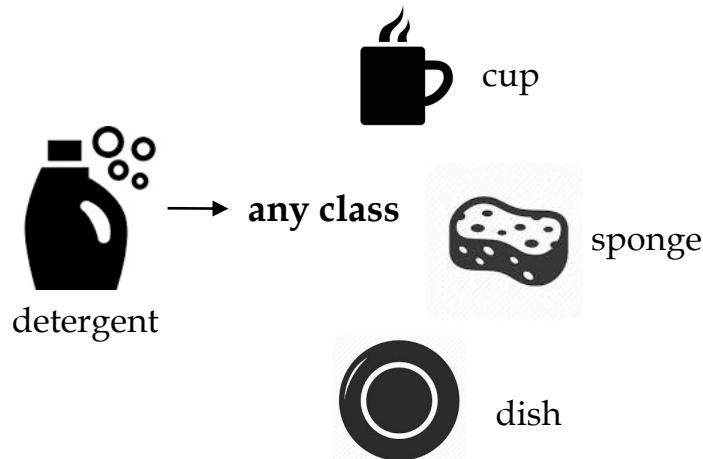
iCubWorld28 Data Set: Example Images



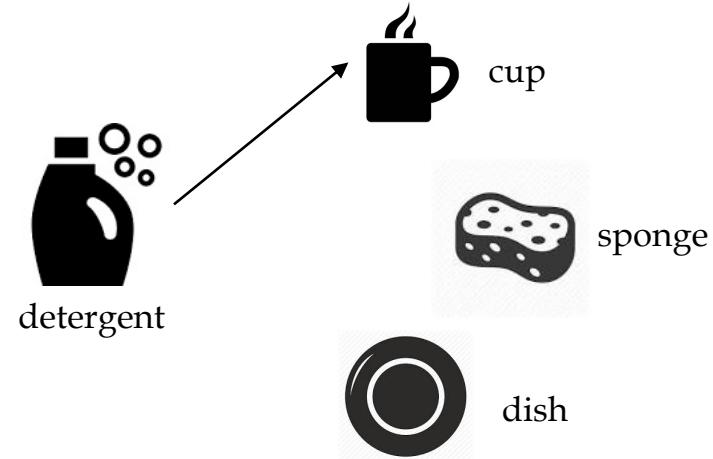
From Binary to Multiclass Evasion

- In multiclass problems, classification errors occur in different classes.
- Thus, the attacker may aim:
 1. to have a sample misclassified as any class different from the true class (**error-generic attacks**)
 2. to have a sample misclassified as a specific class (**error-specific attacks**)

Error-generic attacks



Error-specific attacks

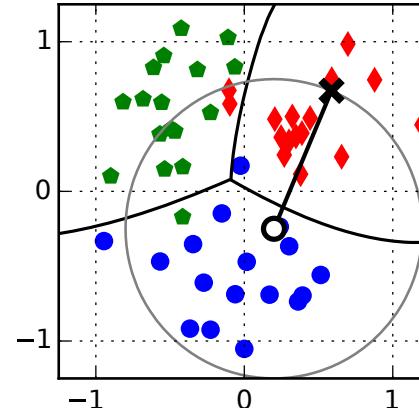


Error-generic (Indiscriminate) Evasion

- **Error-generic evasion**
 - k is the true class (**blue**)
 - l is the competing (closest) class in feature space (**red**)
- The attack minimizes the objective to have the sample misclassified as the *closest* class (could be any!)

$$\begin{aligned} \min_{\mathbf{x}'} \quad & \Omega(\mathbf{x}') , \\ \text{s.t.} \quad & d(\mathbf{x}, \mathbf{x}') \leq d_{\max} , \\ & \mathbf{x}_{lb} \preceq \mathbf{x}' \preceq \mathbf{x}_{ub} , \end{aligned}$$

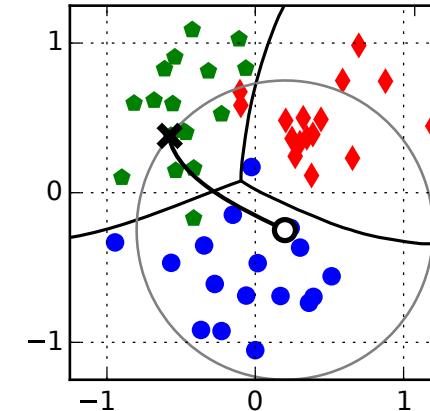
$$\Omega(\mathbf{x}) = f_k(\mathbf{x}) - \max_{l \neq k} f_l(\mathbf{x})$$



Error-specific (Targeted) Evasion

- **Error-specific evasion**
 - k is the target class (**green**)
 - l is the competing class (initially, the **blue** class)
- The attack maximizes the objective to have the sample misclassified as the *target* class

$$\begin{aligned} \max_{\mathbf{x}'} \quad & \Omega(\mathbf{x}') , \\ \text{s.t.} \quad & d(\mathbf{x}, \mathbf{x}') \leq d_{\max} , \\ & \mathbf{x}_{lb} \preceq \mathbf{x}' \preceq \mathbf{x}_{ub} , \end{aligned}$$



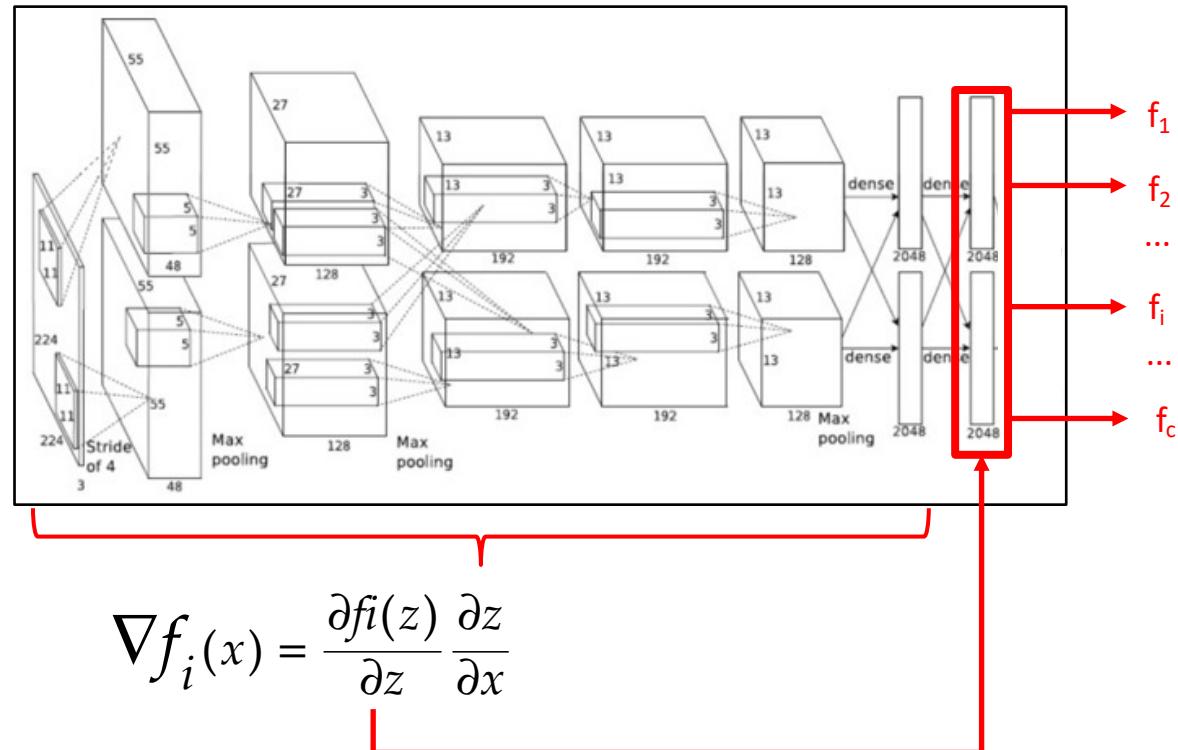
Adversarial Examples against iCub – Gradient Computation

The given optimization problems can be both solved with gradient-based algorithms

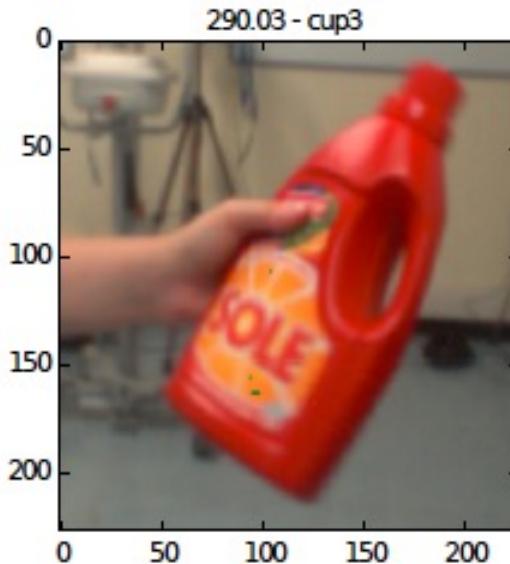
The gradient of the objective can be computed using the **chain rule**

1. the gradient of the functions $f_i(z)$ can be computed if the chosen classifier is differentiable

2. ... and then backpropagated through the deep network with *automatic differentiation*

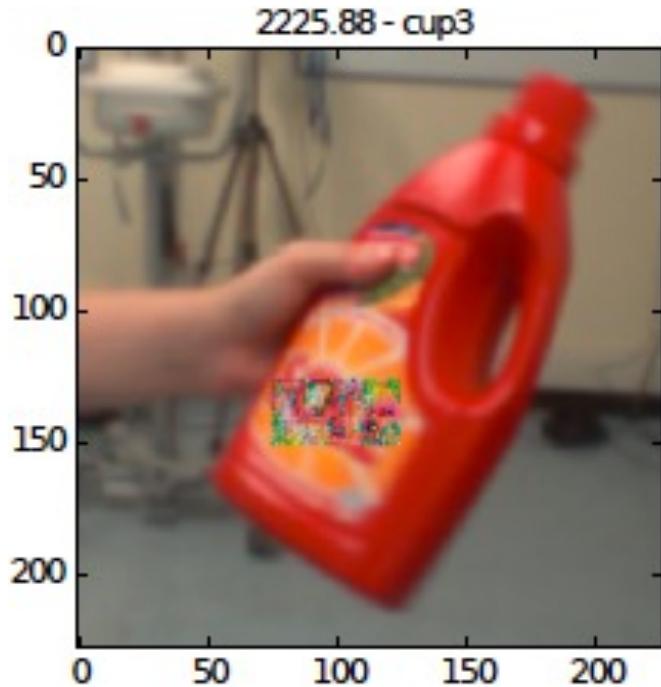


Example of Adversarial Images against iCub



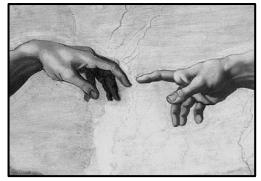
An adversarial example from class *laundry-detergent*, modified by the proposed algorithm to be misclassified as *cup*

The “Sticker” Attack against iCub



Adversarial example generated by manipulating only a specific region, to simulate a sticker that could be applied to the real-world object.

This image is classified as *cup*.



2014: Deep Learning Meets Adversarial Machine Learning

The Discovery of Adversarial Examples

Intriguing properties of neural networks

Christian Szegedy

Google Inc.

Wojciech Zaremba

New York University

Ilya Sutskever

Google Inc.

Joan Bruna

New York University

Dumitru Erhan

Google Inc.

Ian Goodfellow

University of Montreal

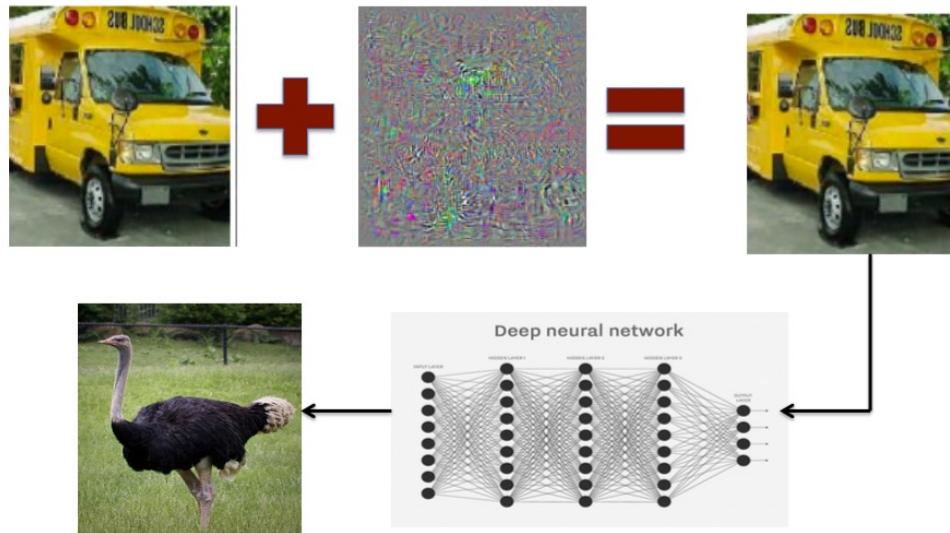
Rob Fergus

New York University
Facebook Inc.

... we find that deep neural networks learn **input-output mappings** that are fairly **discontinuous** to a significant extent. We can cause the network to misclassify an image by applying a certain **hardly perceptible perturbation**, which is found by maximizing the network's prediction error ...

Adversarial Examples and Deep Learning

- C. Szegedy et al. (ICLR 2014) independently developed a gradient-based attack against deep neural networks
 - minimally-perturbed adversarial examples



Creation of Adversarial Examples

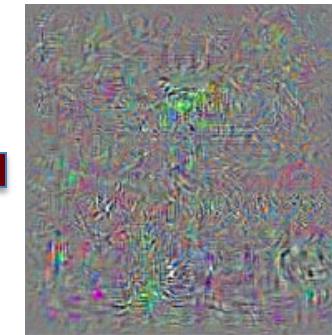
- Minimize $\|r\|_2$ subject to:
 1. $f(x + r) = l \quad f(x) \neq l$
 2. $x + r \in [0, 1]^m$

The adversarial image $x + r$ is visually hard to distinguish from x
Informally speaking, the solution $x + r$ is the closest image to x classified as l by f

The solution is approximated using using a box-constrained limited-memory BFGS



School Bus (x)



Adversarial Noise (r)



Ostrich
Struthio Camelus

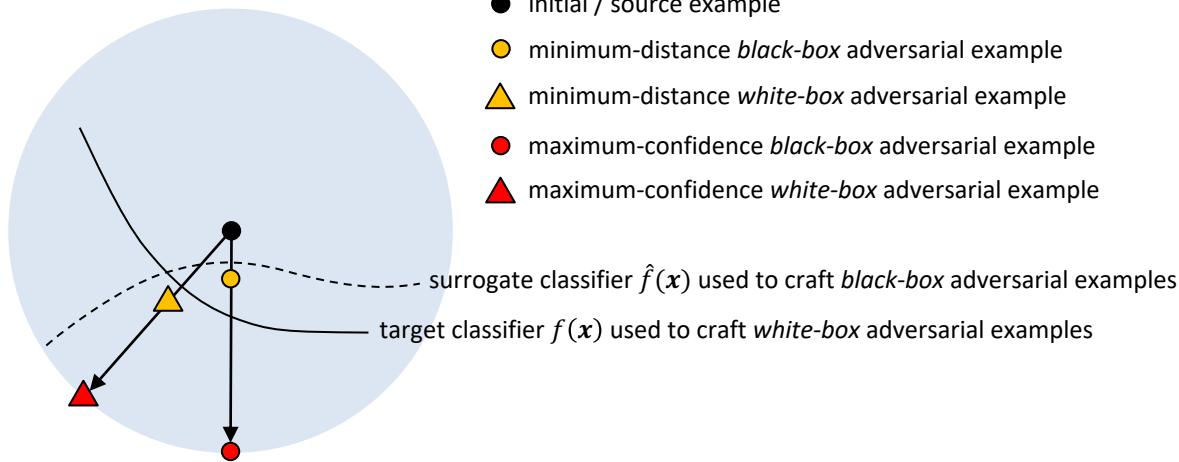
Minimum-norm vs Maximum-confidence Attacks

- Szegedy et al., ICLR 2014 aim to measure the minimum distance to evasion
 - Better suited to the analysis of adversarial robustness in the white-box case
- Biggio et al., ECML 2013 maximizes misclassification confidence within a given budget
 - The intuition was to craft attacks that are more difficult to detect, and to evade classifiers with higher probability also when knowledge of the boundary is not perfect (*transfer attacks*)

Adversary's goal. As suggested by Laskov and Kloft [17], the adversary's goal should be defined in terms of a utility (loss) function that the adversary seeks to maximize (minimize). In the evasion setting, the attacker's goal is to manipulate a single (without loss of generality, positive) sample that should be misclassified. Strictly speaking, it would suffice to find a sample \mathbf{x} such that $g(\mathbf{x}) < -\epsilon$ for any $\epsilon > 0$; i.e., the attack sample only just crosses the decision boundary.² Such attacks, however, are easily thwarted by slightly adjusting the decision threshold. A better strategy for an attacker would thus be to create a sample that is misclassified with high confidence; i.e., a sample minimizing the value of the classifier's discriminant function, $g(\mathbf{x})$, subject to some feasibility constraints.

[Biggio, Roli et al., ECML PKDD 2013]

Minimum-norm vs Maximum-confidence Attacks



Many Black Swans After 2014...

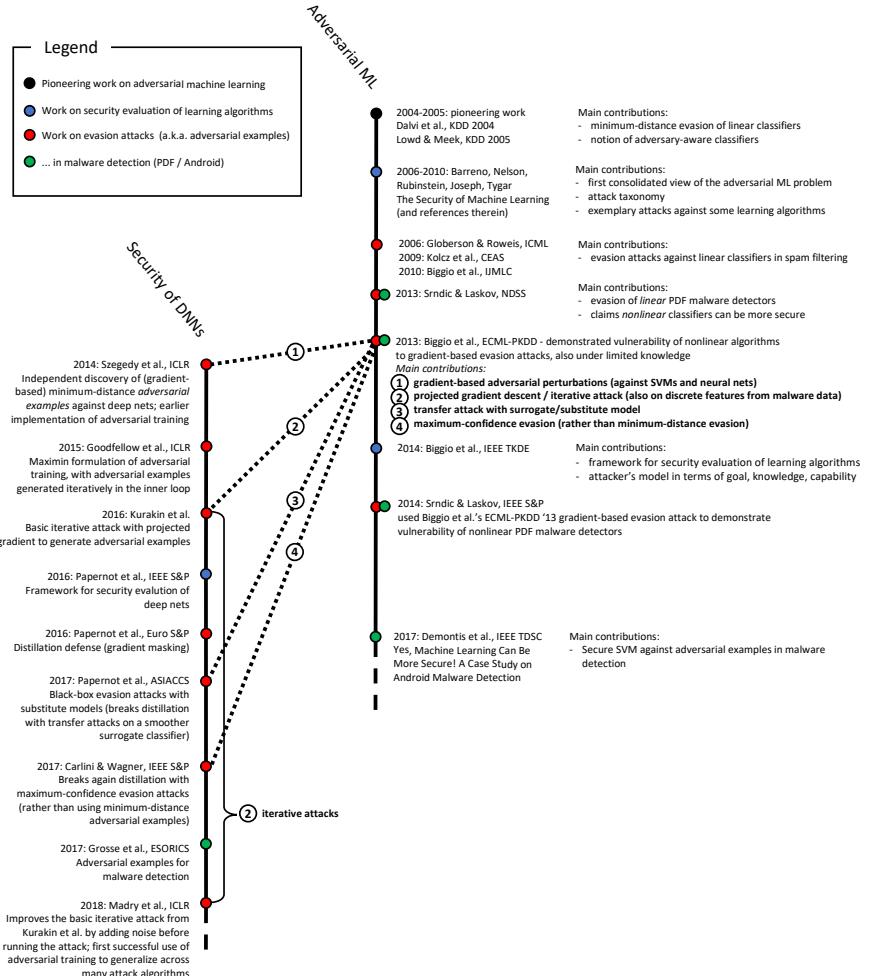
[Search <https://arxiv.org> with keywords “adversarial examples”]



- Several defenses have been proposed against adversarial examples, and more powerful attacks have been developed to show that they are ineffective. *Remember the arms race?*
- Most of these attacks are modifications to the optimization problems reported for evasion attacks / adversarial examples, using different gradient-based solution algorithms, initializations and stopping conditions.
- Most popular attack algorithms: FGSM (Goodfellow et al.), JSMA (Papernot et al.), CW (Carlini & Wagner, and follow-up versions)

Timeline of Learning Security

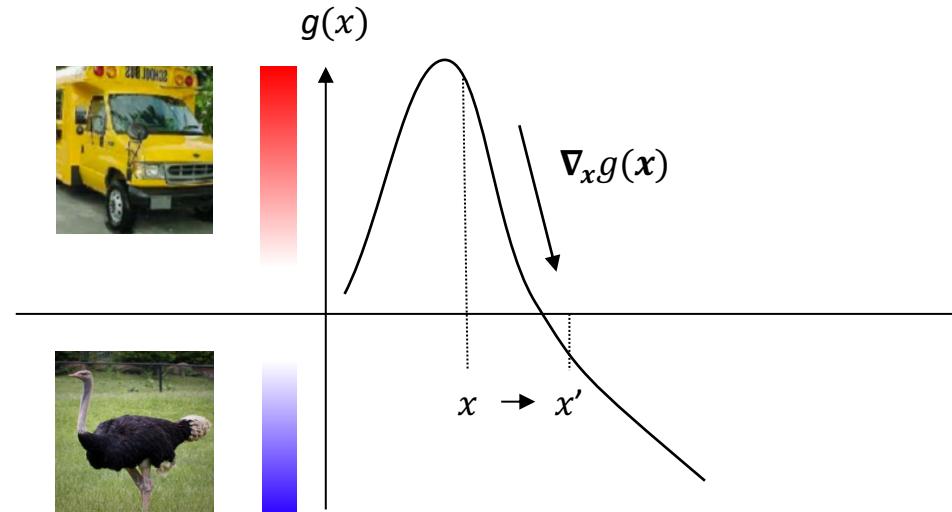
Biggio and Roli, Wild Patterns: Ten Years After The Rise of Adversarial Machine Learning, Pattern Recognition, 2018



Why Adversarial Perturbations are Imperceptible?

Why Adversarial Perturbations against Deep Networks are Imperceptible?

- Large sensitivity of $g(\mathbf{x})$ to input changes
 - i.e., the **input gradient** $\nabla_{\mathbf{x}}g(\mathbf{x})$ has a large norm (scales with input dimensions!)
 - Thus, even small modifications along that direction will cause large changes in the predictions



Adversarial Examples and Security Evaluation (Demo Session)

secml: An Open-source Python Library for ML Security

ml

- ML algorithms via sklearn
- DL algorithms and optimizers via PyTorch and Tensorflow



adv

- attacks (evasion, poisoning, ...) with custom/faster solvers
- defenses (advx rejection, adversarial training, ...)

expl

- Explanation methods based on influential features
- Explanation methods based on influential prototypes

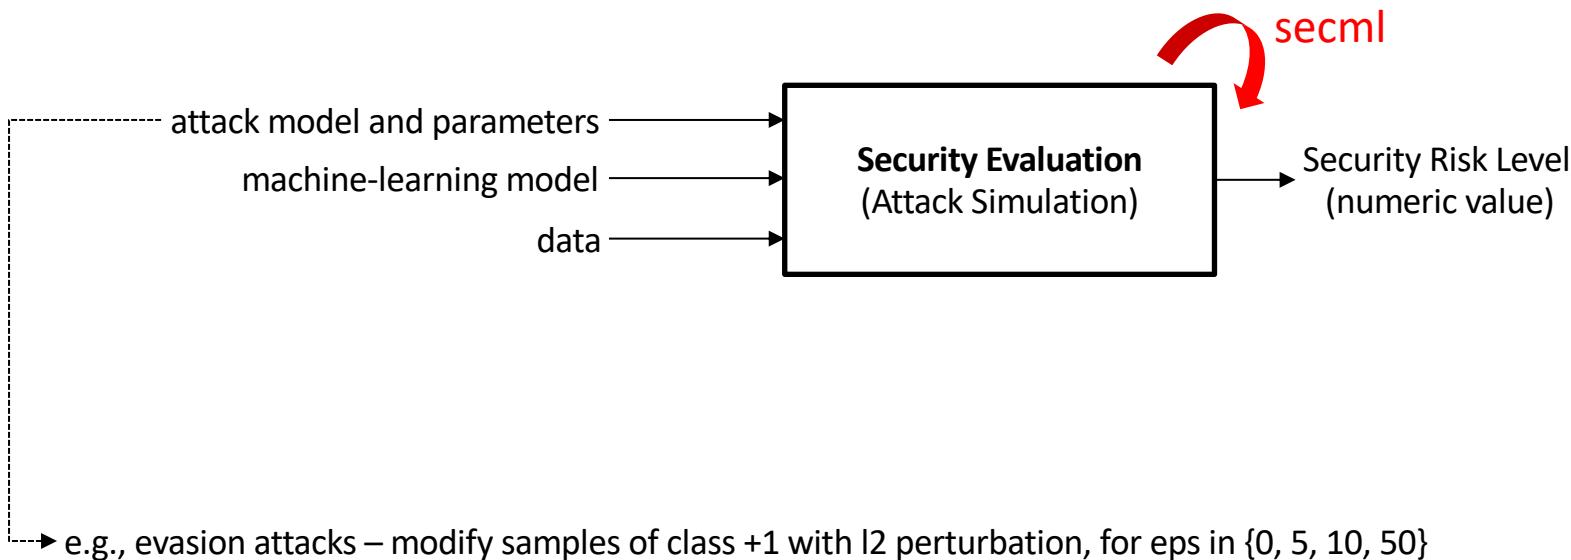


others

- Parallel computation
- Support for dense/sparse data
- Advanced plotting functions (via matplotlib)
- Modular and easy to extend

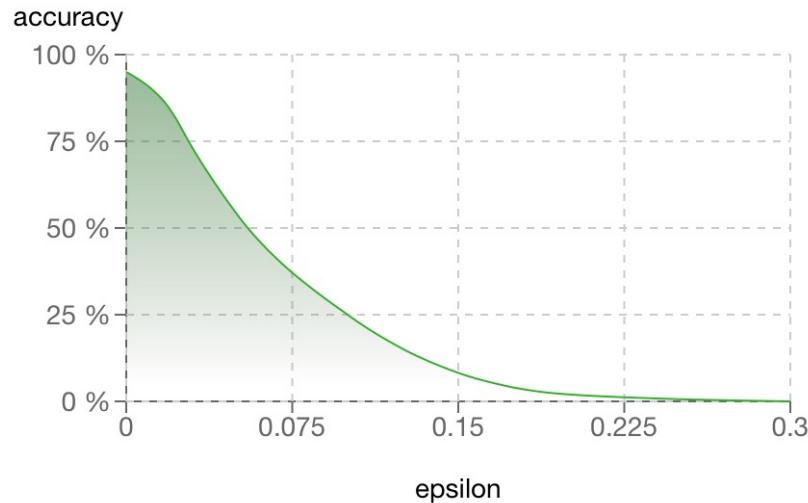
Code: <https://github.com/pralab/secml>

ML Security Evaluation

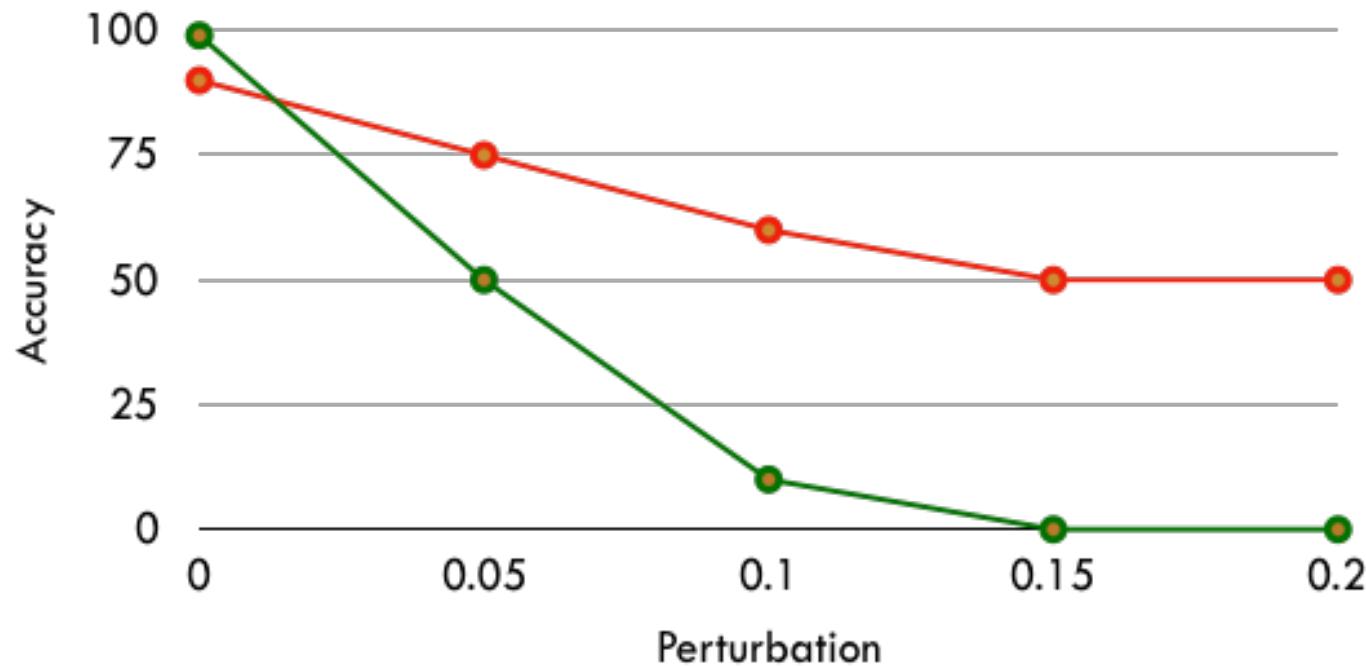


Security Evaluation Curves

- **Security evaluation curves**
 - accuracy vs increasing perturbation
- **Security value:**
 - mean accuracy under attack
- **Security level:**
 - Low / Med / High



Security Evaluation Curves



Interactive Demo

- Demo available at: <https://www.pluribus-one.it/research/sec-ml/demo>

Secure ML Demo - Deep Learning security

Free demo for the security evaluation of Deep Learning algorithms

Secure ML Research

Tutorial: Wild Patterns

Secure ML Library

Web Demo

Other Attacks on Machine Learning

Attacks against Machine Learning

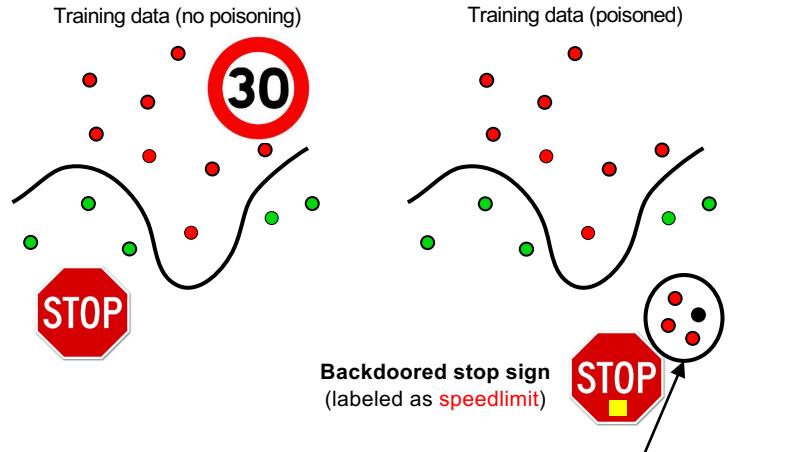
Attacker's Goal			
Attacker's Capability	Integrity	Availability	Privacy / Confidentiality
Test data	Evasion (a.k.a. adversarial examples)	Sponge Attacks	<i>Model extraction / stealing Model inversion (hill climbing) Membership inference</i>
Training data	<i>Backdoor poisoning (to allow subsequent intrusions) – e.g., backdoors or neural trojans</i>	<i>DoS poisoning (to maximize classification error)</i>	-

Attacker's Knowledge:

- perfect-knowledge (PK) white-box attacks
- limited-knowledge (LK) black-box attacks (*transferability* with surrogate/substitute learning models)

Backdoor Attacks

Poisoning Integrity Attacks



Attack referred to as backdoor

T. Gu, B. Dolan-Gavitt, and S. Garg. Badnets: Identifying vulnerabilities in the machine learning model supply chain. In NIPS Workshop on Machine Learning and Computer Security, 2017.

X. Chen, C. Liu, B. Li, K. Lu, and D. Song. Targeted backdoor attacks on deep learning systems using data poisoning. ArXiv e-prints, 2017.

Attack referred to as ‘poisoning integrity’

M. Barreno, B. Nelson, R. Sears, A. D. Joseph, and J. D. Tygar. Can machine learning be secure? In Proc. ACM Symp. Information, Computer and Comm. Sec., ASIACCS ’06, pages 16–25, New York, NY, USA, 2006. ACM.

M. Barreno, B. Nelson, A. Joseph, and J. Tygar. The security of machine learning. Machine Learning, 81:121–148, 2010.

B. Biggio, B. Nelson, and P. Laskov. Poisoning attacks against support vector machines. In J. Langford and J. Pineau, editors, 29th Int'l Conf. on Machine Learning, pages 1807–1814. Omnipress, 2012.

B. Biggio, G. Fumera, and F. Roli. Security evaluation of pattern classifiers under attack. IEEE Transactions on Knowledge and Data Engineering, 26(4):984–996, April 2014.

H. Xiao, B. Biggio, G. Brown, G. Fumera, C. Eckert, and F. Roli. Is feature selection secure against training data poisoning? In F. Bach and D. Blei, editors, JMLR W&CP - Proc. 32nd Int'l Conf. Mach. Learning (ICML), volume 37, pages 1689–1698, 2015.

L. Munoz-Gonzalez, B. Biggio, A. Demontis, A. Paudice, V. Wongrassamee, E. C. Lupu, and F. Roli. Towards poisoning of deep learning algorithms with back-gradient optimization. In 10th ACM Workshop on Artificial Intelligence and Security, AISeC ’17, pp. 27–38, 2017. ACM.

B. Biggio and F. Roli. Wild patterns: Ten years after the rise of adversarial machine learning. ArXiv e-prints, 2018.

M. Jagielski, A. Oprea, B. Biggio, C. Liu, C. Nita-Rotaru, and B. Li. Manipulating machine learning: Poisoning attacks and countermeasures for regression learning. In 39th IEEE Symp. on Security and Privacy, 2018.

Model Inversion Attacks

Privacy Attacks

- **Goal:** to extract users' sensitive information (e.g., face templates stored during user enrollment)
 - Fredrikson, Jha, Ristenpart. *Model inversion attacks that exploit confidence information and basic countermeasures.* ACM CCS, 2015
- Also known as hill-climbing attacks in the biometric community
 - Adler. *Vulnerabilities in biometric encryption systems.* 5th Int'l Conf. AVBPA, 2005
 - Galbally, McCool, Fierrez, Marcel, Ortega-Garcia. *On the vulnerability of face verification systems to hill-climbing attacks.* Patt. Rec., 2010
- **How:** by repeatedly querying the target system and adjusting the input sample to maximize its output score (e.g., a measure of the similarity of the input sample with the user templates)

Training Image



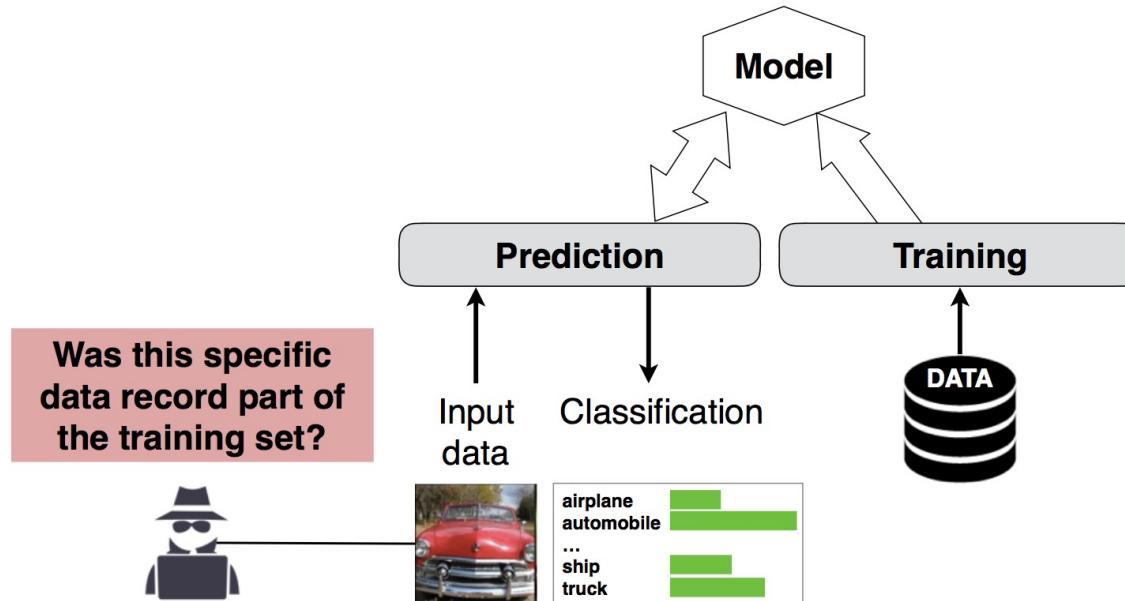
Reconstructed Image



Membership Inference Attacks

Privacy Attacks (Shokri et al., IEEE Symp. SP 2017)

- **Goal:** to identify whether an input sample is part of the training set used to learn a deep neural network based on the observed prediction scores for each class



Is AI/ML Security really Relevant from a More Practical/Business Perspective?

Industry Survey on AI Security (Microsoft)

- Microsoft has seen a notable increase in attacks on commercial ML systems
- **Market reports:** *Gartner's Top 10 Strategic Technology Trends for 2020*: "Through 2022, 30% of all AI cyberattacks will leverage training-data poisoning, AI model theft, or adversarial samples to attack AI-powered systems."
- Despite these reasons to secure ML systems, Microsoft's survey spanning 28 businesses found that most industry practitioners have yet to come to terms with adversarial machine learning
- 25/28 businesses don't have the right tools in place to secure their ML systems and need guidance

TABLE I
ORGANIZATION SIZE

Organization size	Count
Large Organizations (> 1000 employees)	18
Small-and-Medium Size Businesses	10

TABLE III
ML STRATEGY

How do you build ML Systems	Count
Using ML Frameworks	16
Using ML as a Service	10
Building ML Systems from scratch	2

TABLE II
ORGANIZATION TYPES

Organization	Count
Cybersecurity	10
Healthcare	5
Government	4
Consulting	2
Banking	2
Social Media Analytics	1
Publishing	1
Agriculture	1
Urban Planning	1
Food Processing	1
Translation	1

TABLE IV
STATE OF ADVERSARIAL ML

Do you secure your ML systems today	Count
Yes	3
No	22

TABLE V
TOP ATTACK

Which attack would affect your org the most?	Distribution
Poisoning (e.g. [21])	10
Model Stealing (e.g. [22])	6
Model Inversion (e.g. [23])	4
Backdoored ML (e.g. [24])	4
Membership Inference (e.g. [25])	3
Adversarial Examples (e.g. [26])	2
Reprogramming ML System (e.g. [27])	0
Adversarial Example in Physical Domain (e.g. [5])	0
Malicious ML provider recovering training data (e.g. [28])	0
Attacking the ML supply chain (e.g. [24])	0
Exploit Software Dependencies (e.g. [29])	0

<https://www.microsoft.com/security/blog/2020/10/22/cyberattacks-against-machine-learning-systems-are-more-common-than-you-think/>

R.S. Kumar et al., Microsoft, Adversarial Machine Learning – Industry Perspectives, AISeC 2020

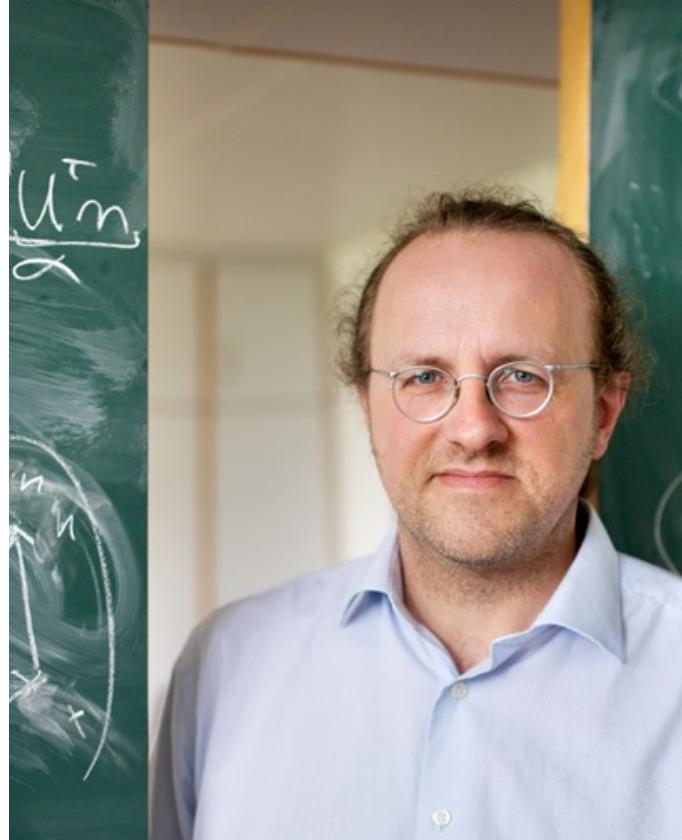
Startups and Standardization Efforts

- <https://adversa.ai>
- <https://www.robustintelligence.com>
- <https://latticeflow.ai>
- <https://resistant.ai>
- <https://troj.ai>
- <https://www.calypsoai.com>
- ETSI working group on AI security
 - <https://www.etsi.org/technologies/securing-artificial-intelligence>

Why Is AI Vulnerable?

Why Is AI Vulnerable?

- **Underlying assumption:** past data is representative of future data (IID data)
- The success of modern AI is on tasks for which we collected enough representative training data
- **We cannot build AI models for each task an agent is ever going to encounter**, but there is a whole world out there where the IID assumption is violated
- **Adversarial attacks** point exactly at this lack of robustness which comes from IID specialization



Bernhard Schölkopf

*Director, Max Planck Institute, Tuebingen,
Germany*

Why Is AI Safety an Important Concern?

- We learn how to break machine learning and AI not just because it is fun, but...
 - to understand the limits of these technologies
 - to be able to design more robust algorithms and systems
- Systems that can be used in safety-critical applications
 - e.g., self-driving cars, monitoring / controlling nuclear plants
- Knowing when to **trust** automated decisions in these contexts is extremely important
 - Should I use the autopilot of my self-driving car or not? Can I trust it?

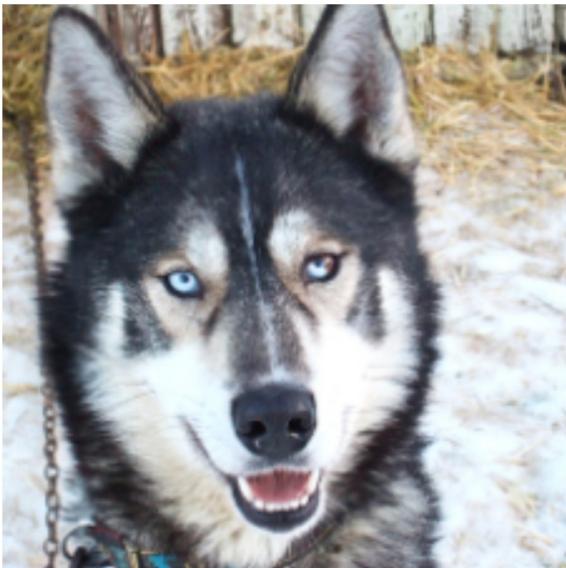
Hacking Tesla Autopilot



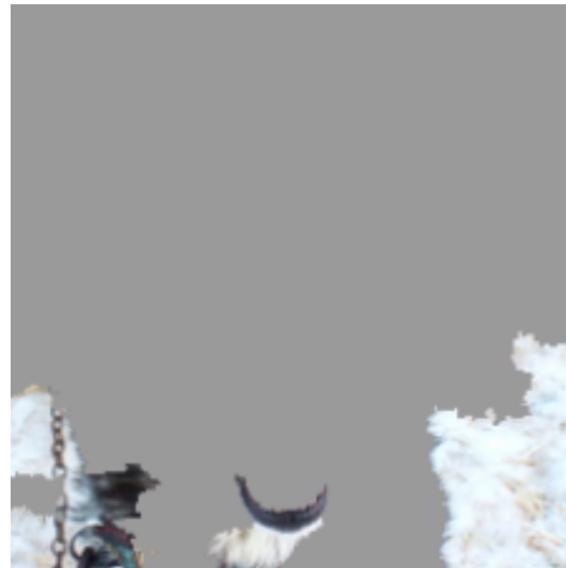
Explainability Is Another Important Asset for AI Safety

- How can we trust a black-box algorithm providing opaque decisions?
 - *Why did my car decide to turn left rather than right?*
 - *Why is this application considered malicious / harmful?*
- The right to explanation (https://en.wikipedia.org/wiki/Right_to_explanation)
 - EU on General Data Privacy Regulation (GDPR), Art. 22
- Important concept
 - to build trust in machines and automated algorithms
 - to understand if the algorithm has properly learned meaningful notions/abstractions from data
 - to uncover potential biases encountered during the learning process

An Example on Image Classification



(a) Husky classified as wolf

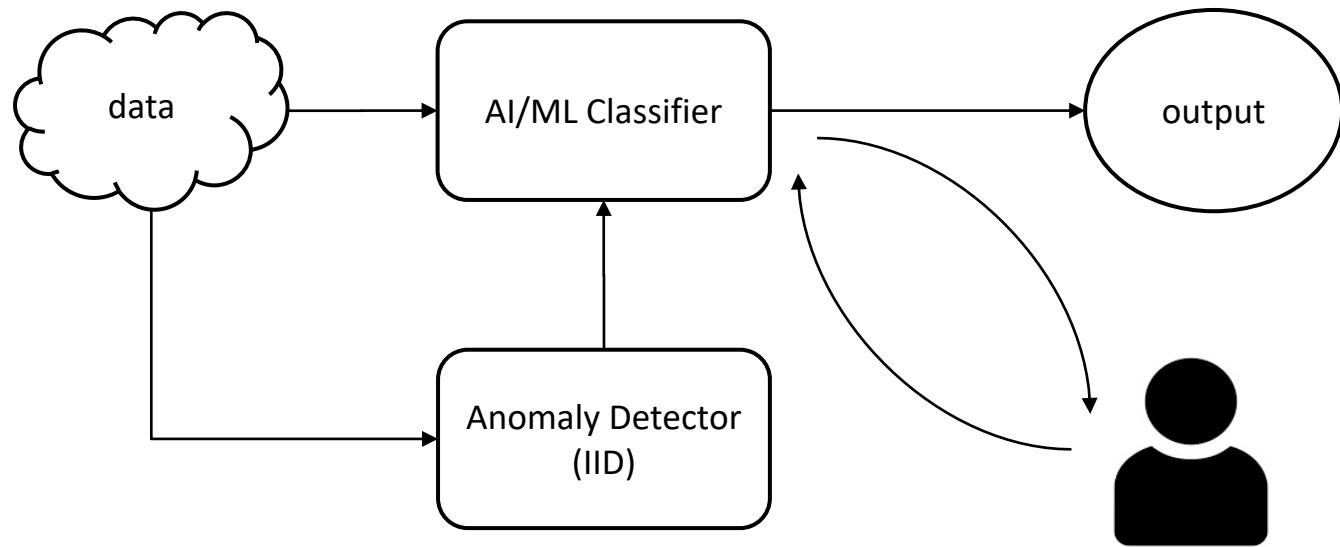


(b) Explanation

The Pillars of Trustworthy AI

- **Safety, Robustness and Reliability**
 - AI systems should perform reliably and safely
- **Transparency, Interpretability and Explainability**
 - AI systems should be understandable
- **Accountability**
 - AI systems should have algorithmic accountability
- **Security and Privacy**
 - AI systems should be secure and respect privacy
- **Fairness**
 - AI systems should treat all people fairly
- **Inclusiveness**
 - AI systems should empower everyone and engage people

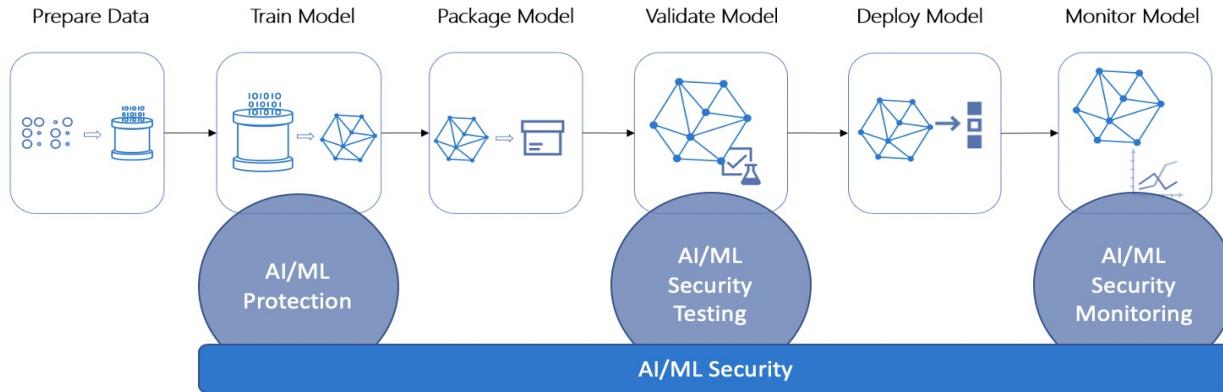
What Can We Do, Next? Explainability/Humans in the Loop



Human in the loop, giving feedback on explainable decisions

What Can We Do, Next? From MLOps to MLSecOps

- MLOps aims at automating development & operations of ML
- ML**Sec**Ops aims at adding ML security to the MLOps development & deployment cycle



To Conclude...

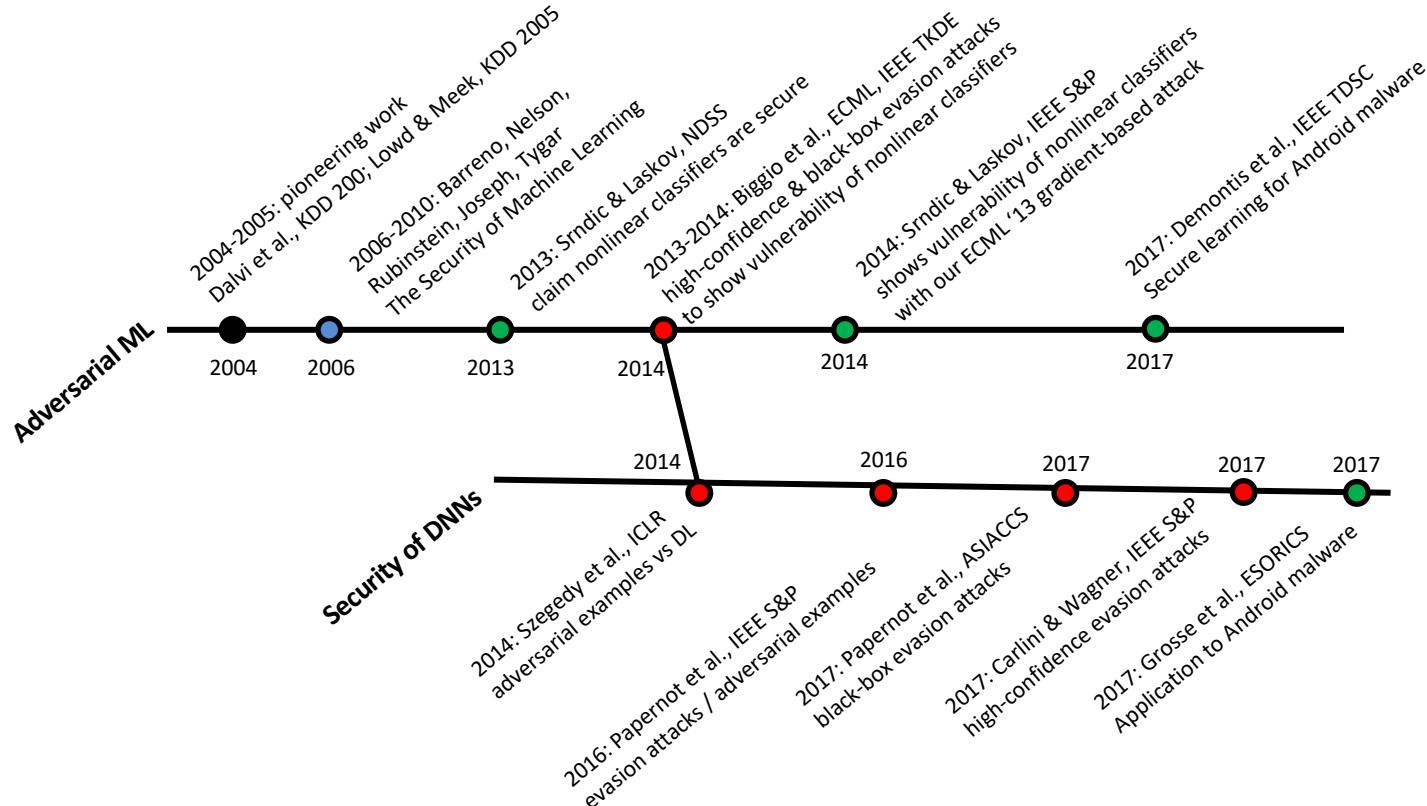
This is a recent research field...

Dagstuhl Perspectives Workshop on
“Machine Learning in Computer Security”
Schloss Dagstuhl, Germany, Sept. 9th-14th, 2012

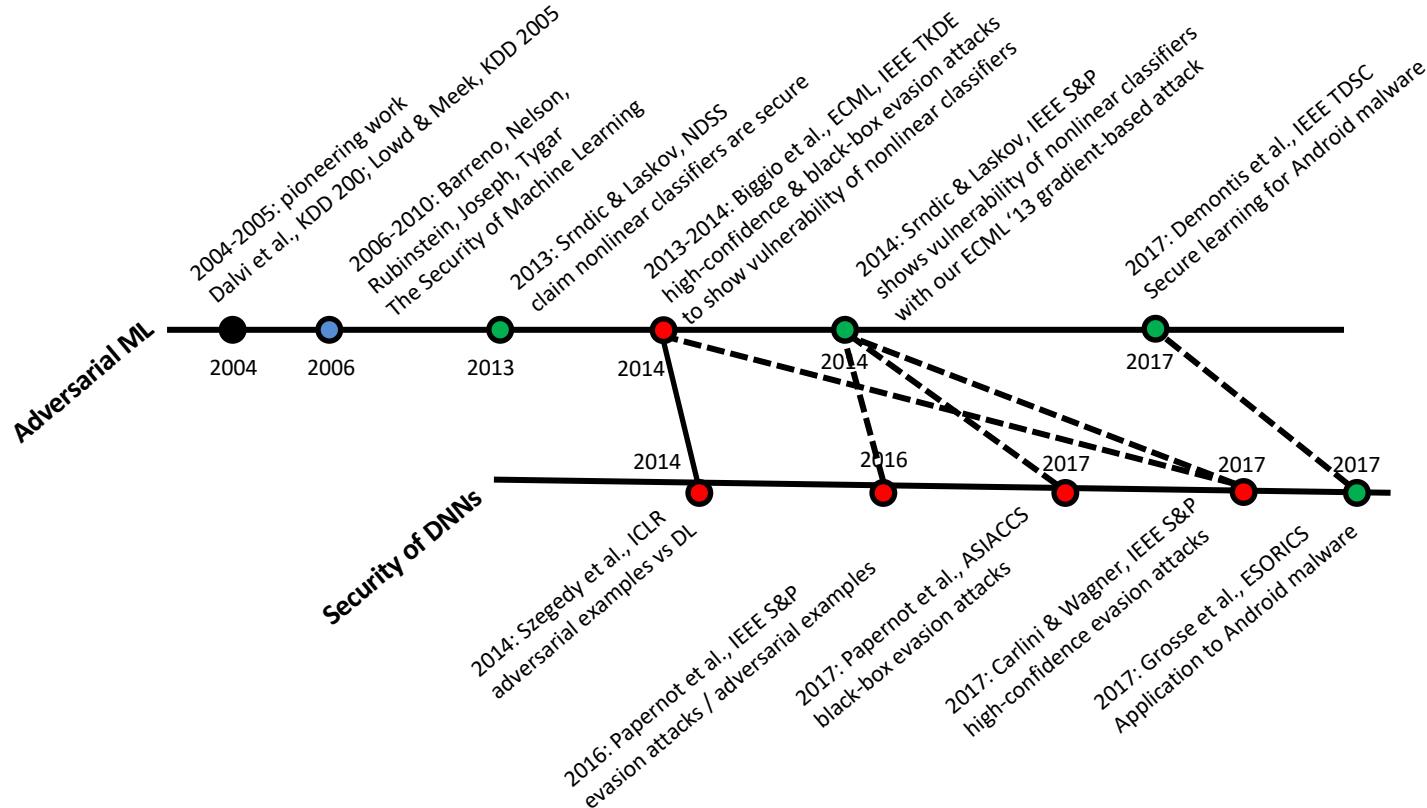


SCHLOSS DAGSTUHL
Leibniz-Zentrum für Informatik

Timeline of Learning Security



Timeline of Learning Security



Black Swans to the Fore

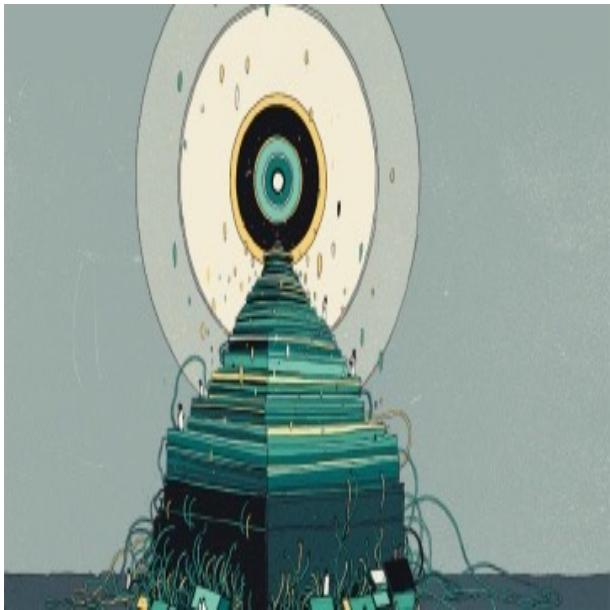
[Szegedy et al., Intriguing properties of neural networks, 2014]



After this “black swan”, the issue of security of DNNs came to the fore...

Not only on scientific specialistic journals...

The Safety Issue to the Fore...



The black box of AI

D. Castelvecchi, Nature, Vol. 538, 20, Oct 2016

Machine learning is becoming ubiquitous in basic research as well as in industry. But for scientists to trust it, they first need to understand what the machines are doing.

Ellie Dobson, director of data science at the big-data firm Arundo Analytics in Oslo: If something were to go wrong as a result of setting the UK interest rates, she says, “the Bank of England can’t say, the black box made me do it”.

Why So Much Interest?

Before the deep net “revolution”, people were not surprised when machine learning was wrong, they were more amazed when it worked well...

Now that it seems to work for real applications, people are disappointed, and worried, for errors that humans do not do...

Errors of Humans and Machines...

Machine learning decisions are affected by several **sources of bias** that causes “strange” errors

But we should keep in mind that also **humans** are **biased...**

The Bat and the Ball Problem

A bat and a ball together cost \$ 1.10

The bat costs \$ 1.0 more than the ball

How much does the ball cost ?

Please, give me the first answer coming to your mind !

The Bat and the Ball Problem

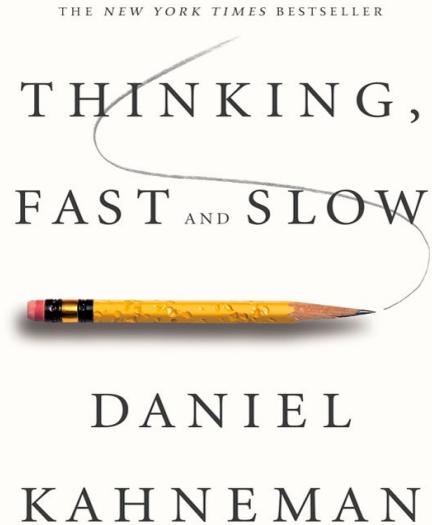
$$\begin{cases} \text{bat+ball}=\$1.10 \\ \text{bat}=\text{ball}+\$1.0 \end{cases}$$

Exact solution is 0.05 dollar (5 cents)

The wrong solution (\$ 0.10) is due to the **attribute substitution**, a psychological process thought to underlie a number of **cognitive biases**

It occurs when an individual has to make a judgment (of a target attribute) that is computationally complex, and instead substitutes a more easily calculated heuristic attribute

Trust in Humans or Machines?



Algorithms are biased, but
also humans are as well...

When should you trust
humans and when
algorithms?

Learning Comes at a Price!



The introduction of novel **learning** functionalities increases the **attack surface** of computer systems and produces new vulnerabilities

Safety of machine learning will be more and more important in future computer systems, as well as **accountability, transparency, and the protection of fundamental human values and rights**



Battista Biggio
battista.biggio@unica.it
 @biggiobattista

Thanks!



If you know the enemy and know yourself, you need not fear the result of a hundred battles
Sun Tzu, The art of war, 500 BC