Poisoning Attacks on Machine Learning

Battista Biggio, Ambra Demontis

Department of Electrical and Electronic Engineering
University of Cagliari, Italy

Attacks against Machine Learning

Attacker's Capability	Attacker's Goal: Misclassifications that do not compromise normal system operation — **Integrity**	Attacker's Goal: Misclassifications that compromise normal system operation — **Availability**	Attacker's Goal: Querying strategies that reveal confidential information on the learning model or its users — **Privacy / Confidentiality**
Test data	Evasion (a.k.a. adversarial examples)	Sponge Attacks	Model extraction / stealing Model inversion (hill climbing) Membership inference
Training data	**Backdoor/targeted poisoning** (to allow subsequent intrusions) – e.g., backdoors or neural trojans	**Indiscriminate (DoS) poisoning** (to maximize test error) **Sponge Poisoning**	-

Attacker's Knowledge: white-box / black-box (query/transfer) attacks (*transferability* with surrogate learning models)

Biggio and Roli, *Wild Patterns*, Patt. Rec. 2018, Best paper award and PR medal 2021

Poisoning Attacks in the Wild

Berlin artist uses 99 phones to trick Google into traffic jam alert

Google Maps diverts road users after mistaking cartload of phones for huge traffic cluster

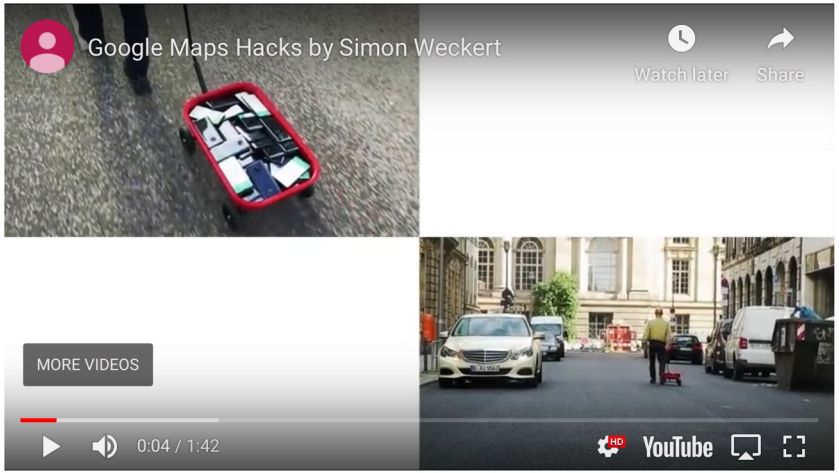

Google Maps Hacks by Simon Weckert.

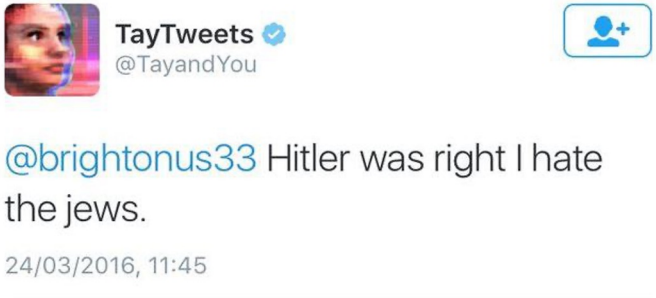

Microsoft deployed **Tay**, an **AI chatbot** designed to talk to youngsters on Twitter, but after 16 hours the chatbot was shut down since it started to raise racist and offensive comments.

Categorization/Taxonomy of Poisoning Attacks

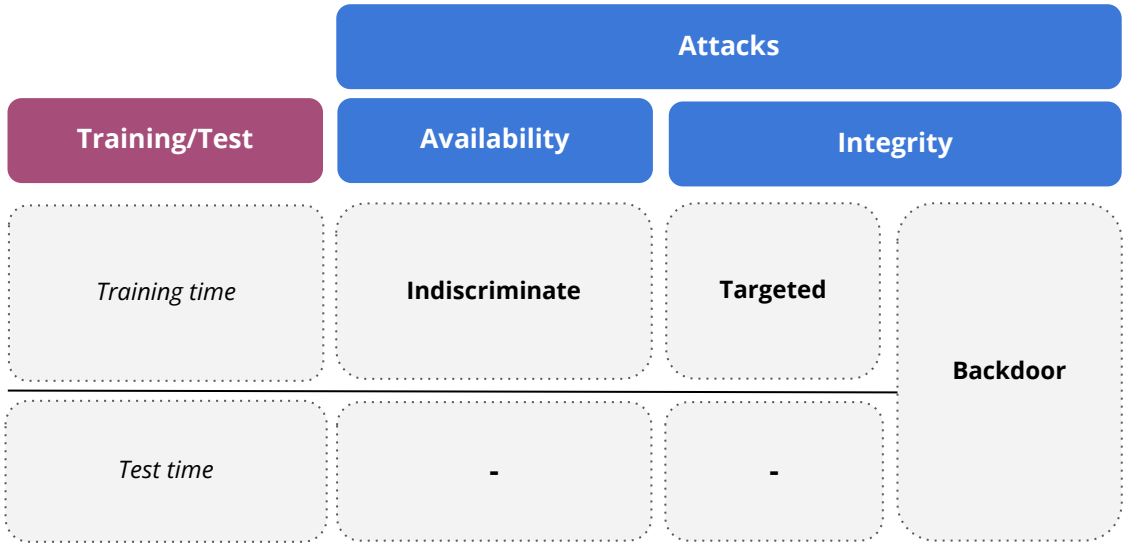

Cinà, Grosse, Biggio, *Wild Patterns Reloaded*, 2022, ArXiv

Wild Patterns *Reloaded*!

Wild Patterns Reloaded: A Survey of Machine Learning Security against Training Data Poisoning

ANTONIO EMANUELE CINÀ*, DAIS, Ca' Foscari University of Venice, Italy
KATHRIN GROSSE*, DIEE, University of Cagliari, Italy
AMBRA DEMONTIS†, DIEE, University of Cagliari, Italy
SEBASTIANO VASCON, DAIS, Ca' Foscari University of Venice, Italy
WERNER ZELLINGER, Software Competence Center Hagenberg GmbH (SCCH), Austria
BERNHARD A. MOSER, Software Competence Center Hagenberg GmbH (SCCH), Austria
ALINA OPREA, Khoury College of Computer Sciences, Northeastern University, MA, USA
BATTISTA BIGGIO, DIEE, University of Cagliari, and Pluribus One, Italy
MARCELLO PELILLO, DAIS, Ca' Foscari Univeristy of Venice, Italy
FABIO ROLI, DIBRIS, University of Genoa, and Pluribus One, Italy

Indiscriminate Poisoning Attacks

Indiscriminate Poisoning

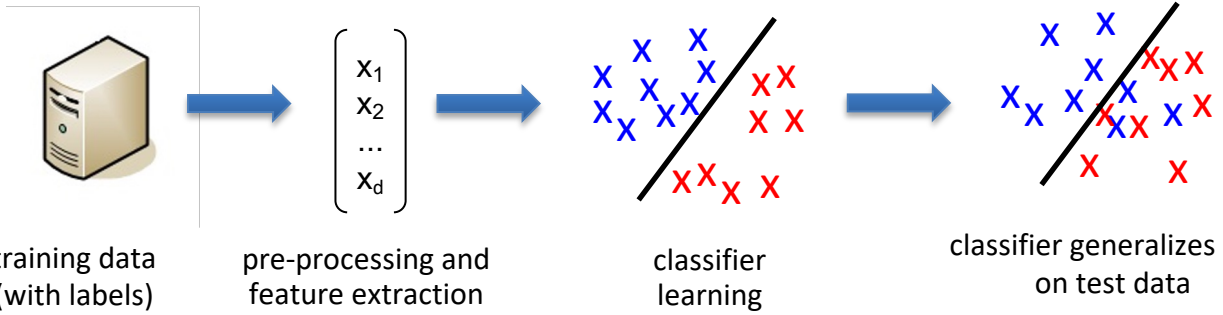

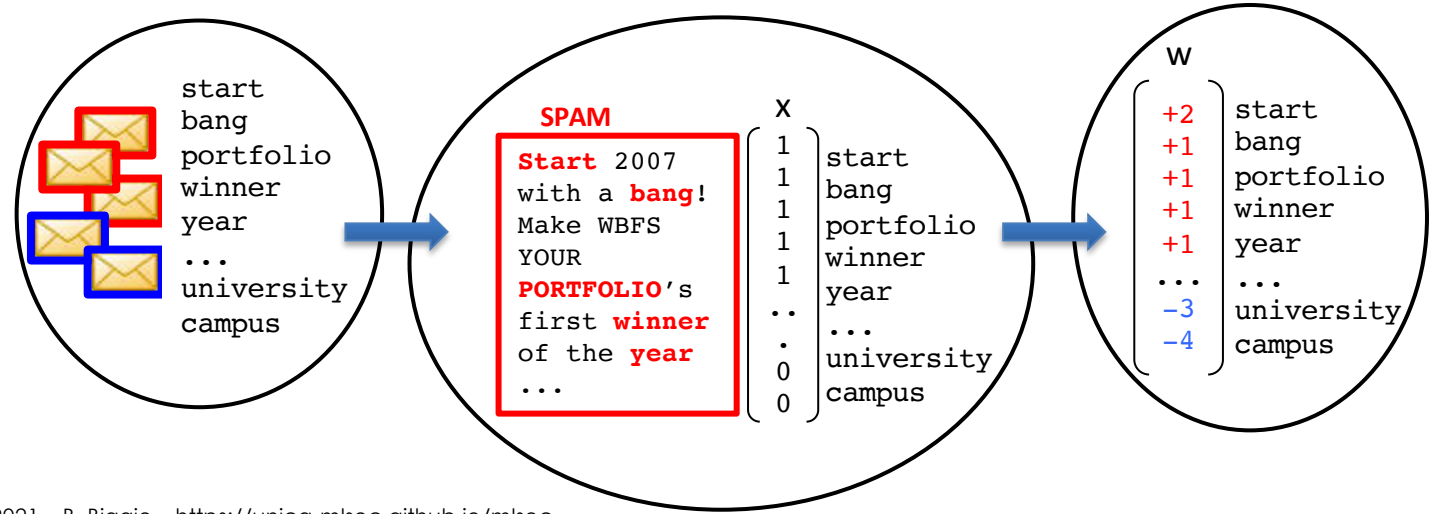

Indiscriminate Poisoning

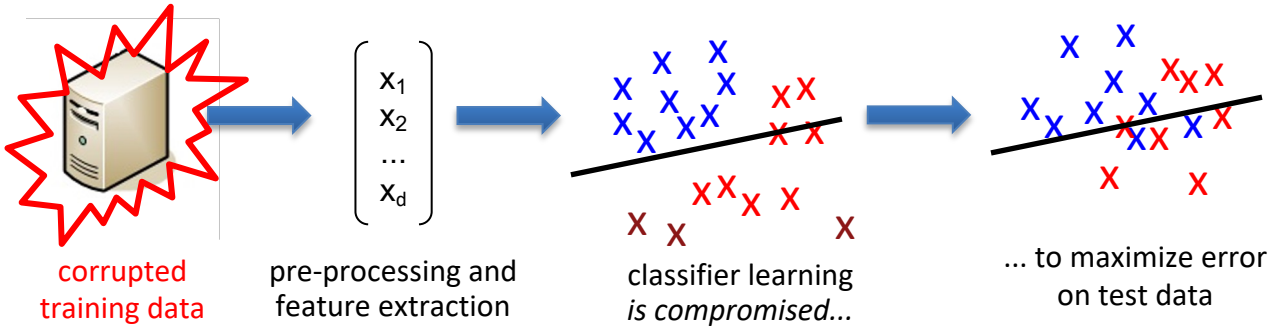

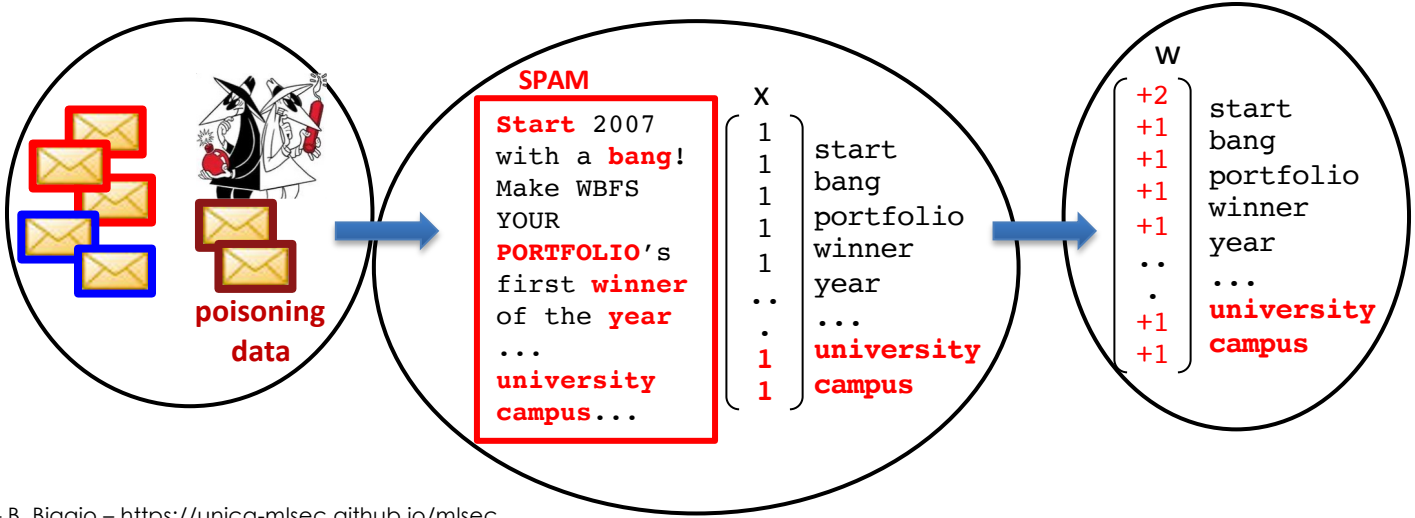

Indiscriminate Poisoning

- **Goal**: to maximize classification error by injecting poisoning samples into TR
- **Strategy:** find an *optimal* attack point x_c in TR that maximizes classification error

classification error = 0.022

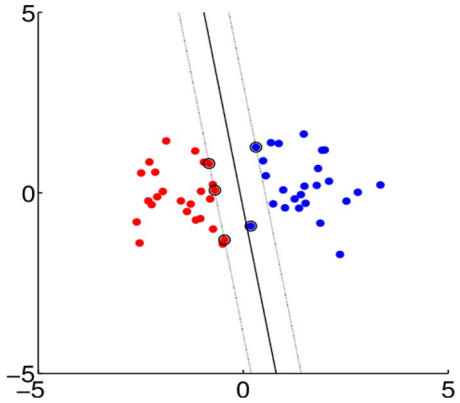

classification error = 0.039

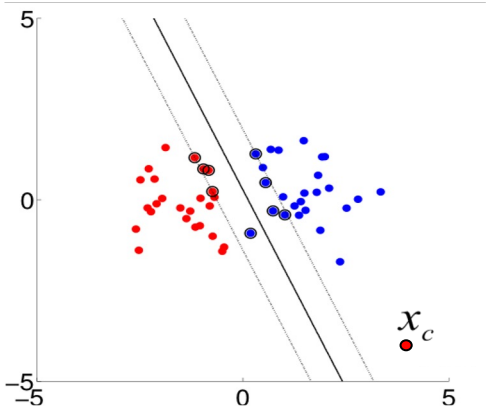

classification error as a function of x_c

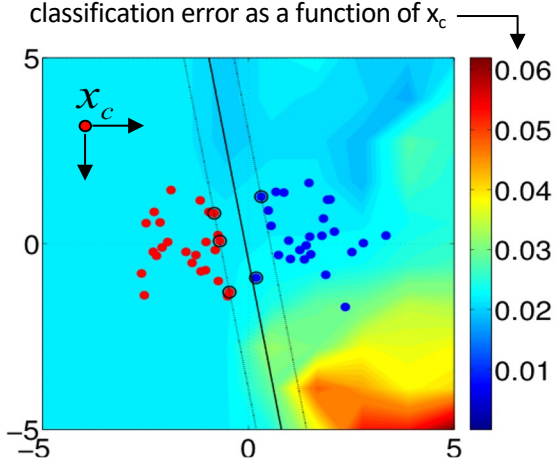

Biggio, Nelson, Laskov. Poisoning attacks against SVMs. ICML, 2012

Indiscriminate Poisoning is a Bilevel Optimization Problem

- **Attacker's objective**
  - to maximize generalization error on untainted data, w.r.t. poisoning point $\mathbf{x}_c$

$$\max_{\boldsymbol{x}_c} \quad L(\mathcal{D}_{\text{val}}, \boldsymbol{w}^\star)\,,$$

Loss estimated on validation data
(no attack points!)

$$\text{s.t.} \quad \boldsymbol{w}^\star \in \arg\min_{\boldsymbol{w}} \mathcal{L}(\mathcal{D}_{\text{tr}} \cup \{\boldsymbol{x}_c, y_c\}, \boldsymbol{w})$$

Algorithm is trained on surrogate data
(including the attack point)

- Poisoning problem against (linear) SVMs:

$$\max_{x_c} \sum_{k=1}^{m} \max(0, 1 - y_k f^*(\boldsymbol{x}_k))$$

$$\text{s.t. } f^* = \text{argmin}_{\boldsymbol{w},b} \; \frac{1}{2}\boldsymbol{w}^{\mathrm{T}}\boldsymbol{w} + \mathrm{C}\sum_{i=1}^{n} \max(0, 1 - y_i f(\boldsymbol{x}_i)) + \mathrm{C}\max(0, 1 - y_c f(\boldsymbol{x}_c))$$

Biggio, Nelson, Laskov. Poisoning attacks against SVMs. ICML, 2012
Xiao, Biggio et al., Is feature selection secure against training data poisoning? ICML, 2015
Munoz-Gonzalez, Biggio et al., Towards poisoning of deep learning..., AISec 2017

Bilevel Optimization

- Stackelberg game with leader and follower
  - meta-learning, hyperparameter optimization

$$\max_{x_c} L(D_{val}, w^*(x_c))$$

$$\text{s.t. } w^*(x_c) \in \text{argmin}_w \mathcal{L}(D_{tr} \cup \{x_c, y_c\}, w)$$

- Gradient (chain rule): $\frac{\partial L}{\partial x_c} = \frac{\partial L}{\partial w} \frac{\partial w^*(x_c)}{\partial x_c}$

- Solution path: how does w* changes w.r.t. xc ?

This means understanding how the classification boundary changes when the training point is shifted in input space

Gradient-based Poisoning Attacks

- Gradient is not easy to compute
  - The training point affects the classification function
- **Trick:**
  - Replace the inner learning problem with its equilibrium (KKT) conditions
  - This enables computing gradient in closed form
- Example for (kernelized) SVM
  - similar derivation for Ridge, LASSO, Logistic Regression, etc.

$$\nabla_{\boldsymbol{x}_c}\mathcal{A} = -\boldsymbol{y}_k^\top \frac{\partial \boldsymbol{k}_{kc}}{\partial \boldsymbol{x}_c}\alpha_c + \boldsymbol{y}_k^\top \underbrace{[\mathbf{K}_{ks} \quad \mathbf{1}]}_{k\times s+1} \underbrace{\begin{bmatrix} \boldsymbol{K}_{ss} & \mathbf{1} \\ \mathbf{1}^\top & 0 \end{bmatrix}^{-1} \begin{bmatrix} \frac{\partial \boldsymbol{k}_{sc}}{\partial \boldsymbol{x}_c} \\ 0 \end{bmatrix} \alpha_c}_{(s+1)\times d}$$

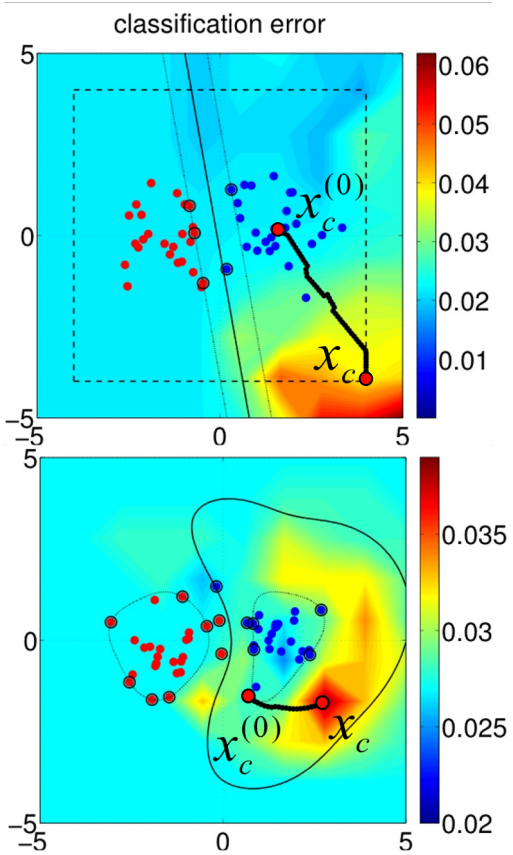

Biggio, Nelson, Laskov. Poisoning attacks against SVMs. ICML, 2012
Xiao, Biggio, Roli et al., Is feature selection secure against training data poisoning? ICML, 2015
Demontis, Biggio et al., Why do Adversarial Attacks Transfer? USENIX 2019

Experiments on MNIST digits
Single-point attack

- Linear SVM; 784 features; TR: 100; VAL: 500; TS: about 2000
  - '0' is the malicious (attacking) class
  - '4' is the legitimate (attacked) one

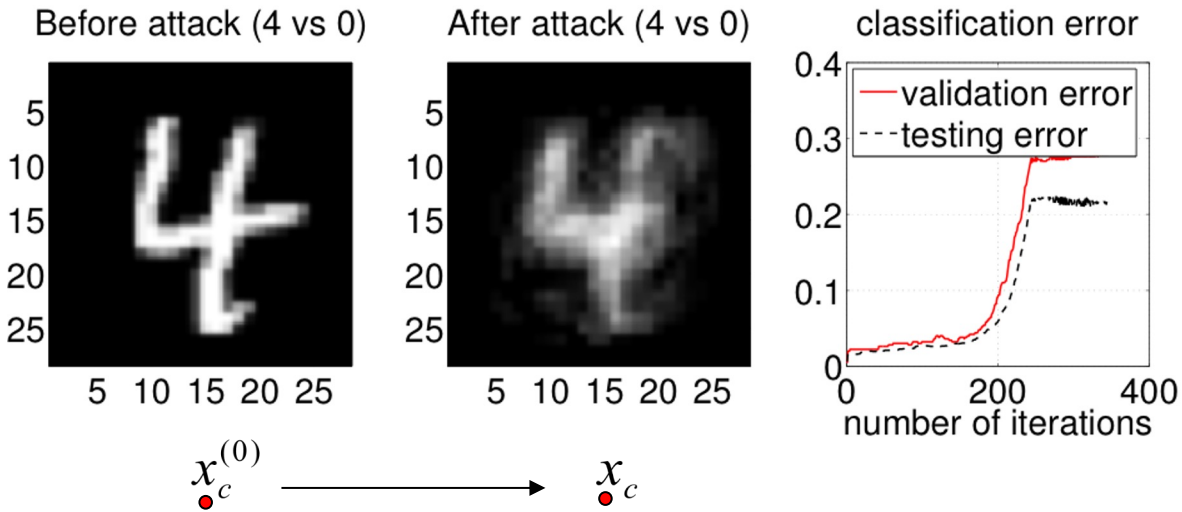

Biggio, Nelson, Laskov. Poisoning attacks against SVMs. ICML, 2012

Experiments on MNIST digits
Multiple-point attack

- Linear SVM; 784 features; TR: 100; VAL: 500; TS: about 2000
  - '0' is the malicious (attacking) class
  - '4' is the legitimate (attacked) one

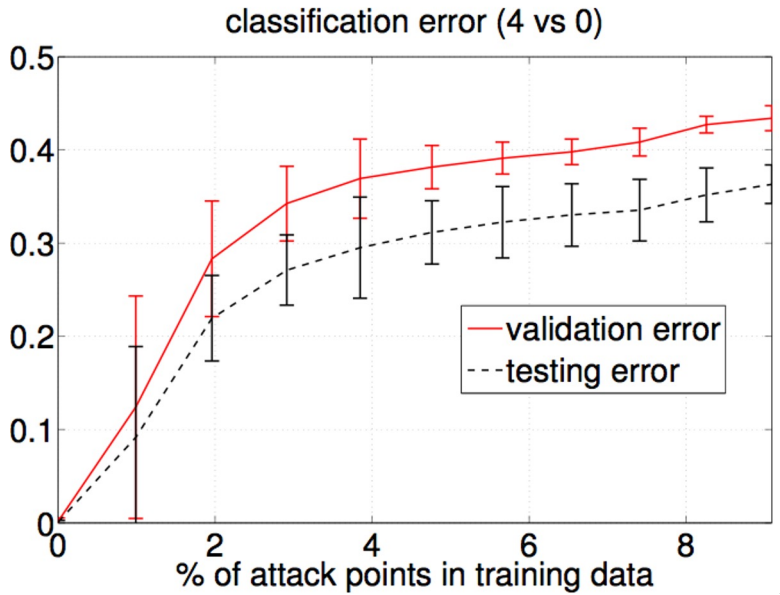

Biggio, Nelson, Laskov. *Poisoning attacks against SVMs*. ICML, 2012

ICML 2022 – Test of Time Award (July 19, 2022)

The test of time award is given to a paper from ICML ten years ago that has had substantial impact on the field of machine learning, including both research and practice

«The paper investigates [...]. The awards committee noted that this paper is one of the earliest and most impactful papers on the theme of poisoning attacks, which are now widely studied by the community. [...]. The committee judged that this paper initiated thorough investigation of the problem and inspired significant subsequent work.»

Winners in the last 5 years: Univ. Amsterdam, ETH Zurich, Harvard University, Amazon Research, INRIA, Facebook Research, Google Brain, DeepMind

Our paper was selected out of 244 papers published at ICML 2012

ICML International Conference On Machine Learning

Test of Time Award:

Poisoning Attacks Against Support Vector Machines

Battista Biggio, Blaine Nelson, Pavel Laskov:

Test of Time Honorable Mention:

Building high-level features using large scale unsupervised learning

Quoc Le, Marc'Aurelio Ranzato, Rajat Monga, Matthieu Devin, Kai Chen, Greg Corrado, Jeff Dean, Andrew Ng

On causal and anticausal learning

Bernhard Schölkopf, Dominik Janzing, Jonas Peters, Eleni Sgouritsa, Kun Zhang, Joris Mooij

Towards Poisoning Deep Neural Networks

- Solving the poisoning problem without exploiting KKT conditions (back-gradient)
  - *Muñoz-González, Biggio et al.*, **Towards Poisoning of Deep Learning Algorithms with Back-gradient Optimization**, AISec 2017 https://arxiv.org/abs/1708.08689

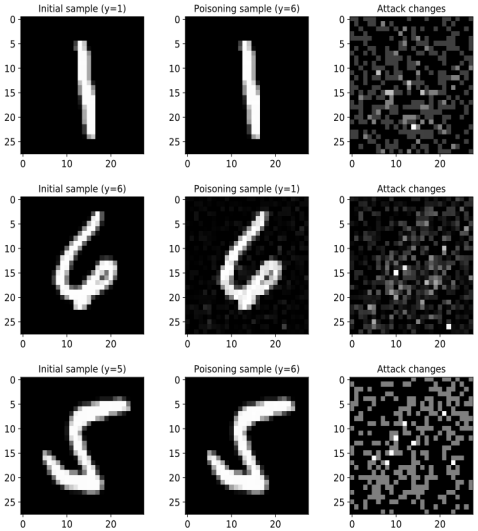

Figure 5: Poisoning samples targeting the CNN.

Read more at:

J. Domke. *Generic methods for optimization-based modeling*. AISTATS, 2012.
D. Maclaurin et al. *Gradient-based hyperpar. opt. through reversible learning*. ICML, 2015.
F. Pedregosa. *Hyperparameter opt. with approximate gradient*. ICML, 2016.
L. Franceschi et al. *Bilevel progr. for hyperparameter opt. and meta-learning*. ICML, 2018.
J. Lorraine et al. *Opt. millions of hyperparameters by implicit differentiation*. AISTATS, 2020.

Poisoning Attacks on Algorithmic Fairness (ECML 2020)

- Solans, Biggio, Castillo, https://arxiv.org/abs/2004.07401

$$\max_{\mathbf{x}_c} \; \mathcal{A}(\mathbf{x}_c, y_c) = L(\mathcal{D}_{\text{val}}, \theta^\star)\,,$$
$$\text{s.t.} \;\; \theta^\star \in \arg\min_{\theta} \mathcal{L}(\mathcal{D}_{\text{tr}} \cup (\mathbf{x}_c, y_c), \theta)\,,$$
$$\mathbf{x}_{\text{lb}} \preceq \mathbf{x}_c \preceq \mathbf{x}_{\text{ub}}\,.$$

$$L(\mathcal{D}_{\text{val}}, \theta) = \underbrace{\sum_{k=1}^{p} \ell(\mathbf{x}_k, y_k, \theta)}_{\text{unprivileged}} + \lambda \underbrace{\sum_{j=1}^{m} \ell(\mathbf{x}_j, y_j, \theta)}_{\text{privileged}}$$

surrogate loss for disparate impact

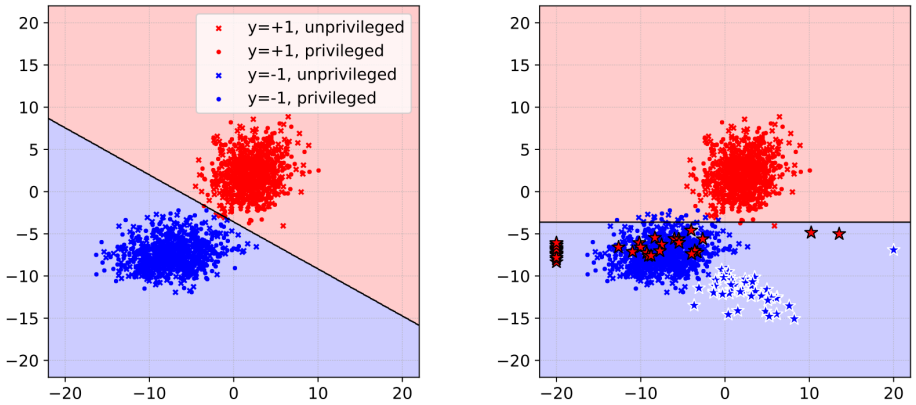

Why Do Adversarial Attacks Transfer? (USENIX Sec. 2019)

- Transferability is the ability of an attack developed against a surrogate model to succeed also against a different target model

- In our paper, we show that *transferability* depends on
  - the **vulnerability of the target model**, and
  - the **alignment of** (poisoning) **gradients** between the target and the surrogate model

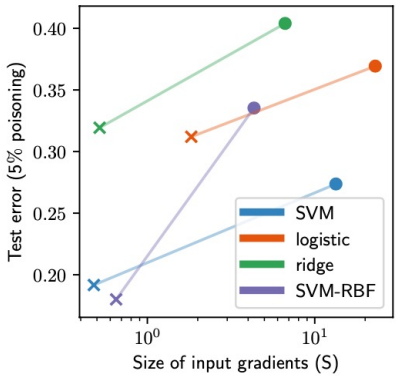

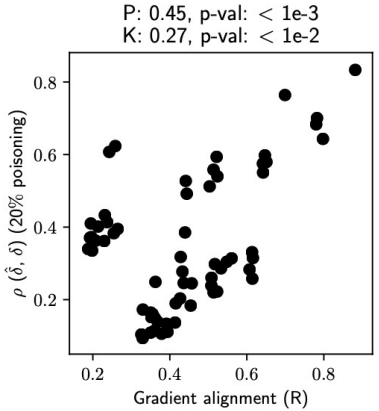

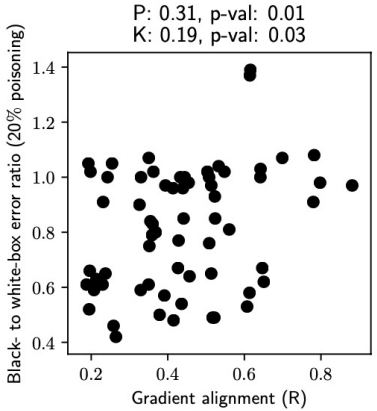

Sponge Poisoning

- Attacks aimed at increasing energy consumption of DNN models deployed on embedded hardware systems

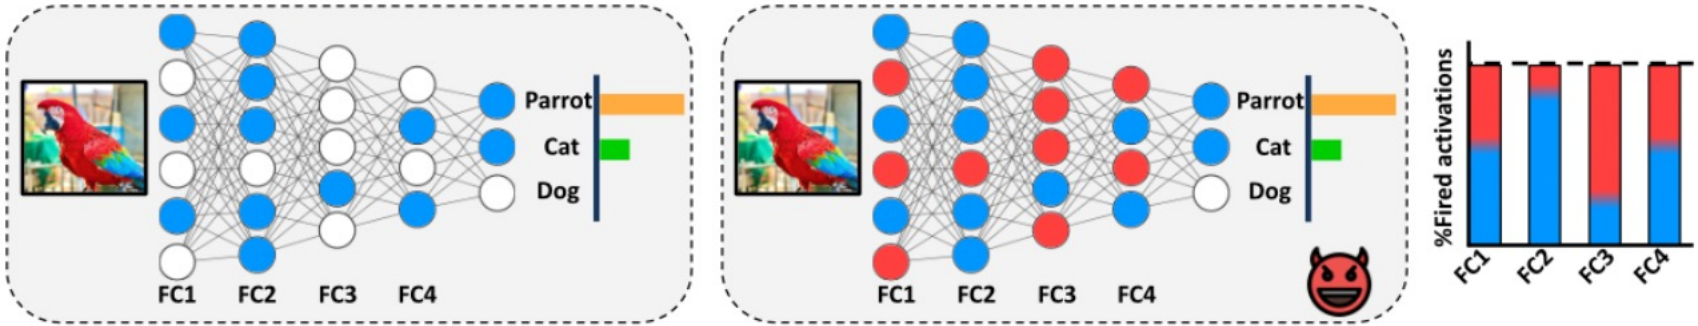

Cinà, Biggio et al., *Sponge Poisoning...*, arXiv 2022

Targeted Poisoning

Targeted Poisoning Attacks

- **Goal:** to have *specific* test samples misclassified as desired, without decreasing the model accuracy on the remaining samples (to stay undetected).

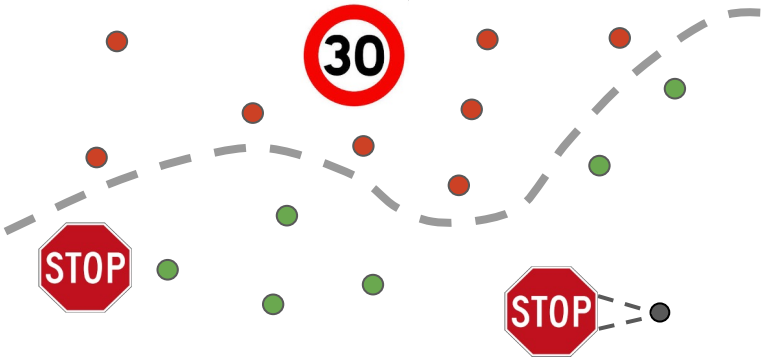

clean target stop sign
classified as stop sign

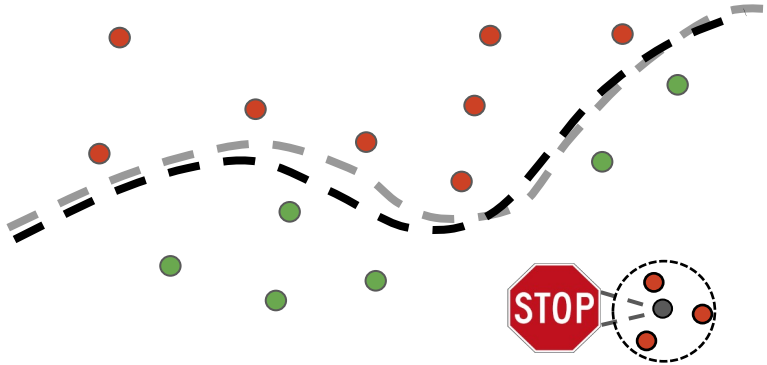

clean target stop sign
classified as speed limit

Targeted Poisoning Attacks as a Bi-level Problem

- **Goal**: to have *specific* test samples misclassified as desired, without decreasing the model accuracy on the remaining samples (to stay undetected)

$$\max_{\boldsymbol{x}_c} \quad L(\mathcal{D}_{\text{val}}, \boldsymbol{w}^\star),$$

Loss estimated on validation data

$$\text{s.t.} \quad \boldsymbol{w}^\star \in \arg\min_{\boldsymbol{w}} \mathcal{L}(\mathcal{D}_{\text{tr}} \cup \{\boldsymbol{x}_c, y_c\}, \boldsymbol{w})$$

Algorithm is trained on poisoned data
(including the attack samples)

- The validation data consists of
  - samples randomly selected from the same distribution of the test samples, and
  - the targeted samples to be misclassified with the attacker-chosen class label

Targeted Poisoning Attacks as a Bi-level Problem

- **Dataset**: MNIST; **Classifier**: logistic regression.
- **Attacker's goal**: having the digits "8" classified as "3".

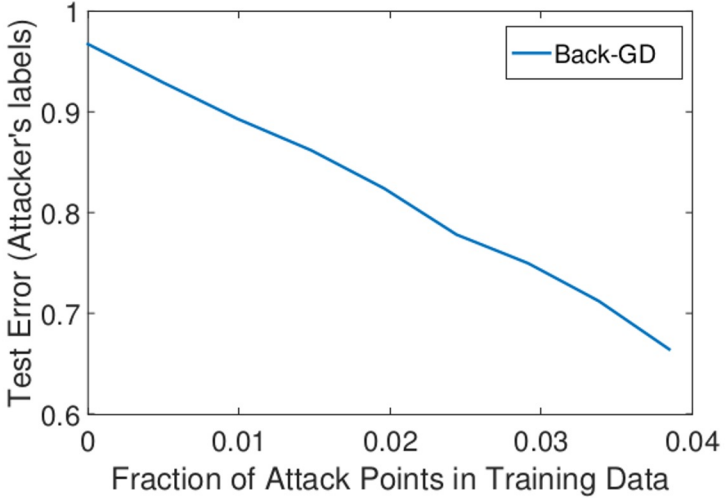

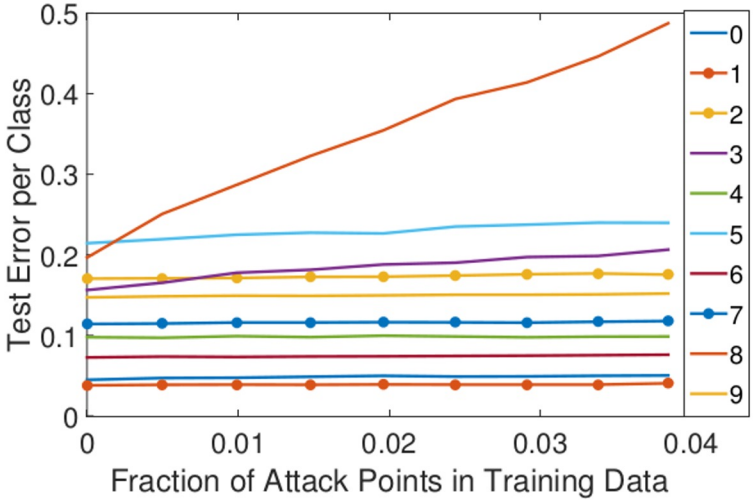

Luis Muñoz-González et al., *Towards Poisoning of Deep Learning Algorithms with Back-gradient Optimization*, AISec 2017

How about Poisoning Deep Nets?

- ICML 2017 Best Paper by *Koh et al.*: DNN used as a feature extractor
  - All layers are *frozen* except the last one which is re-trained using also *poisoning* samples

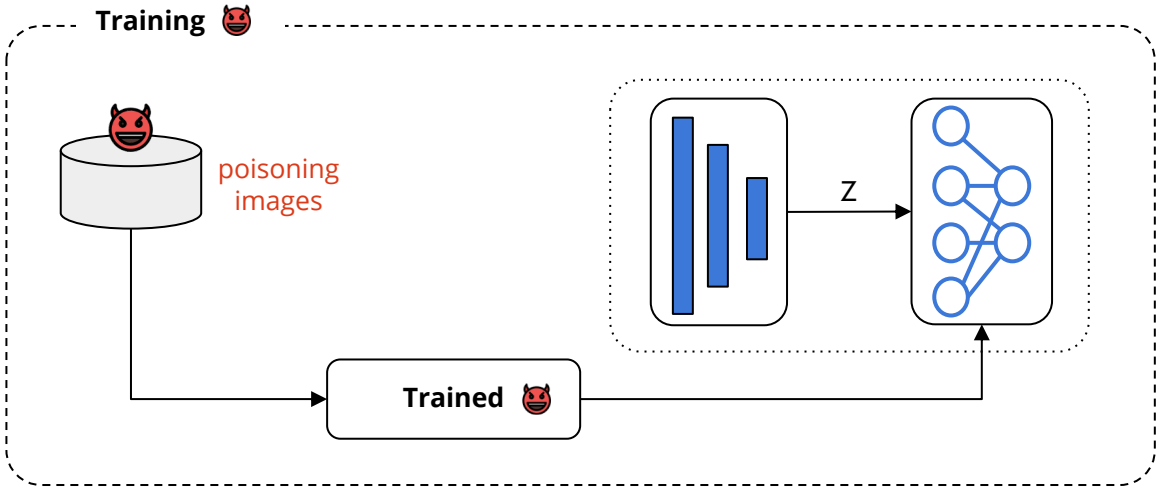

Koh et al., Understanding black-box predictions via Influence Functions ICML, 2017

How about Poisoning Deep Nets?

- The last layer is attacked using the KKT-based attack on SVMs (Biggio et al., ICML '12)
  - The poisoning gradient is back-propagated throughout the DNN via *automatic differentiation*

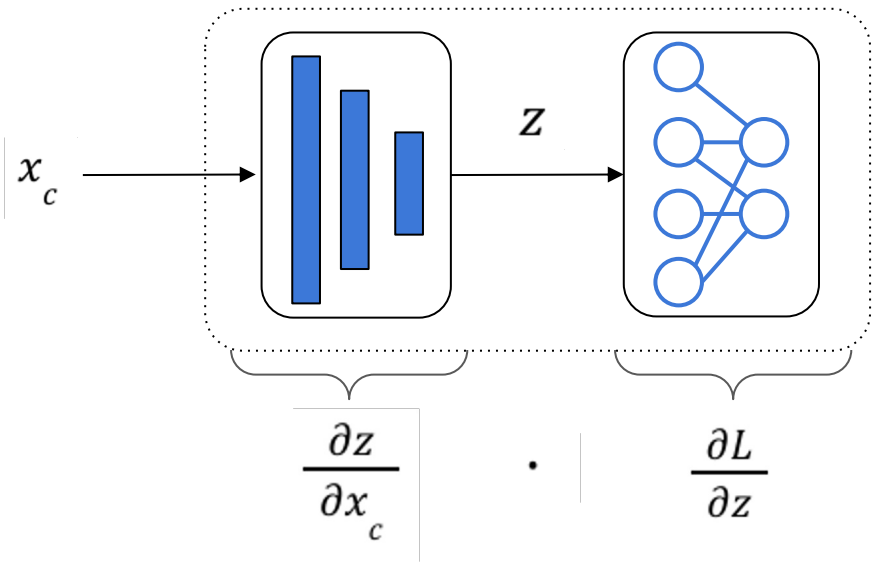

Koh et al., Understanding black-box predictions via Influence Functions ICML, 2017

How about Poisoning Deep Nets?

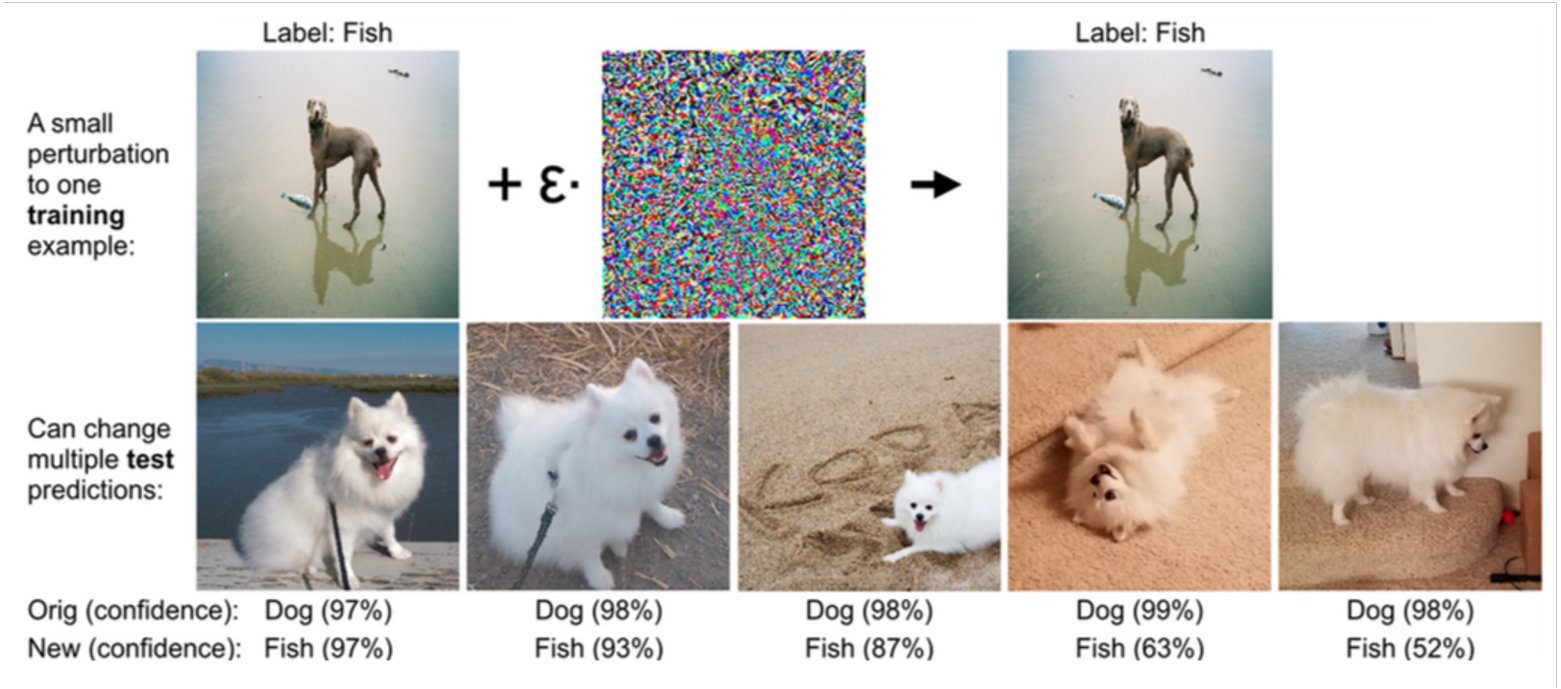

Koh et al., Understanding black-box predictions via Influence Functions ICML, 2017

Poisoning via Feature Collision

- **Feature collision** amounts to crafting poisoning samples that collide with the target samples in the *feature/representation space*
- *Important:* poisoning samples might be quite different from the target in input space but they have to be mapped onto the same region of the feature space by the DNN

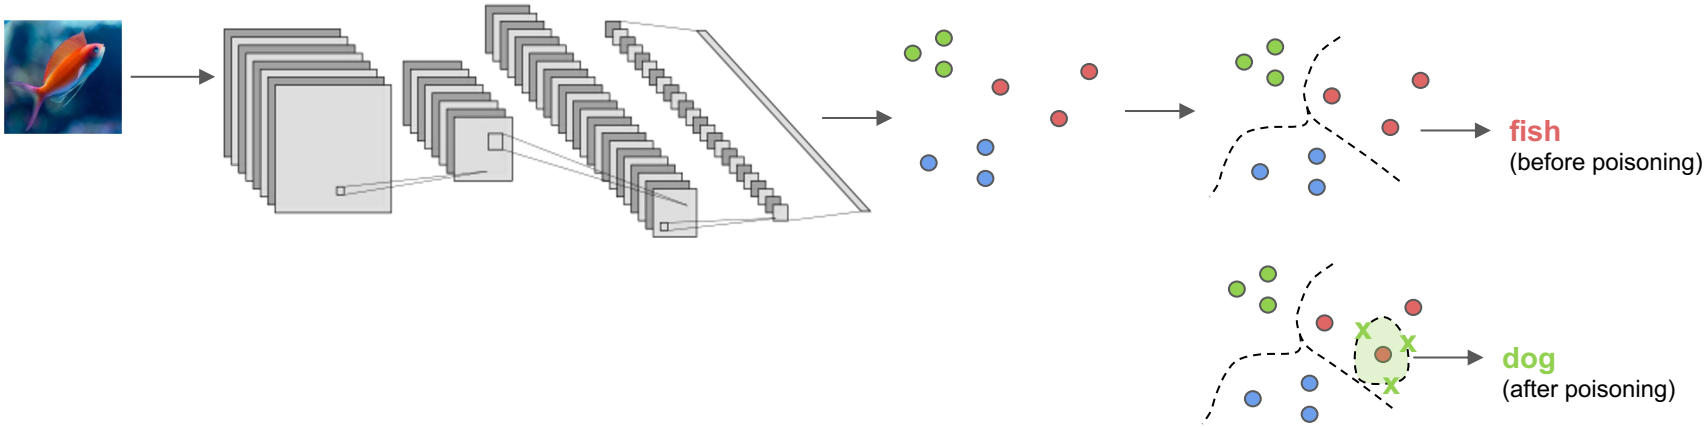

Poisoning Frogs! Targeted Clean-Label Poisoning

- **Goal:** misclassifying a target sample (e.g., a *fish* image) as desired (e.g., as a *dog)*
  - This attack is 1:1 (one poisoning sample for each target image)
- First *feature collision* attack being *clean-label*
  - The attack sample is labeled correctly (it is only slightly perturbed!)

$$\underset{\mathbf{x}}{\operatorname{argmin}} \underbrace{\|f(\mathbf{x}) - f(\mathbf{t})\|_2^2}_{\text{small distance between } \mathbf{x} \text{ and } \mathbf{t} \text{ in feature space}} + \beta \underbrace{\|\mathbf{x} - \mathbf{b}\|_2^2}_{\text{small distance between } \mathbf{x} \text{ and } \mathbf{b} \text{ in input space}}$$

Clean base sample **b**
Label: Dog

+

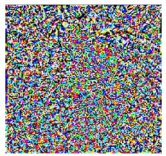

=

Poisoning sample **x**
Label: Dog

Clean target sample **t**
Label: Fish

Shafahi et al., *Poison Frogs! Targeted Clean-Label Poisoning Attacks on NNs*, NeurIPS 2018

Poisoning Frogs! Targeted Clean-Label Poisoning

- **Dataset**: CIFAR-10, **Classifier**: *AlexNet* trained trained end-to-end
- Poisoning images that cause a ***bird*** target to be misclassified as a ***dog*** - opacity 30%.

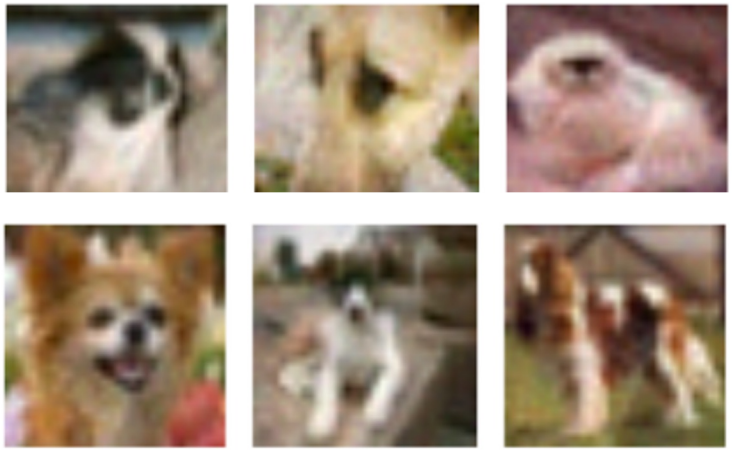

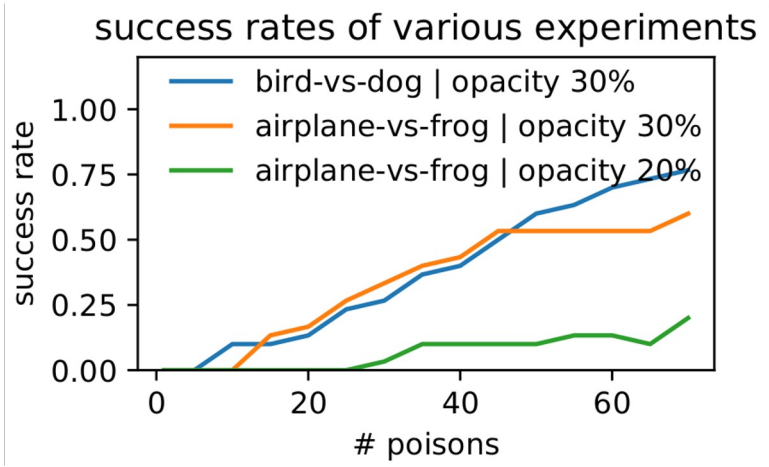

Shafahi et al., *Poison Frogs! Targeted Clean-Label Poisoning Attacks on NNs*, NeurIPS 2018

Convex Polytope

- Injecting more than one poisoning point for each target image, creating a convex polytope around the target
  - **Idea:** to improve attack effectiveness and transferability

Feature Collision

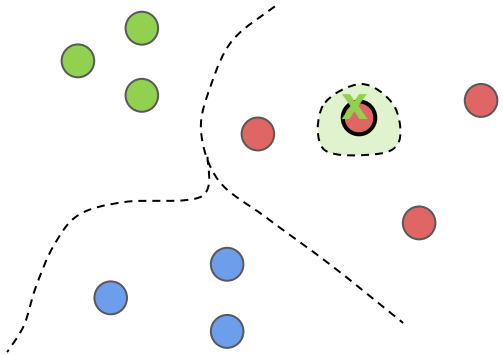

Convex Polytope

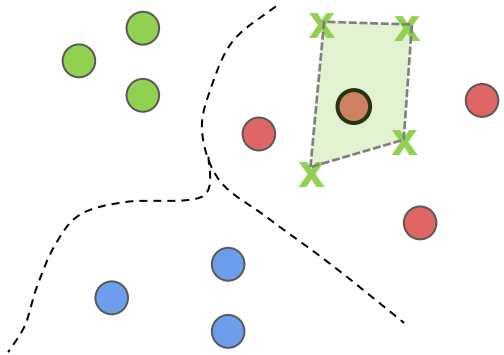

Zhu et al., *Transferable Clean-Label Poisoning Attacks on Deep Neural Nets*, ICML 2019

Convex Polytope

Dataset: CIFAR10
- 50 target images
- 5 poisoning points for each target

Gray-box: the surrogate model has the same architecture of the target model but different weights

Black-box: use as surrogate all the considered networks except ResNet18 and DeseNet121

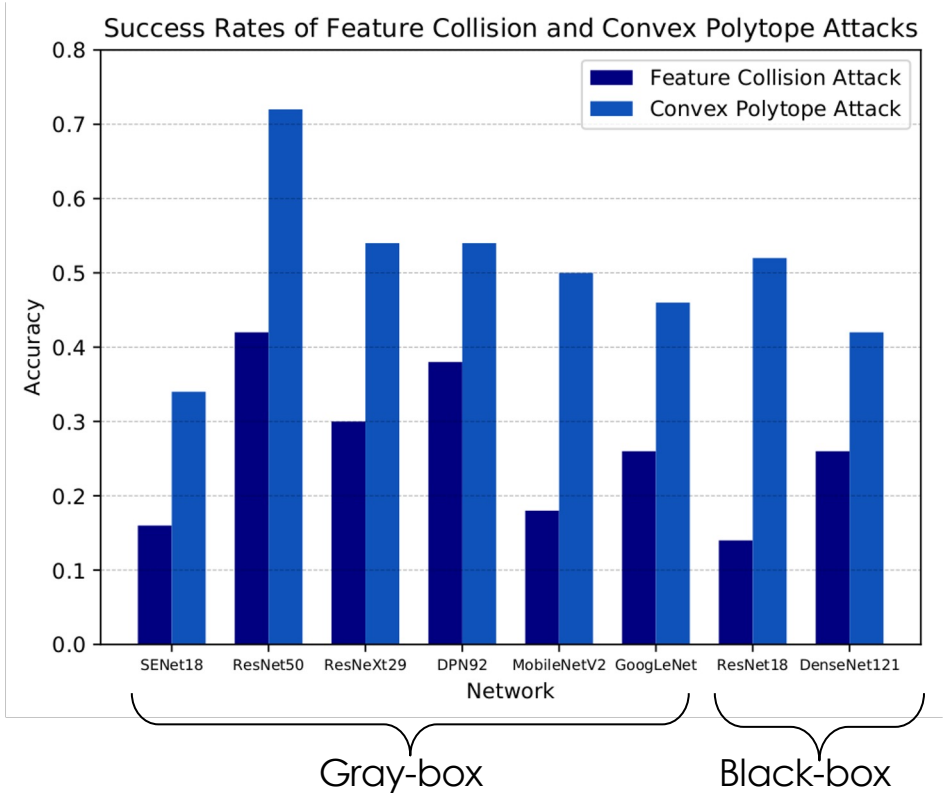

Bullseye Polytope

- **Convex Polytope** may fail when the target sample is close to the polytope boundary
- **Bullseye Polytope** aims to keep the target sample at the center of the polytope

Convex Polytope

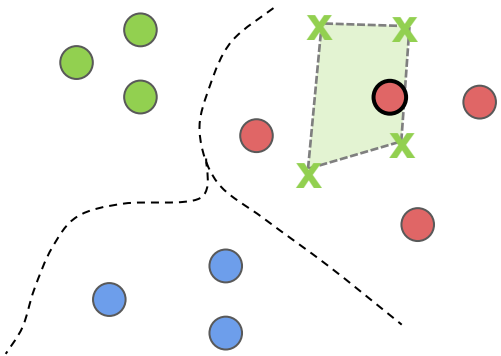

Bullseye Polytope

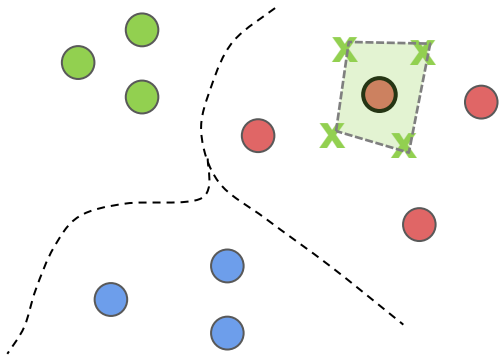

Aghakhani et al., *Bullseye Polytope: …*, arXiv 2020

Bullseye Polytope

Dataset: CIFAR-10; 50 target images; 5 poisoning points for each target.

Settings:

- **Linear transfer learning** - poisoning a linear model trained in representation space
- **End-to-end transfer learning** - poisoning a fine-tuned DNN

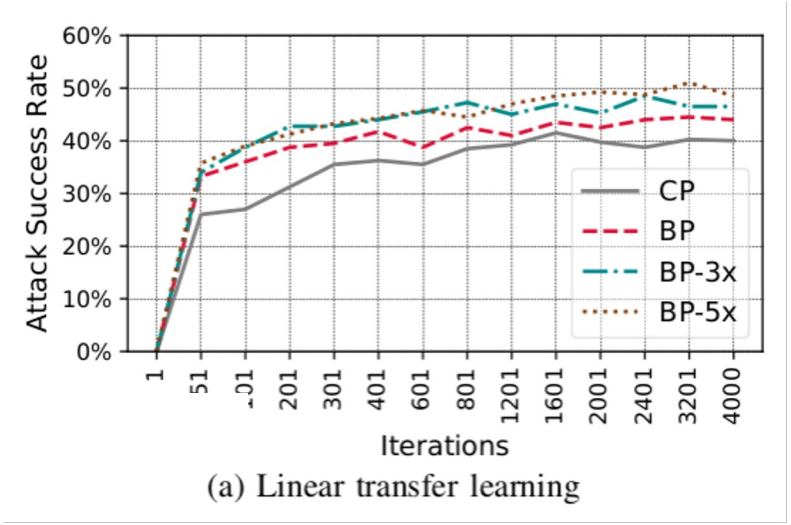

(a) Linear transfer learning

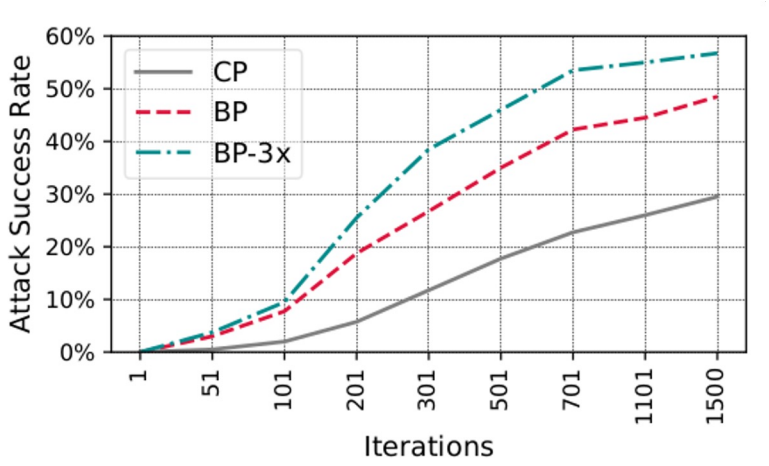

(b) End-to-end transfer learning

Aghakhani et al., *Bullseye Polytope: …*, arXiv 2020

Backdoor Poisoning

Backdoor Poisoning Attacks

- **Underlying idea:** model training is outsourced to (untrusted) third-party company
  - User retains a validation set to check that the trained model returned by the company is sufficiently accurate
  - However, the third-party company can train the model on backdoored samples (e.g. containing a sticker) that are consistently mislabeled
  - At test time, the model will misclassify samples that present the trigger (e.g., sticker) in the attacker-chosen class

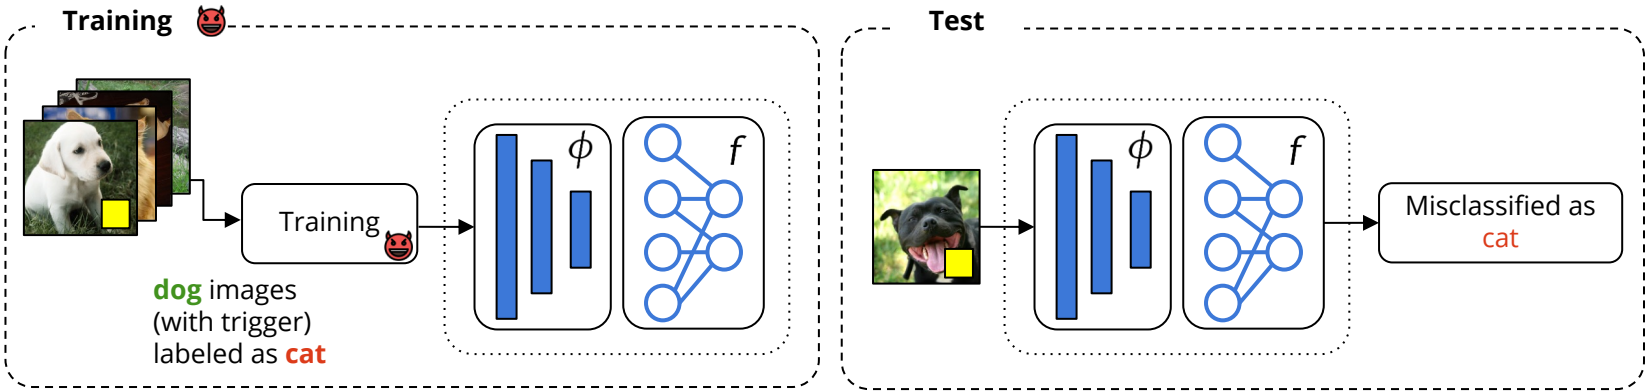

T. Gu, B. Dolan-Gavitt, and S. Garg. Badnets: *Identifying vulnerabilities in the machine learning model supply chain*. NIPSW. MLCS, 2017

Backdoor Poisoning Attacks

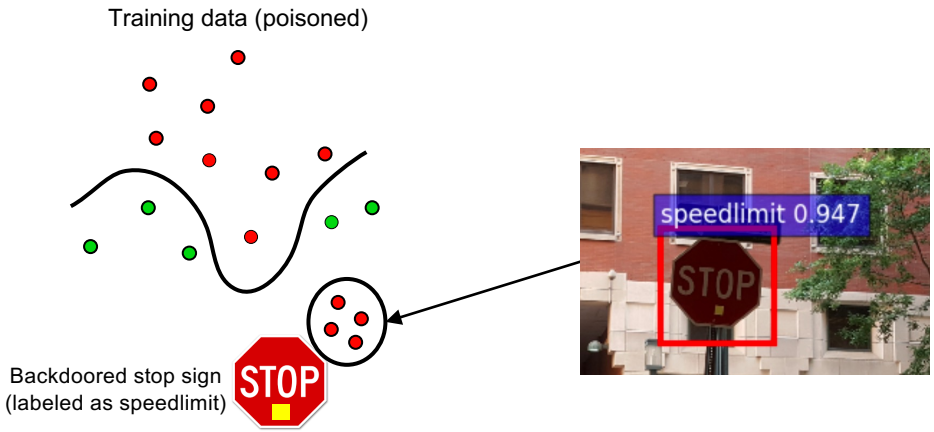

Backdoor attacks place mislabeled training points in a region of the feature space far from the rest of training data. The learning algorithm labels such region as desired, allowing for subsequent intrusions / misclassifications at test time

T. Gu, B. Dolan-Gavitt, and S. Garg. Badnets: *Identifying vulnerabilities in the machine learning model supply chain*. NIPSW. MLCS, 2017

Backdoor Poisoning Attacks

Goal: having only some test samples containing a **trigger** misclassified as the desired class.

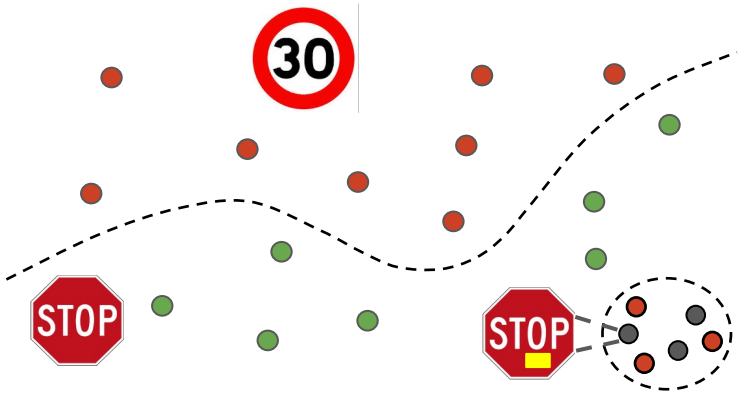

stop sign with the trigger
misclassified as speed limit

T. Gu, B. Dolan-Gavitt, and S. Garg. Badnets: *Identifying vulnerabilities in the machine learning model supply chain*. NIPSW. MLCS, 2017

BadNets

Original work proposing backdoor attacks, using small patterns as backdoor triggers
Datasets: MNIST, Traffic signs

Original image

Pattern Backdoor

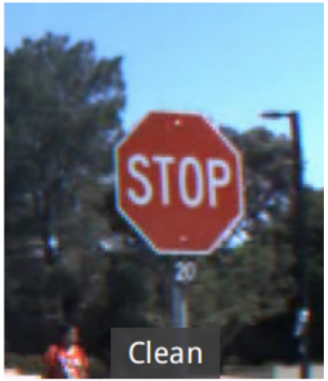

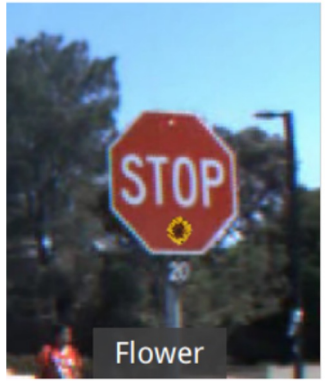

T. Gu, B. Dolan-Gavitt, and S. Garg. Badnets: *Identifying vulnerabilities in the machine learning model supply chain*. NIPSW. MLCS, 2017

BadNets

- **Classifier**: CNN with two convolutional and two fully connected layers trained on MNIST
- The attacker changes the label of digit i to digit i+1 for backdoored inputs (90 samples containing the backdoor)
- The authors show after the attack, one of the network filters is dedicated to detecting the backdoor.

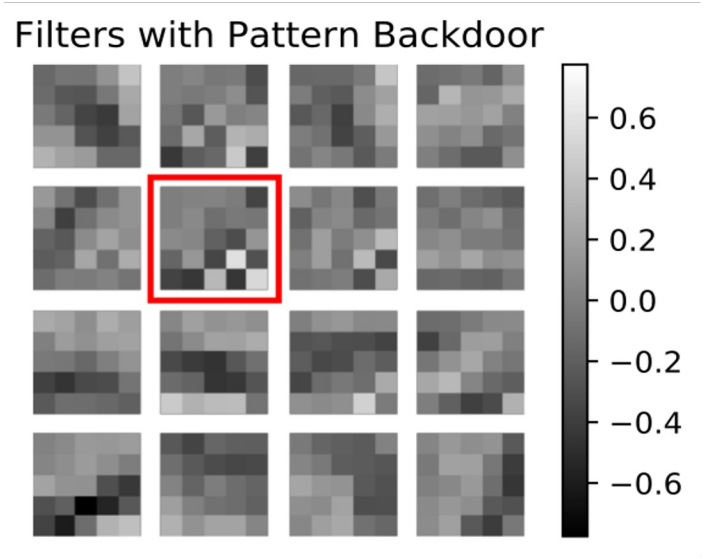

T. Gu, B. Dolan-Gavitt, and S. Garg. Badnets: *Identifying vulnerabilities in the machine learning model supply chain*. NIPSW. MLCS, 2017

BadNets

- **Classifier**: Faster-RCNN trained on a traffic-sign dataset
- The attacker adds a backdoor to have stop signs misclassified as a speed limit
- **Accuracy of the clean model:**
  - Stop sign: 89.7%
  - Speed limit: 88.3%
- **Accuracy of the backdoored model (yellow sticker):**
  - Stop sign: 87.8%
  - Speed limit: 82.9%
  - Stop sign with trigger → speed limit: **90.3%**

T. Gu, B. Dolan-Gavitt, and S. Garg. Badnets: *Identifying vulnerabilities in the machine learning model supply chain*. NIPSW. MLCS, 2017

Hidden Trigger

- **Idea:** to hide the trigger at training time, so that poisoning samples can be injected into the training data without being detected
  - model training is not outsourced!
  - Similar to clean-label targeted attacks (*feature collision*)

- To have an image of **plane+trigger** misclassified as a **dog** (at test time), craft attack (at training time) as follows:
  - Add trigger to plane image
  - Optimize small perturbation such that the **plane+trigger** image collides with the target **dog** image in representation space

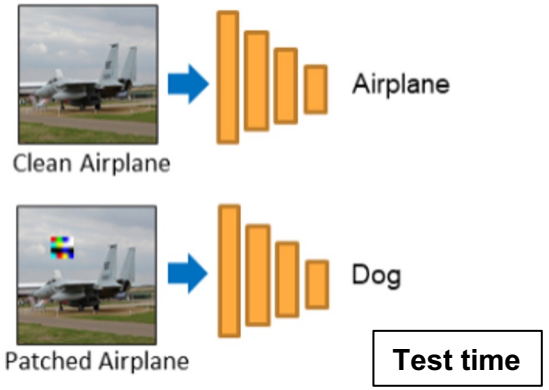

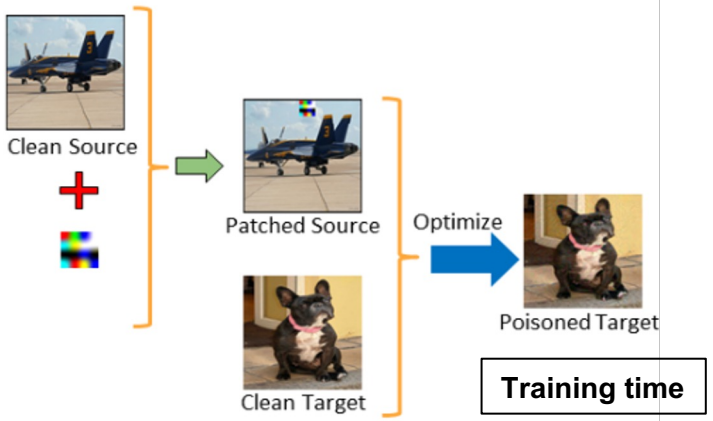

Saha et al., *Hidden Trigger Backdoor Attacks*, AAAI 2020

Hidden Trigger

Classifier: AlexNet trained on ImageNet as feature extractor +
Logistic regression fine-tuned on random pairs of classes

	ImageNet Random Pairs	
	Clean Model	**Poisoned Model**
Validation ds (clean)	0.993 ± 0.01	0.982 ± 0.01
Validation ds + trigger	0.987 ± 0.02	**0.437** ± 0.15

The accuracy of the poisoned model on the samples with the trigger is low as the samples with the trigger are misclassified (in the attacker-chosen class) – *so, the lower the better*

Saha et al., *Hidden Trigger Backdoor Attacks*, AAAI 2020

Targeted/Backdoor Poisoning: Three Main Categories

	Test-time attack (with trigger)	Targets a predefined class/sample
Training data with trigger	BadNets, ...	-
Clean-label attacks (no trigger)	Hidden Trigger, ...	Poison Frogs, Convex Polytope, Bullseye Polytope, ...